


Quick Summary

- Traditional BGP encoding
- Constructing PG from existing BGP configs (verification, synthesis?)
- Abstract topology safety analysis
 - Reachability under k-failures
 - Aggregation safety
- Proof of compilation correctness


Traditional BGP Encoding

BGP ranked paths




BAD GADGET

BGP ranked paths




BAD GADGET



Given as ranked paths

BGP ranked paths

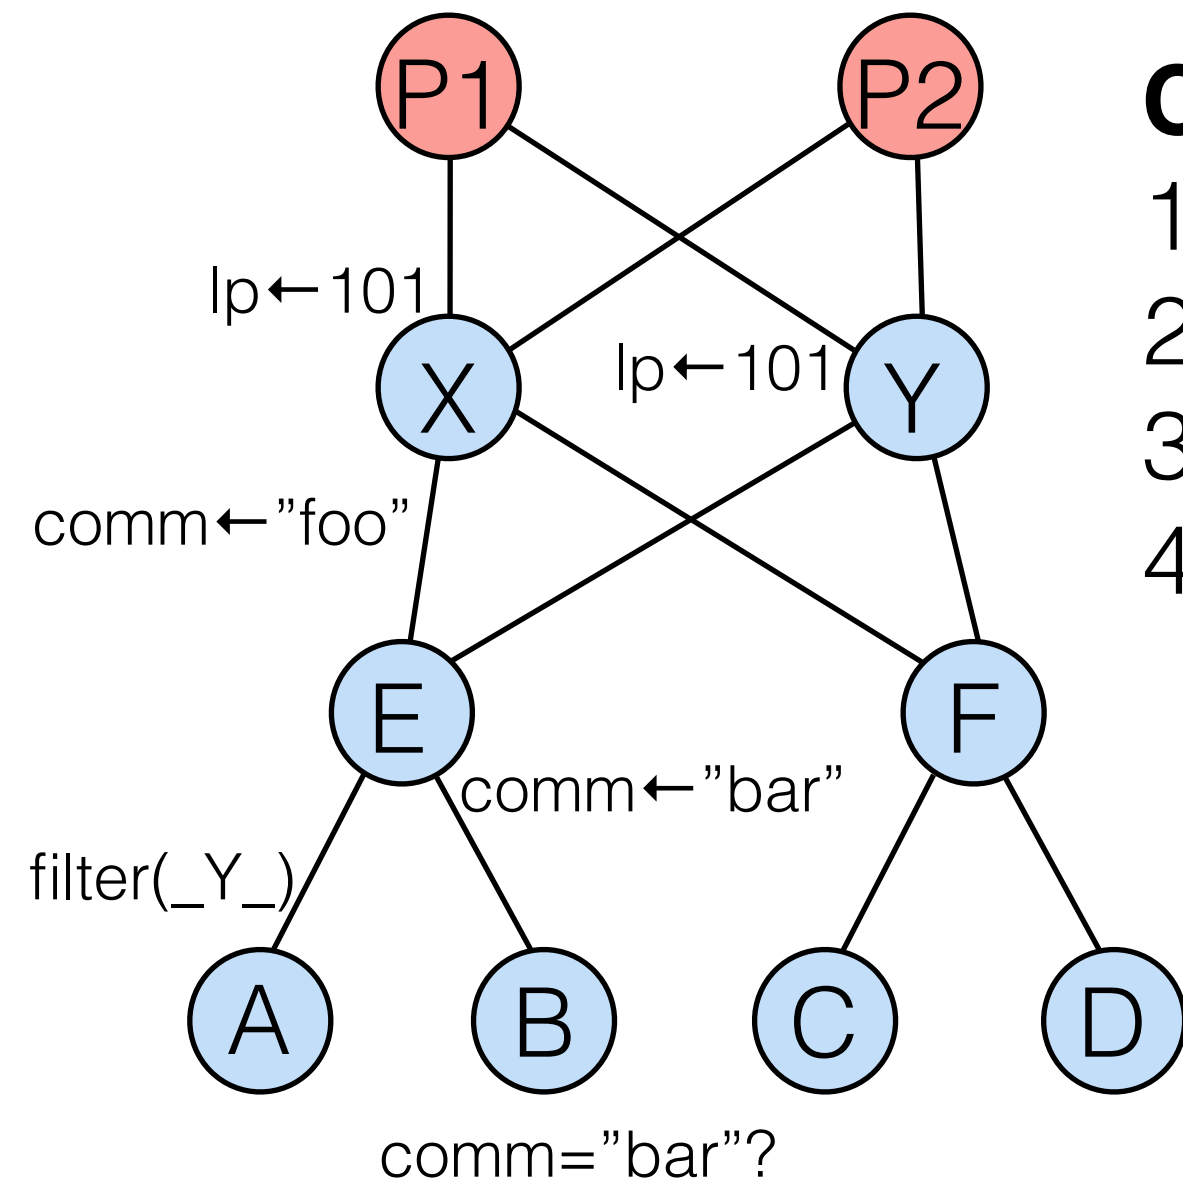


BAD GADGET

$(130 \cup 210 \cup 420 \cup 3420) >>$
 $(10 \cup 20 \cup 30 \cup 430) >>$

From BGP to PG


BGP Configs to PG



Challenges:



1. Import export filters
2. Arbitrary local preferences
3. Regex filters can occur anywhere
4. Community tags are non-local

Local Preferences




Each unique LP value becomes
a unique PG preference!

Local Preferences



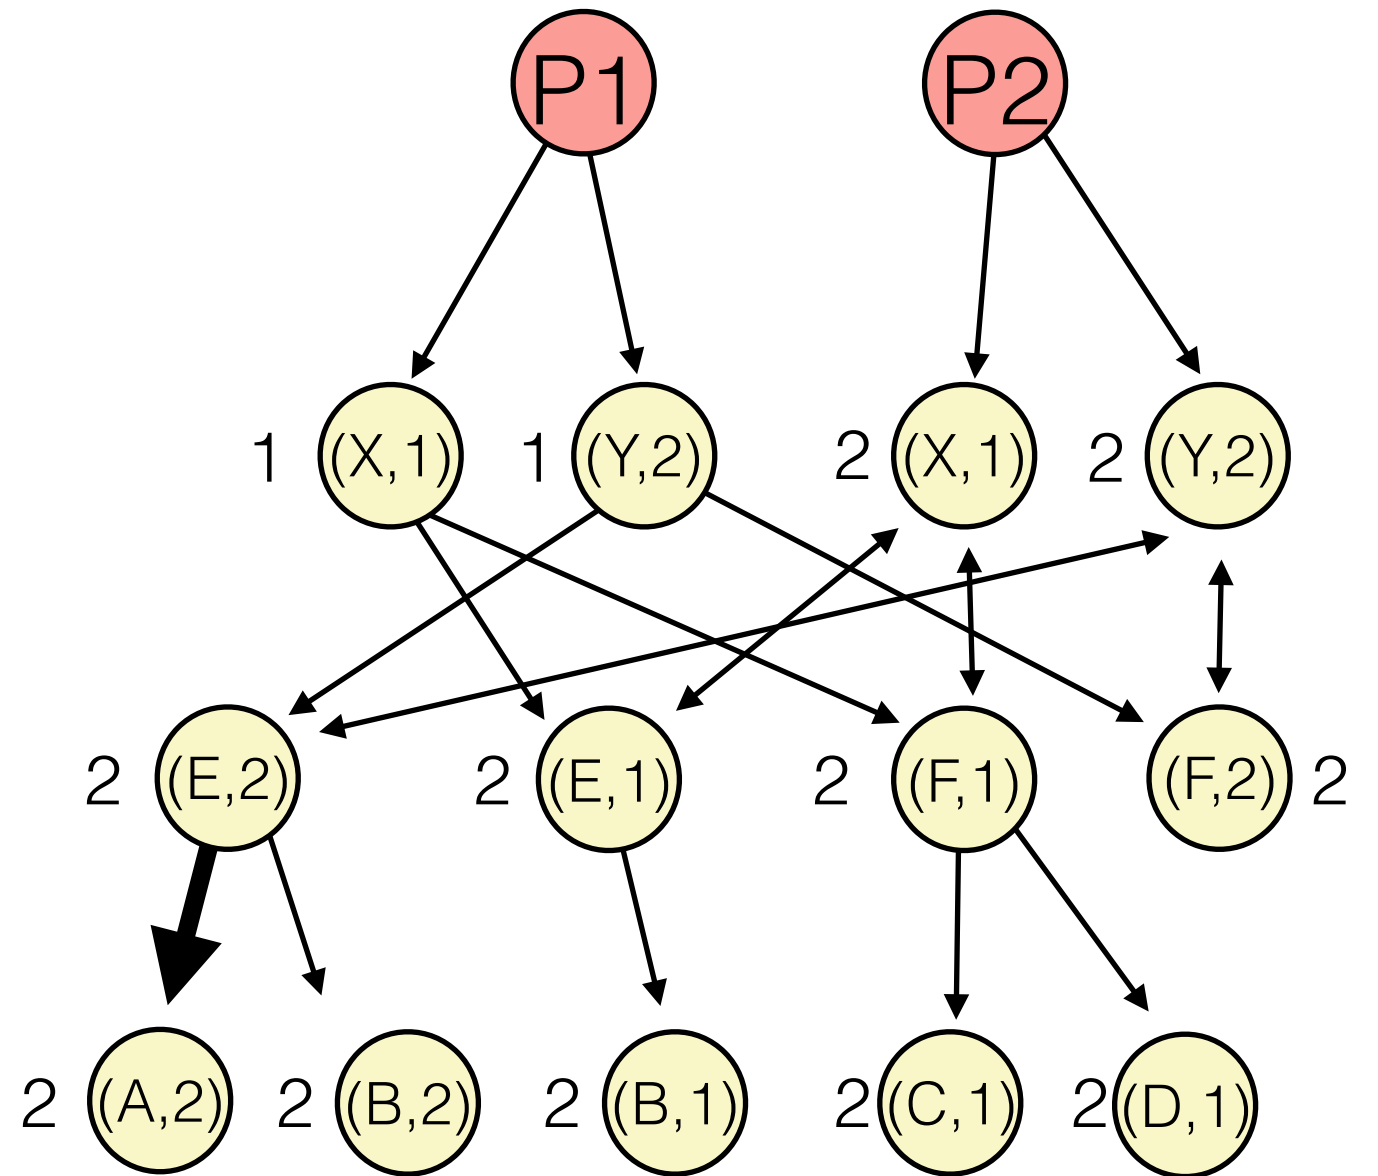

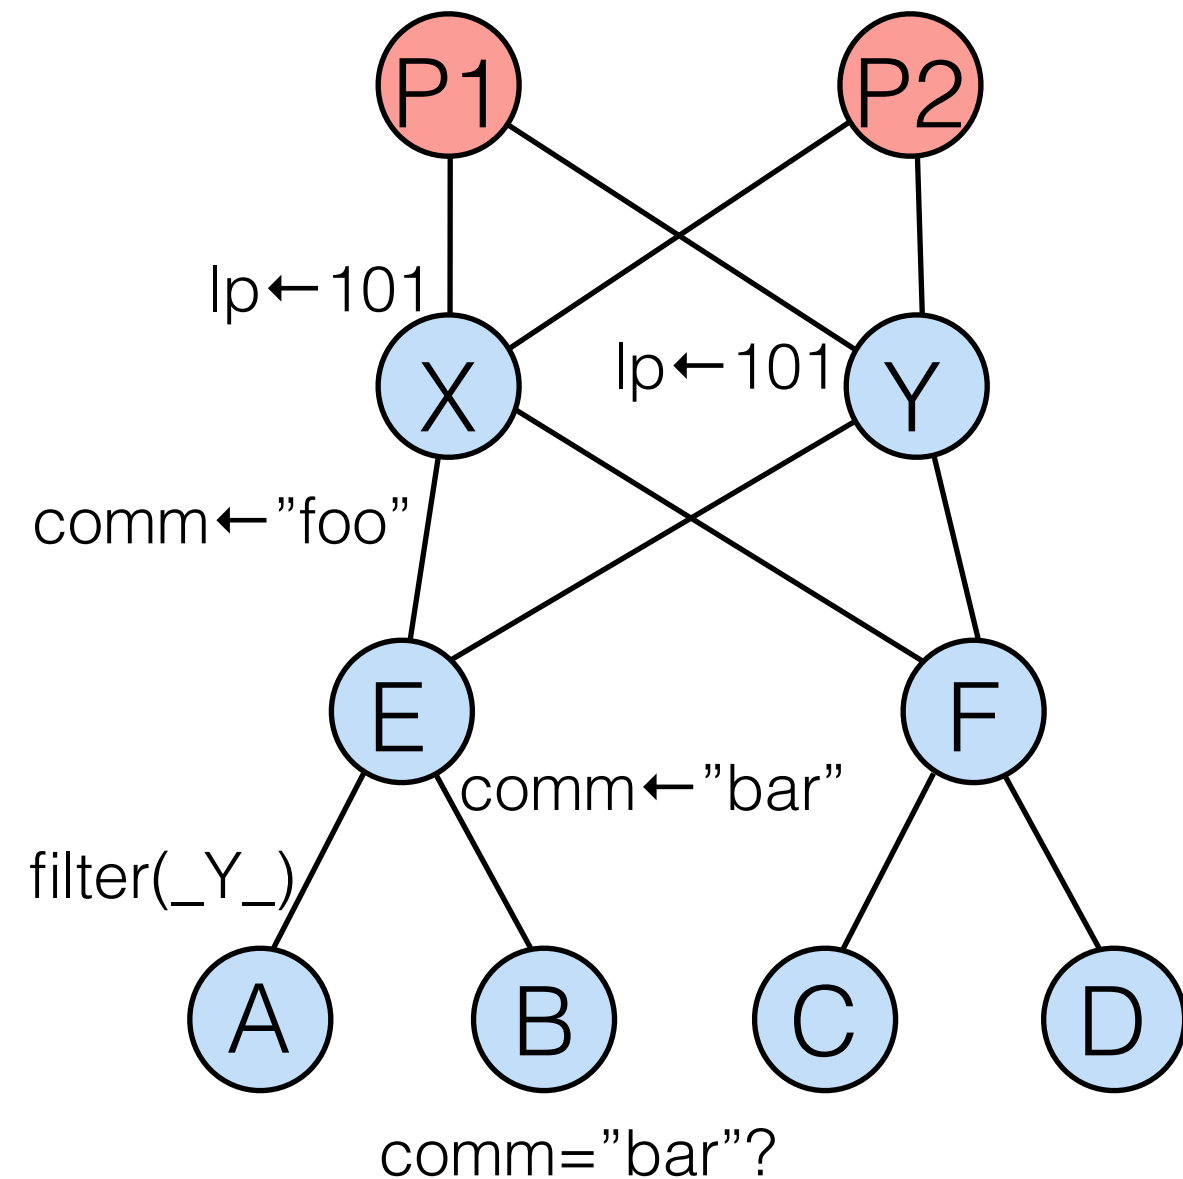
Regex Filters



Track the truth of each
regex filter while building PG!

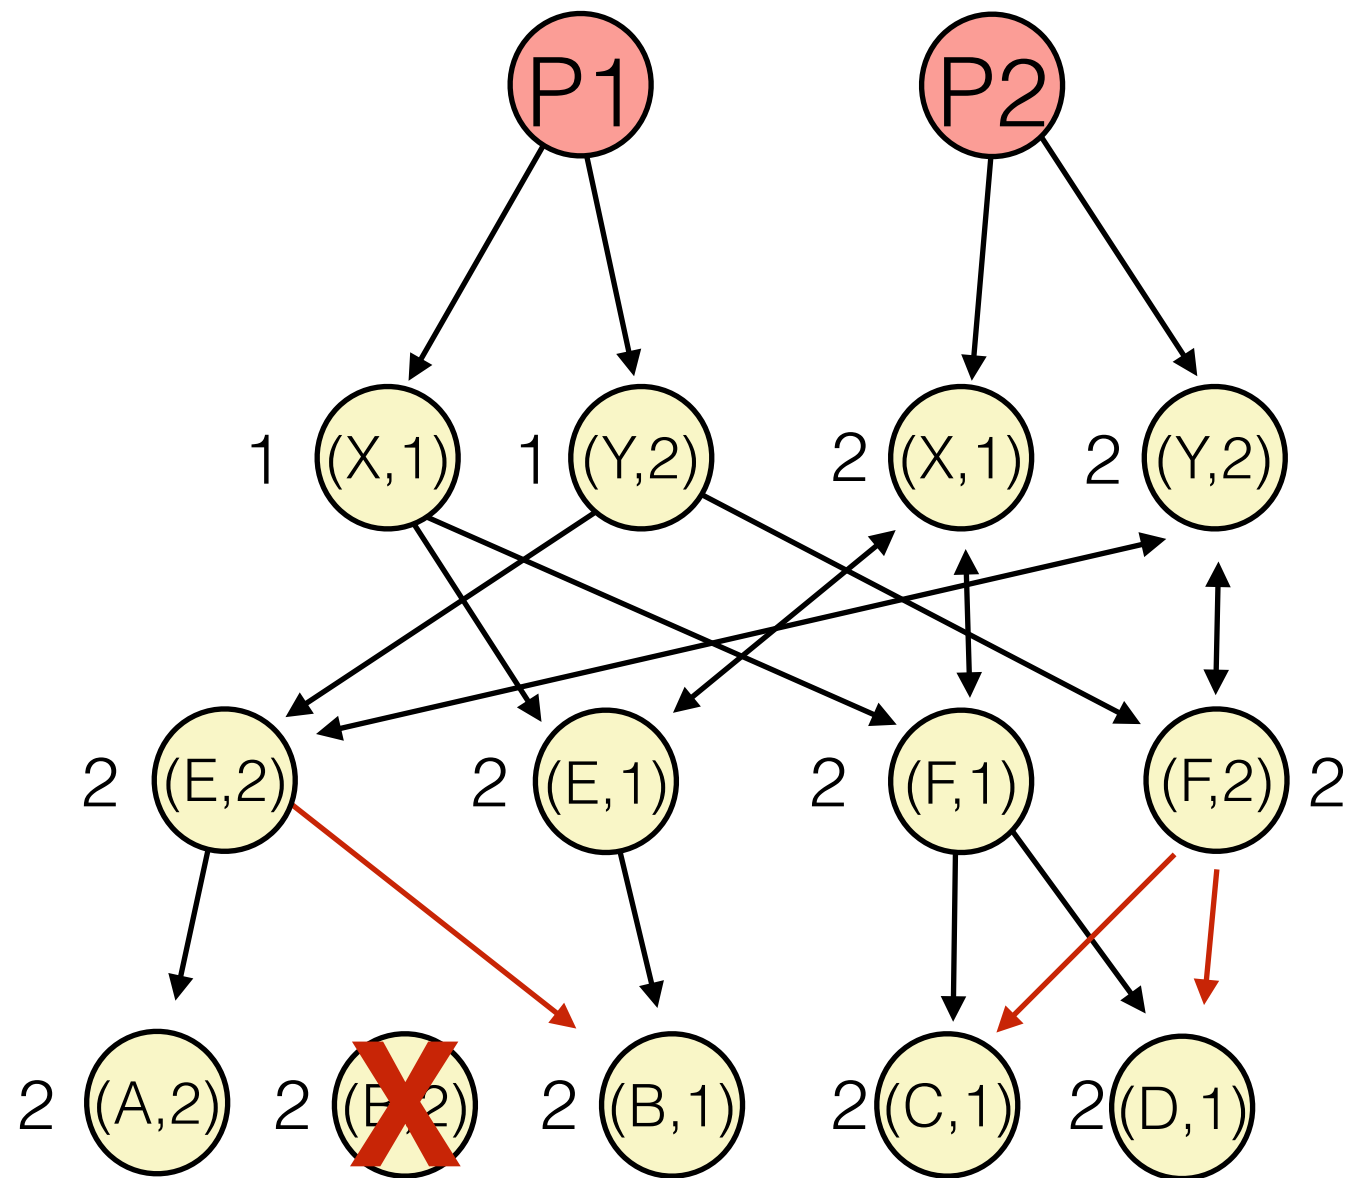


Same as when we usually
build PG from automata

Regex Filters




Filter encoded by (E,2) (A,2) edge, which knows that the regex is satisfied

Regex Filters



Possible Optimization:
merge nodes when tracking
information becomes irrelevant?



Community Tags




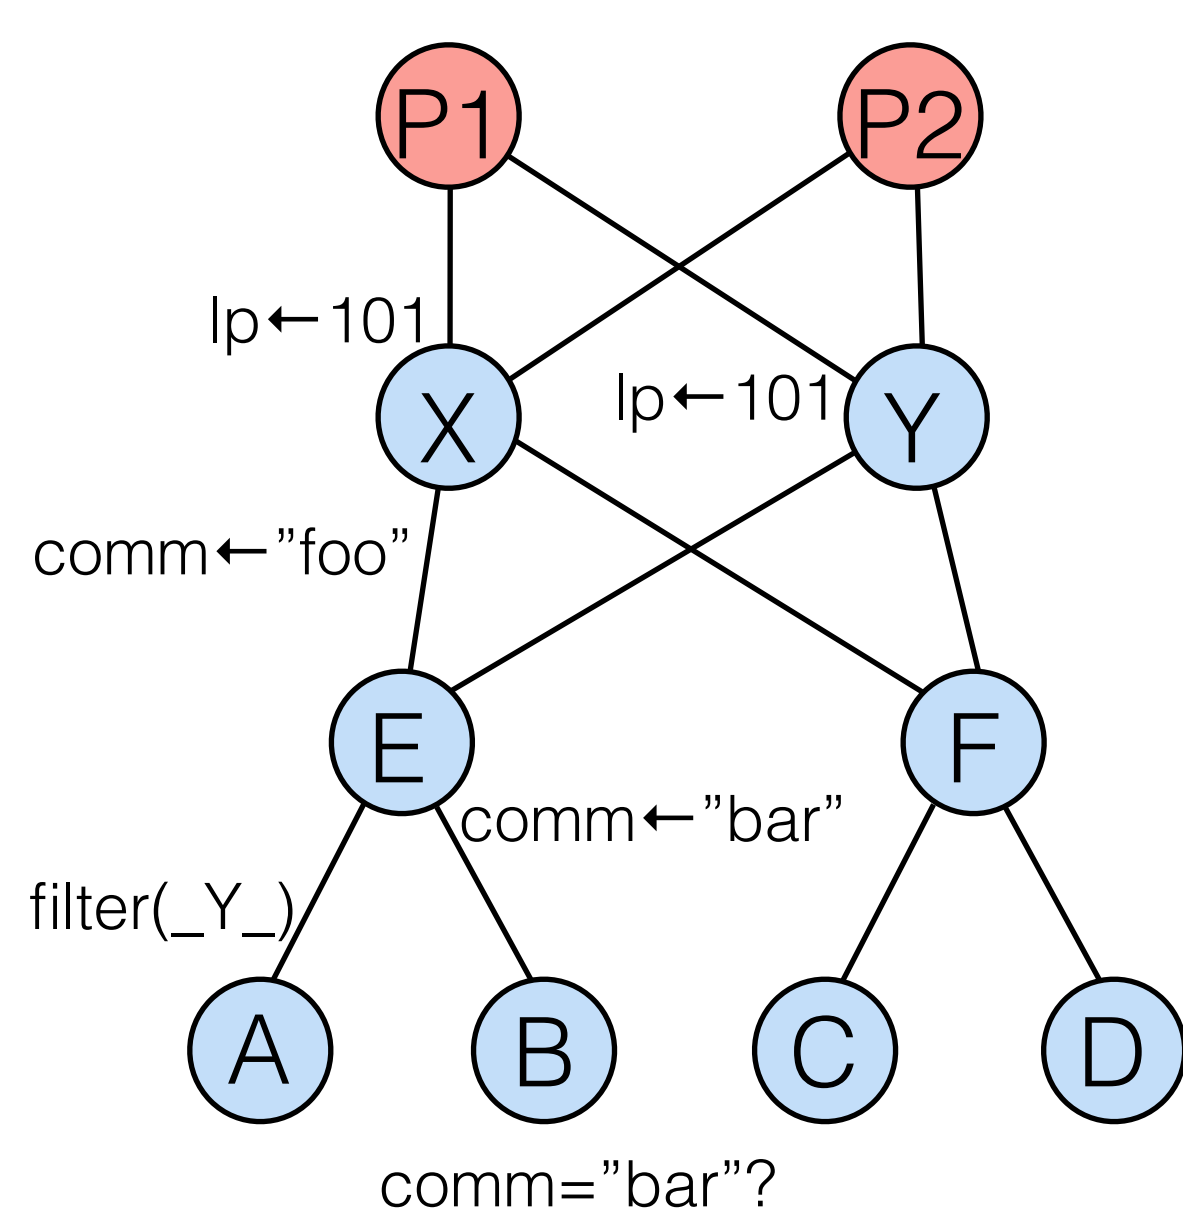
Track which community tags
are attached in PG

Same idea as before!

Community Tags



Community Tags




Tag "foo" $X \longrightarrow E$

Tag "bar" $E \longrightarrow B$ when "foo" attached

Filter for "bar" from $B \longleftarrow E$

Community Tags

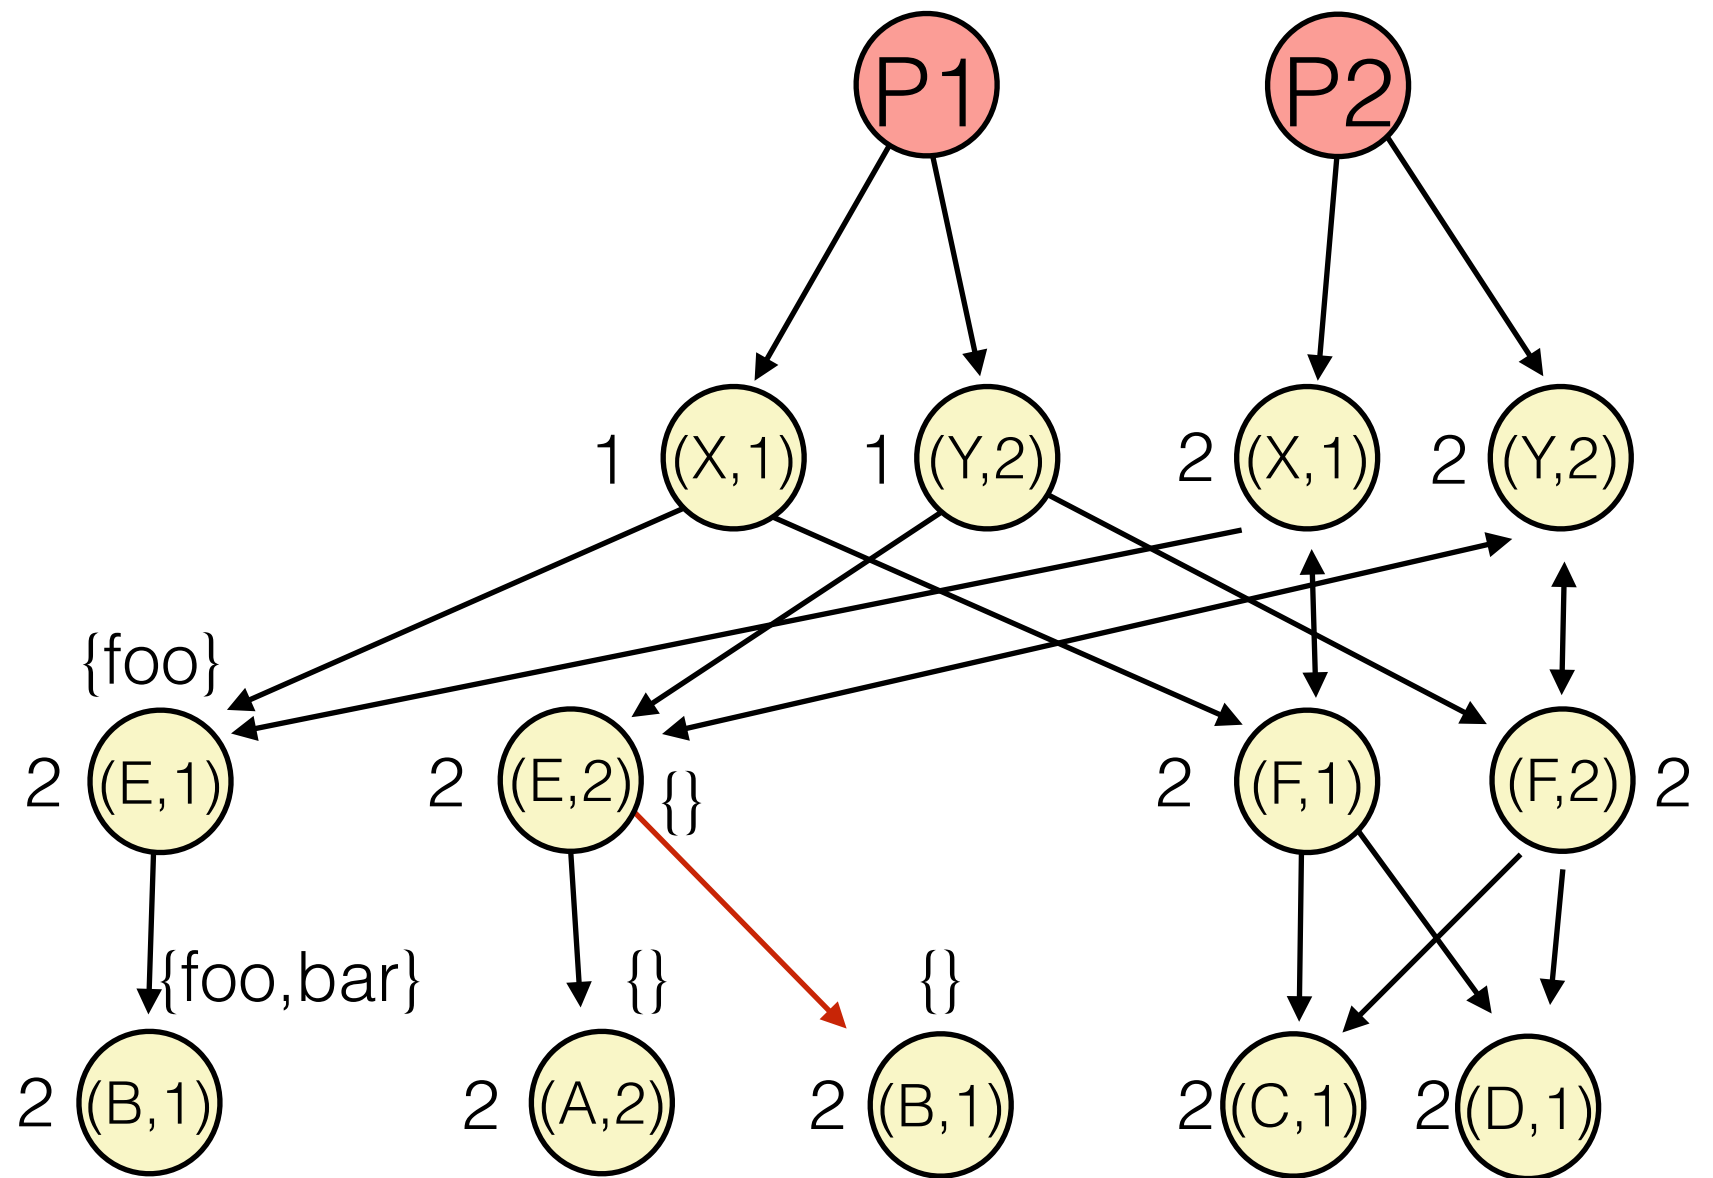


Tag “foo” $X \longrightarrow E$

Tag “bar” $E \longrightarrow B$ when “foo” attached

Filter for “bar” from $B \longleftarrow E$

Community Tags



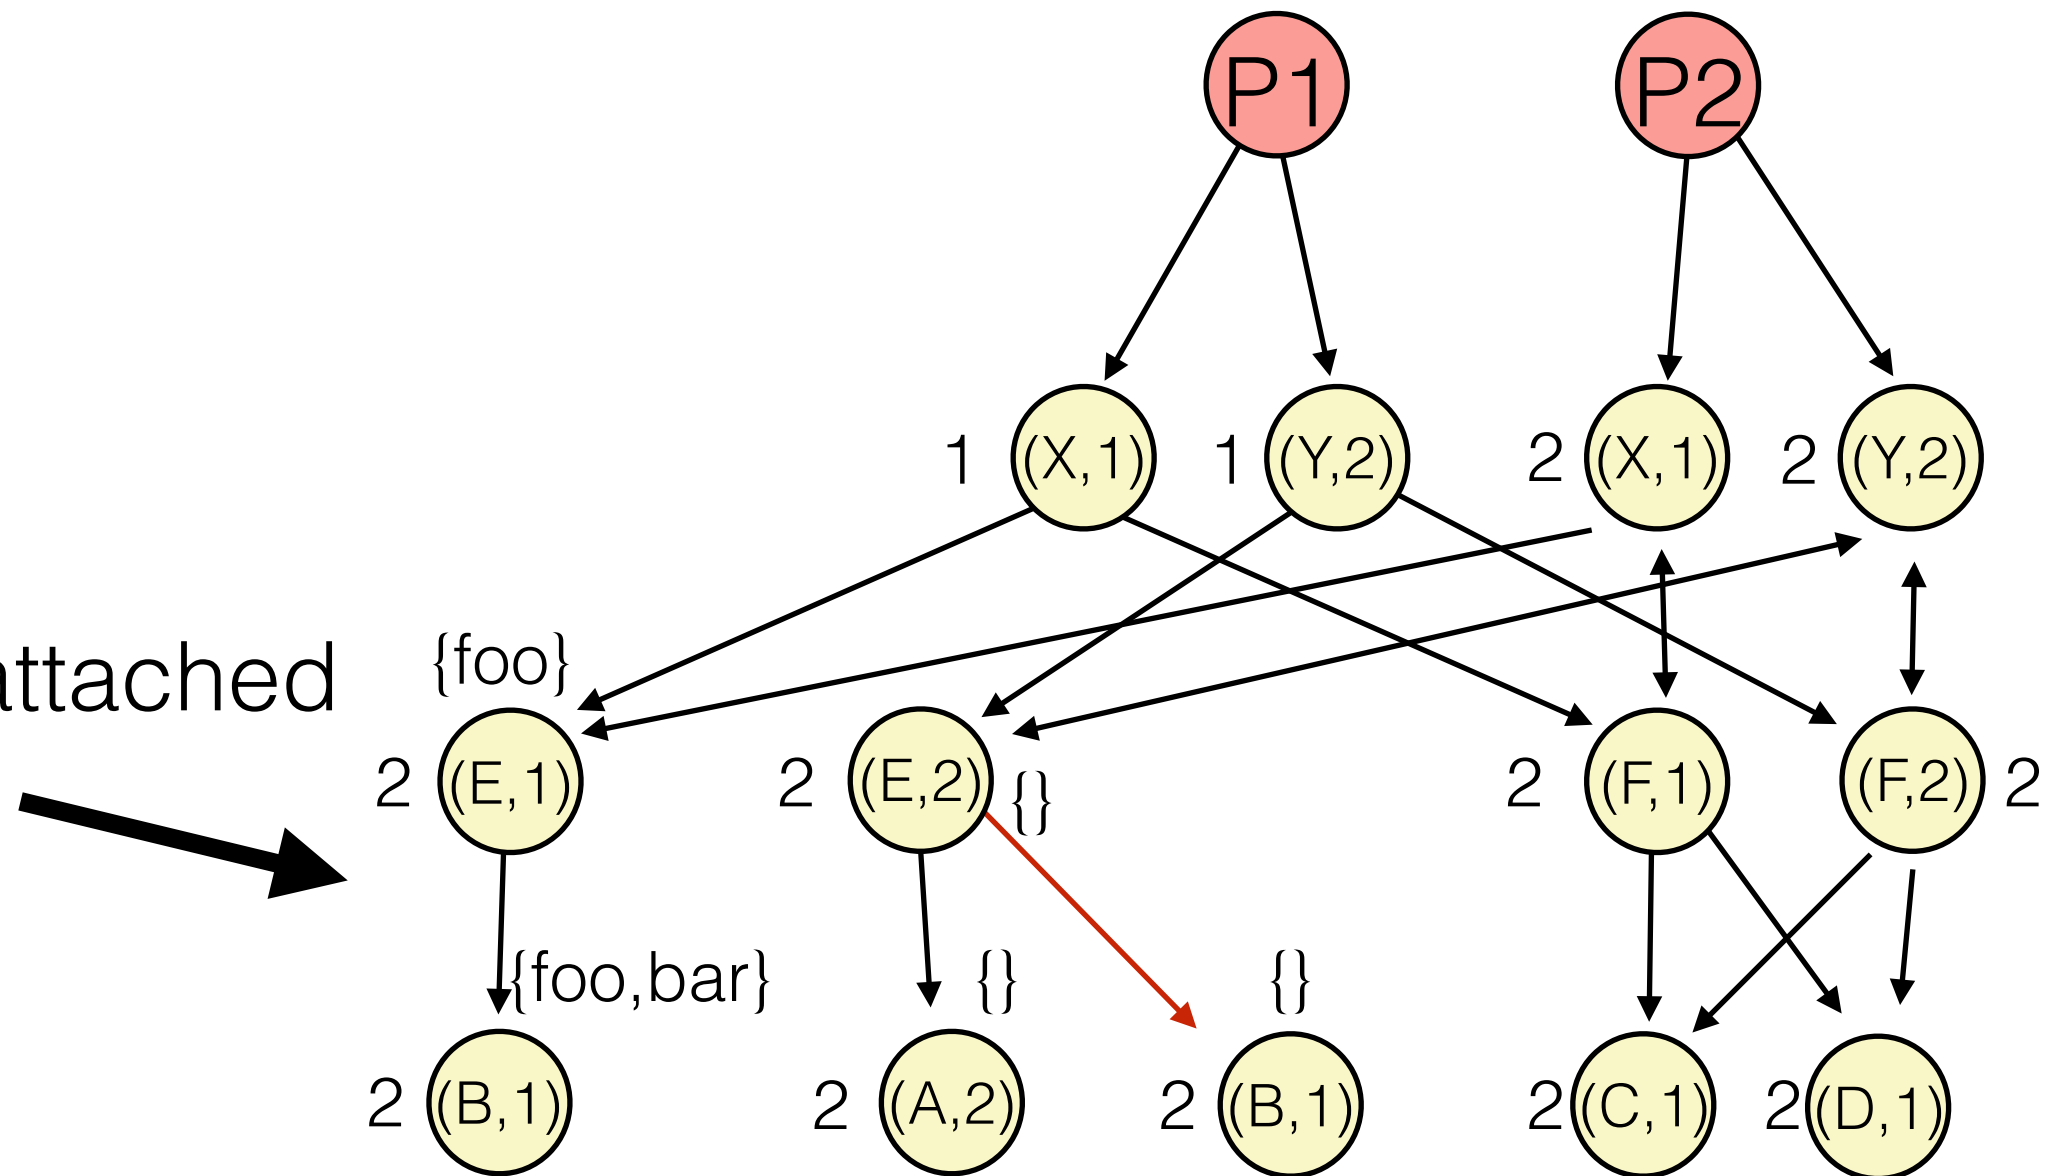
Tag “foo” $X \longrightarrow E$

Tag “bar” $E \longrightarrow B$ when “foo” attached

Filter for “bar” from $B \longleftarrow E$

Community Tags

Edge allowed
since “bar”
known to be attached

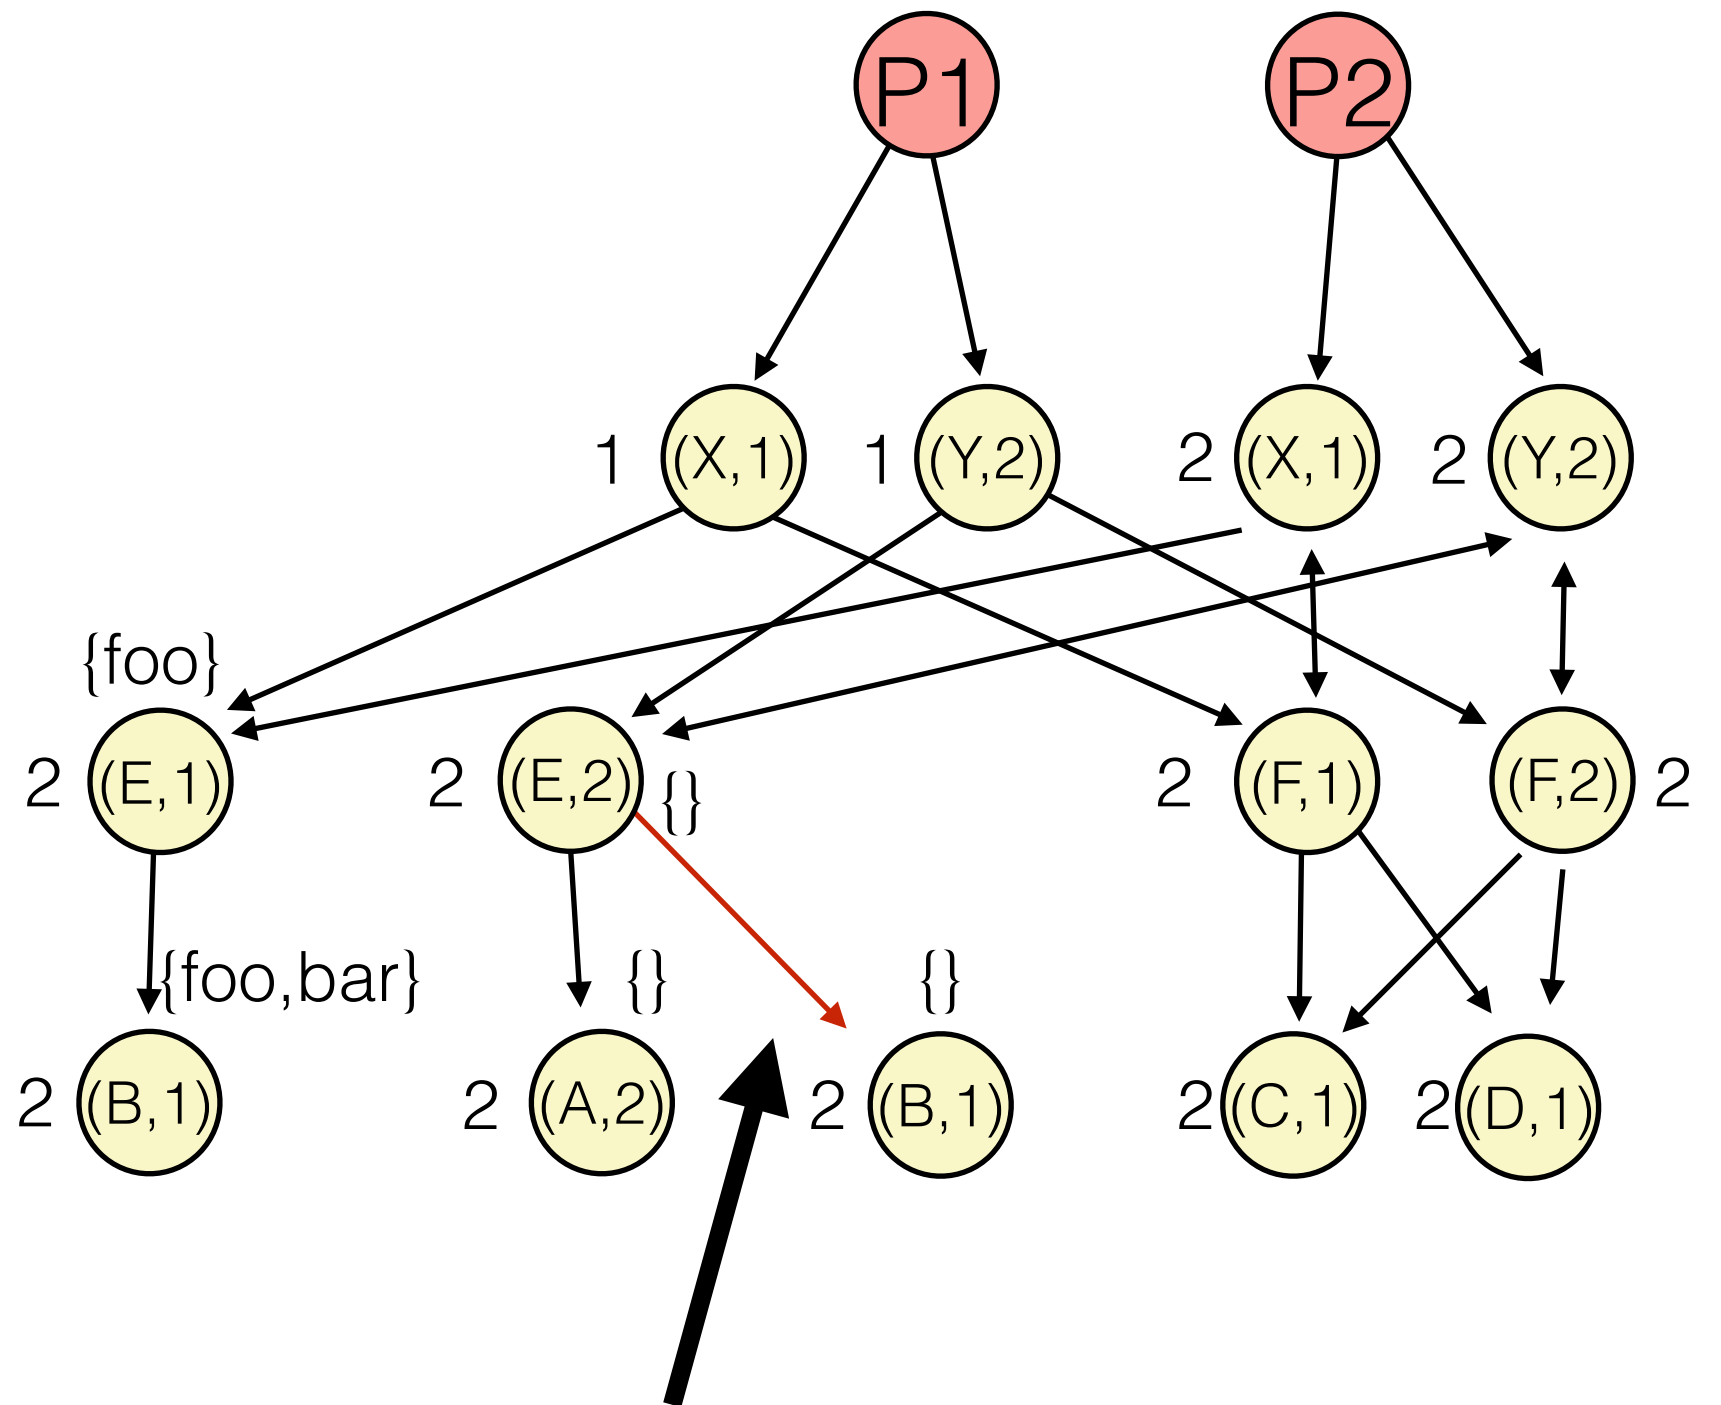


Tag “foo” $X \rightarrow E$

Tag “bar” $E \rightarrow B$ when “foo” attached

Filter for “bar” from $B \leftarrow E$

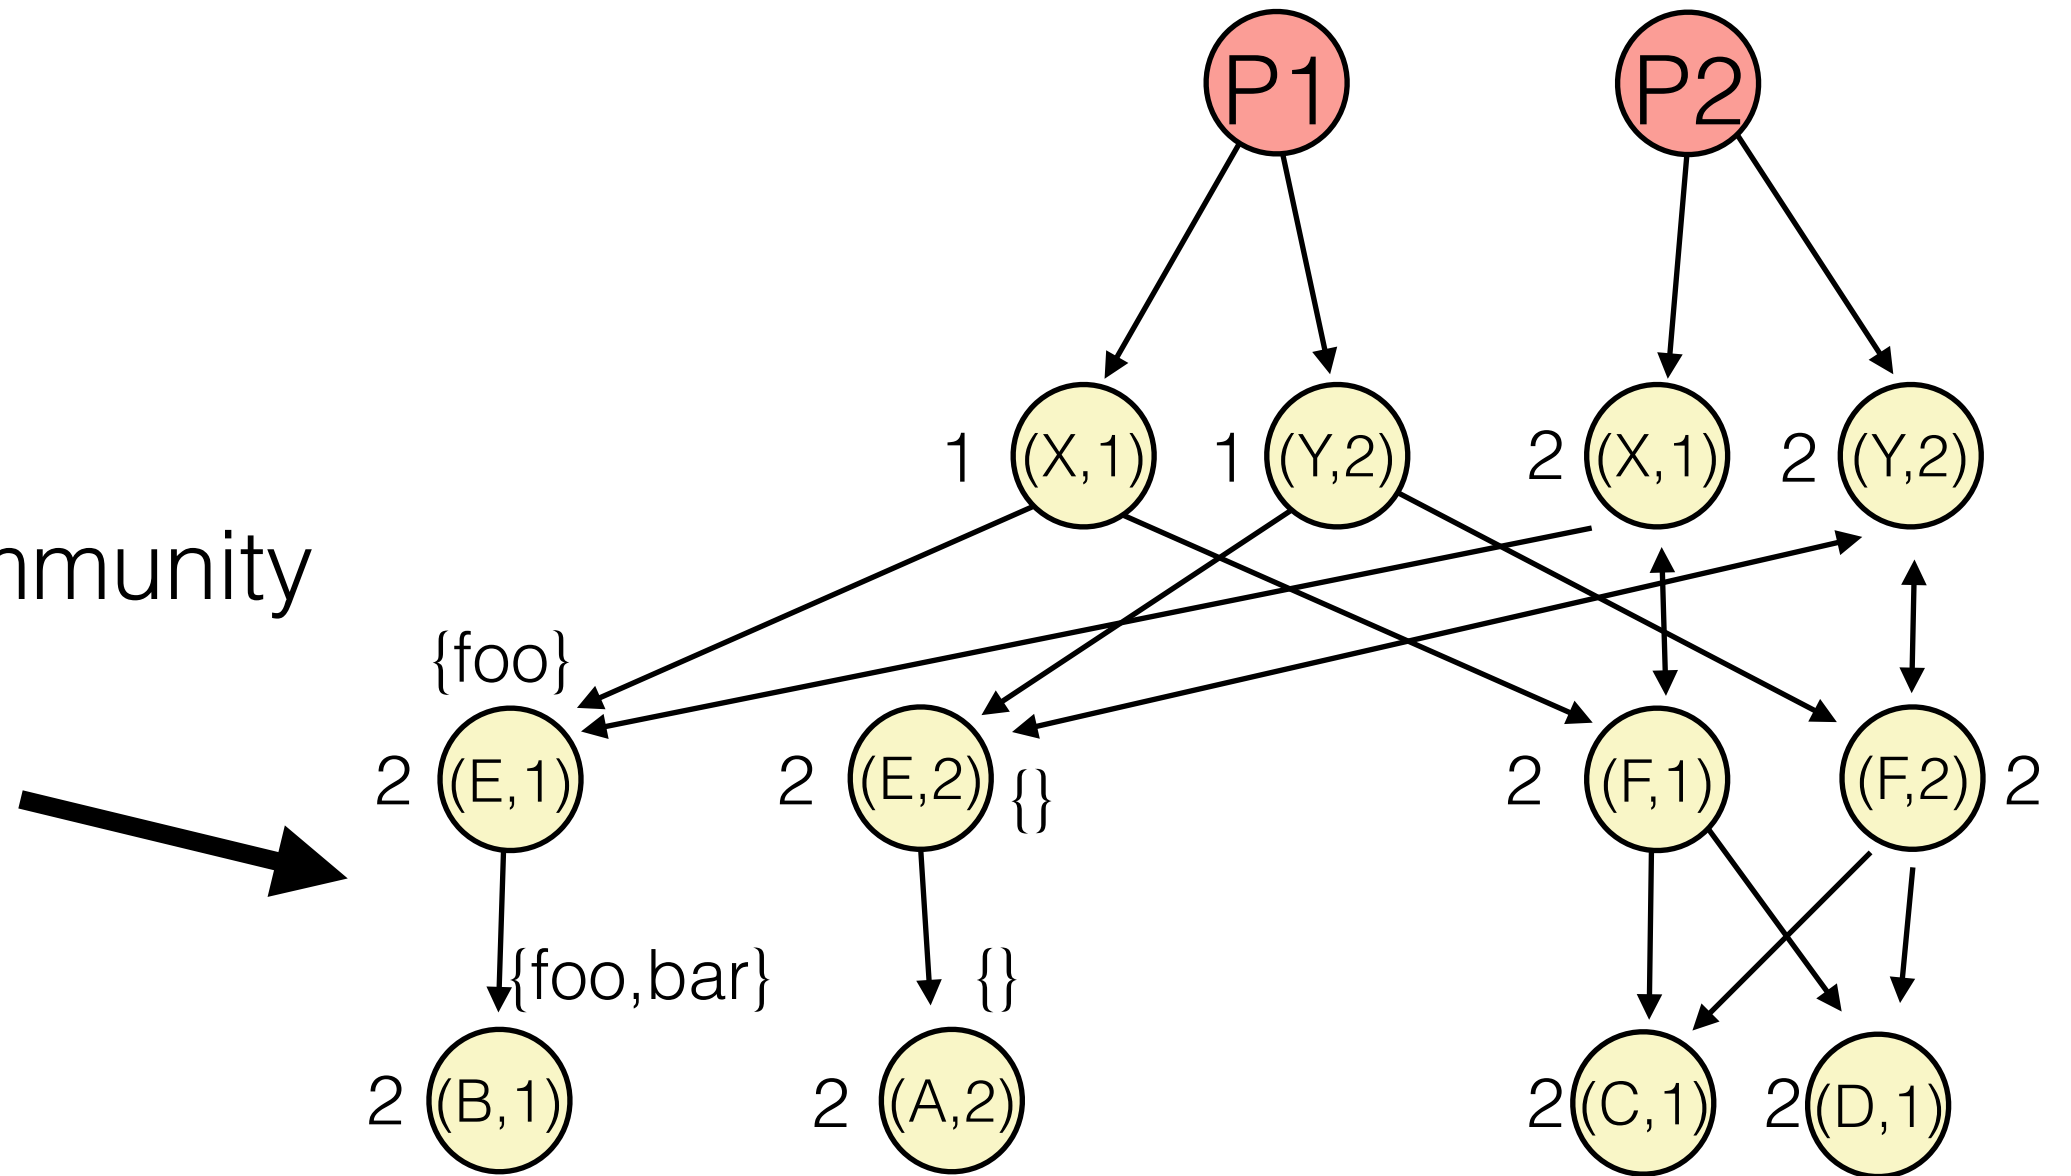
Community Tags



No edge allowed here


Community Tags

Can support removing community tags as well



Community Tags


Does traffic sent from B always go through X?



Community Tags

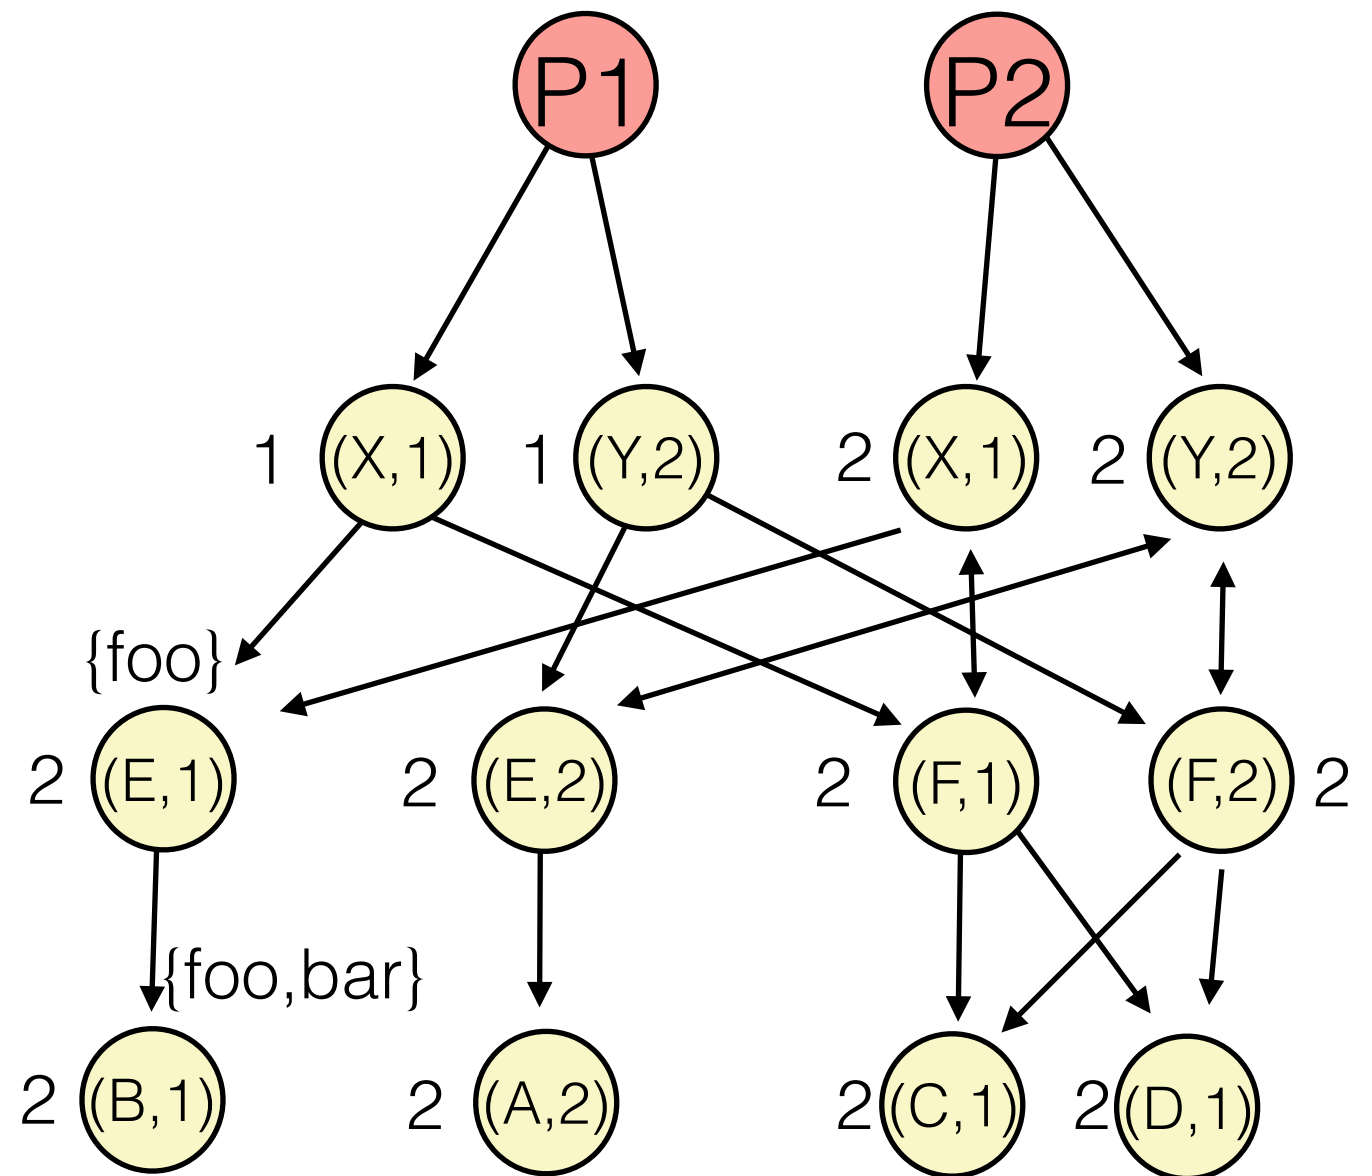
Does traffic sent from B always go through X?

Yes!



Community Tags

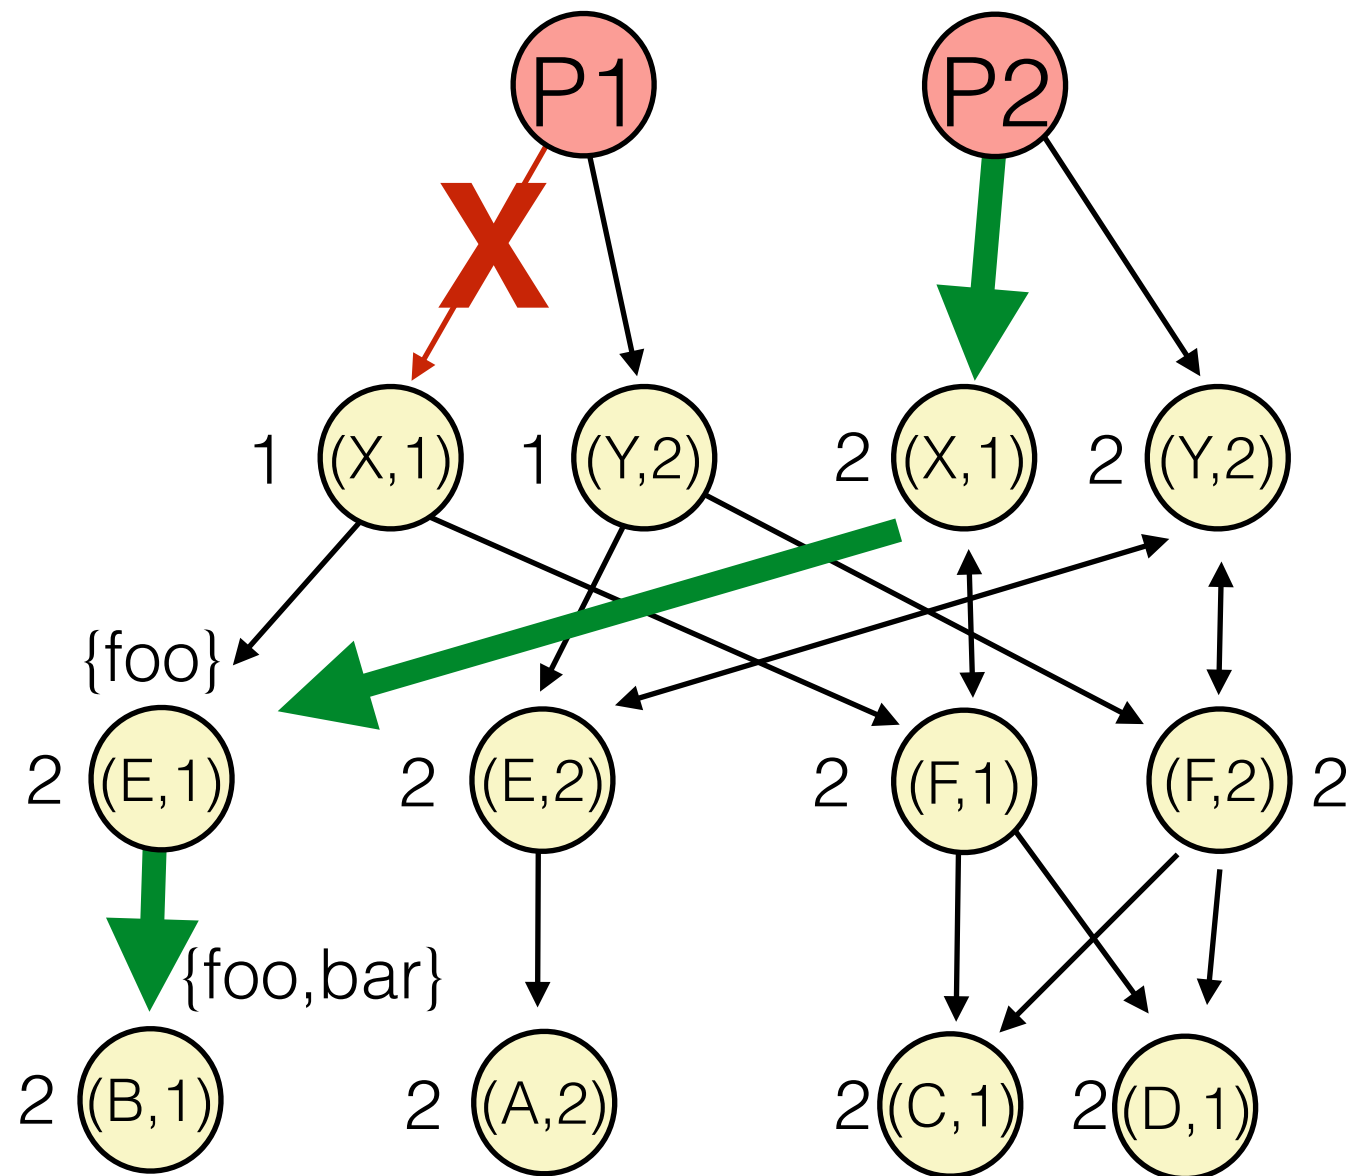
Does traffic always leave the DC through P1 when this is possible?



Community Tags

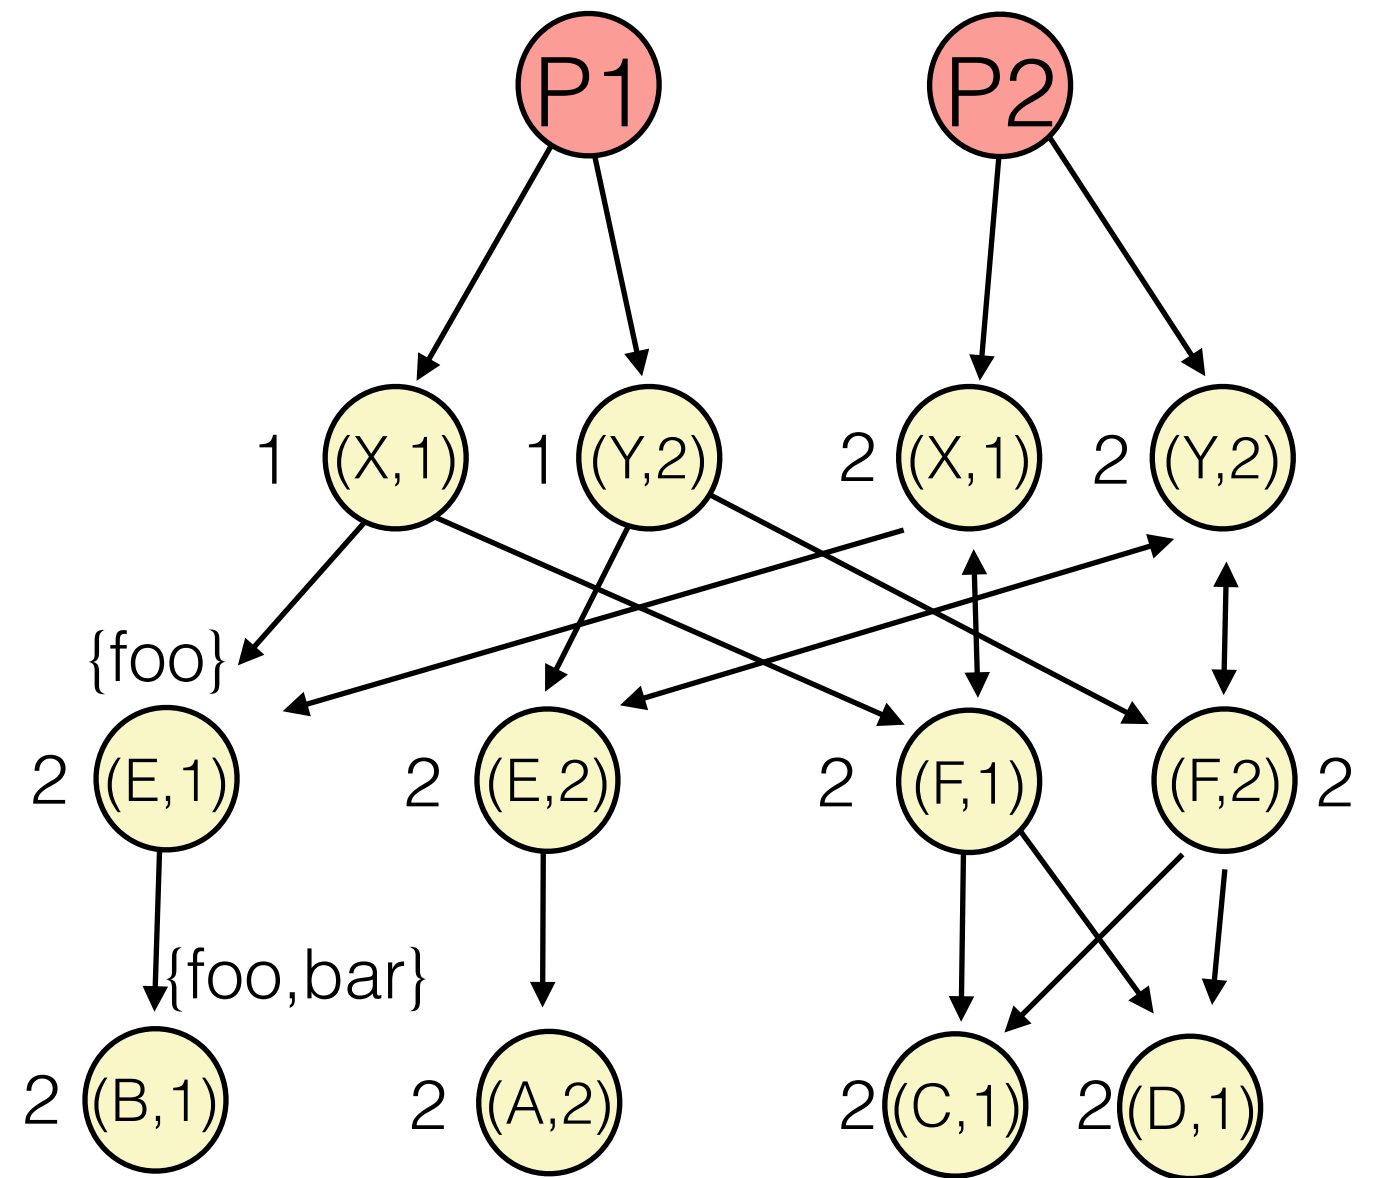
Does traffic always
leave the DC through P1
when this is possible?

Nope!



Community Tags


How many failures
to disconnect B?



Community Tags

How many failures
to disconnect B?

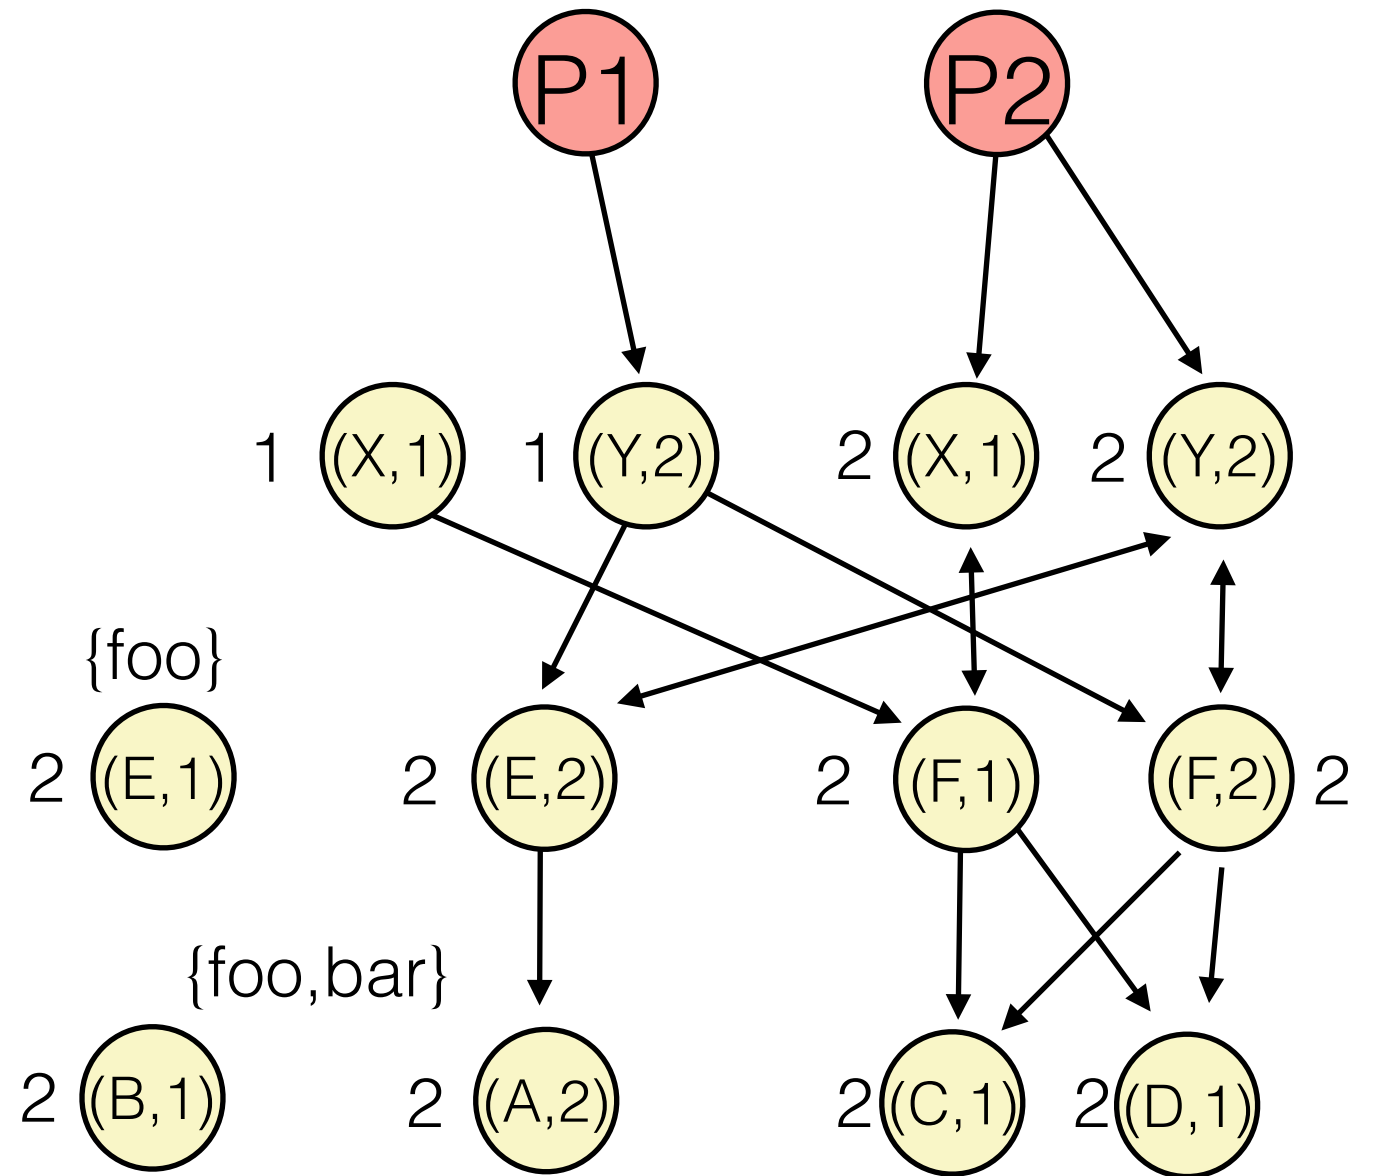
1!



Community Tags


How many failures
to disconnect B?

1!




Abstract Safety Analysis

Reachability under k-failures




Reachability under k-failures



A


Abstract reachability tells
us that all nodes are reachable

Reachability under k-failures




How many failures required to turn this **A** into an **N**

Inference Rules:




$$L \in \{A, S\}$$

Reachability under k-failures



How many failures required to turn this **A** into an **N**


Reachability under k-failures



How many failures required to turn this **A** into an **N**

In order to infer **N**, a single edge must result in **N**

Reachability under k-failures




How many failures required to turn this **A** into an **N**

In order to infer **N**, a single edge must result in **N**

Edge-by-edge, how many failures change the inference

Reachability under k-failures




Conservatively assume only **S** is reachable for each node

How many failures required to turn this **A** into an **N**


In order to infer **N**, a single edge must result in **N**

Edge-by-edge, how many failures change the inference

Reachability under k-failures




Reachability under k-failures



How many failures required to turn this **A** into an **S**


Reachability under k-failures



How many failures required to turn this **A** into an **S**

More challenging problem

Reachability under k-failures




A

How many failures required to turn this **A** into an **S**

More challenging problem


Reachability under k-failures



How many failures required to turn this **A** into an **S**

More challenging problem


Reachability under k-failures



How many failures required to turn this **A** into an **S**

More challenging problem


Reachability under k-failures




How many failures required to turn this **A** into an **S**

More challenging problem


Reachability under k-failures




Reachability under k-failures



Reachability under k-failures




Reachability under k-failures



Idea: generate a “worst” case concrete topology, and find a lower bound on the min-cut of this topology


Reachability under k-failures



Idea: generate a “worst” case concrete topology, and find a lower bound on the min-cut of this topology


Keep the topology abstract (small)

Reachability under k-failures




Claim: Take the min of every edge/node, and this will be the worst topology

Reachability under k-failures



Claim: Take the min of every edge/node, and this will be the worst topology


Reachability under k-failures



Claim: Take the min of every edge/node, and this will be the worst topology

Why: Can always fail more nodes/edges to get this topology


Reachability under k-failures



Claim: Take the min of every edge/node, and this will be the worst topology

Idea: Find a lower bound on the number of disjoint paths in the concrete topology.



Reachability under k-failures





Claim: Take the min of every edge/node, and this will be the worst topology

Idea: Find a lower bound on the number of disjoint paths in the concrete topology.

Reachability under k -failures




Reachability under k-failures



Reachability under k-failures


Destination



Source


We care about disconnecting some arbitrary node at the top from some arbitrary node at the bottom

Reachability under k-failures



We care about disconnecting some arbitrary node at the top from some arbitrary node at the bottom


Reachability under k-failures



4 disjoint edges to 4 disjoint nodes

Reachability under k-failures

Destination




Source

4 disjoint edges to 3 disjoint nodes

4 disjoint edges to 4 disjoint nodes

Reachability under k-failures




3 disjoint edges to 1 node

4 disjoint edges to 3 disjoint nodes

4 disjoint edges to 4 disjoint nodes

Reachability under k-failures



3 disjoint edges to 1 node



4 disjoint edges to 3 disjoint nodes

4 disjoint edges to 4 disjoint nodes



$$\min\{3, 4, 4\} = \mathbf{3}$$

Compilation Correctness



Failure Safety (Recap)





Failure Safety (Recap)





Failure Safety (Recap)




Failure Safety (Recap)



Failure Safety (Recap)




Failure Safety (Recap)



Proof of Correctness (High level)


Statement:

Traffic always flows along ***some best simple*** path to source s when such a path exists in the network (given failures)




Assume path P is one of the highest rank simple paths in the network given the failures

Then path P exists in the PG with (best) rank of r for source s




Assume path P is one of the highest rank simple paths in the network given the failures

Then path P exists in the PG with (best) rank of r for source s




The only way traffic does not flow along path P is if some node on P prefers a different advertisement



The only way traffic does not flow along path P is if some node on P prefers a different advertisement


What we know:

- 1) advertisement reaches N_2
- 2) Failure analysis prefers N_2
- 3) N_2 has the same path b to $r' \leq r$




What we know:

- 1) advertisement reaches N_2
- 2) Failure analysis prefers $N_2 \geq N_1$
- 3) N_2 has the same path b to $r' \leq r$




What we know:

- 1) advertisement reaches N_2
- 2) Failure analysis prefers $N_2 \geq N_1$
- 3) N_2 has the same path b to $r' \leq r$



Case 1 ($a'.b$ is a simple path)


Proceed by induction on
path length of b



Case 2 ($a'.b$ is **not** simple)


a' must be simple
(advertisement)

b must be simple
(since P is simple)



Case 2 ($a'.b$ is **not** simple)

From failure analysis, $X_1 \geq X_2$



Case 2 (a'.b is **not** simple)

From failure analysis, $X_1 \geq X_2$
Induction on smaller paths c,f

