

# Propane: Programming Distributed Control Planes

Ryan Beckett<sup>1</sup>, Todd Millstein<sup>2</sup>, Jitendra Padhye<sup>3</sup>, Ratul Mahajan<sup>3</sup>, David Walker<sup>1</sup>  
<sup>1</sup>Princeton University, <sup>2</sup>University of California Los Angeles, <sup>3</sup>Microsoft Research

## Motivation

It is well known that traditional network configuration is highly error-prone [1,2,3]. Surveys of network operators often point to human error as a leading cause of misconfigurations [4,5].




Figure 1. Operator survey on network misconfiguration causes [4,5]

**High-level Idea:** Traditional networks rely on distributed mechanisms to compute forwarding paths. This means they have many nice properties (e.g., scalability, latency), but configuration is hard because it involves distributed programming (a configuration for each device). SDN simplifies configuration but this comes at the cost of many such properties [6,7].

**Distributed programming:** a config for each device  
**Centralized programming:** a config for the network

**Distributed mechanism:** devices talk to each other  
**Centralized mechanism:** a controller decides

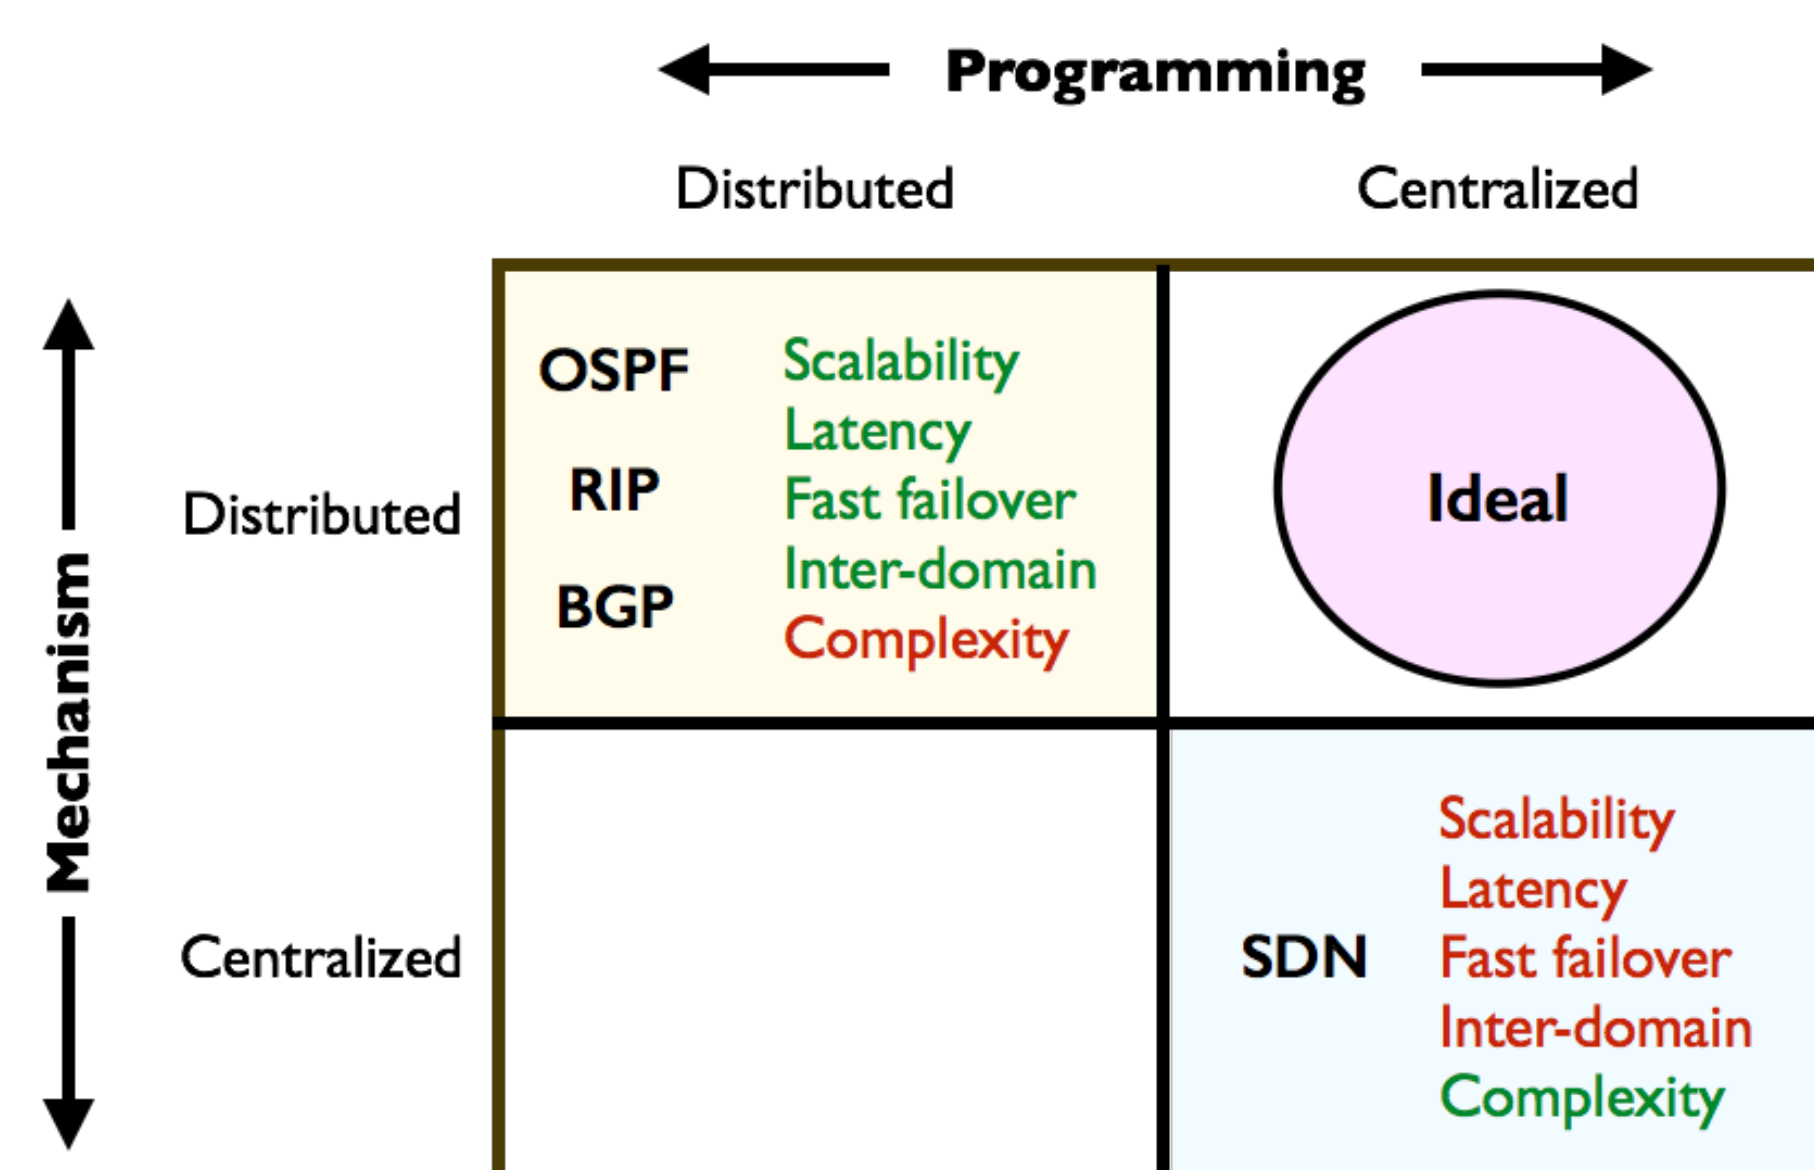



Figure 2. Tradeoffs between traditional networks and SDN

## Propane Language

### Programming Model:

- A fully centralized view of the network
- Including other autonomous system
- High-level constraints and preferences on paths
- Can describe both intra- and inter-domain routes



### Guarantees:

- Complies to pure **distributed BGP** configurations
- Policy-compliant under **all possible failures**
- That is, preserves the centralized abstraction

### Example: Configuring a datacenter with BGP

- BGP running on every router
- Keep local services internal to the datacenter
- Aggregate global prefixes externally
- Prefer leaving through Peer1 over Peer2
- Don't let the datacenter become a transit point

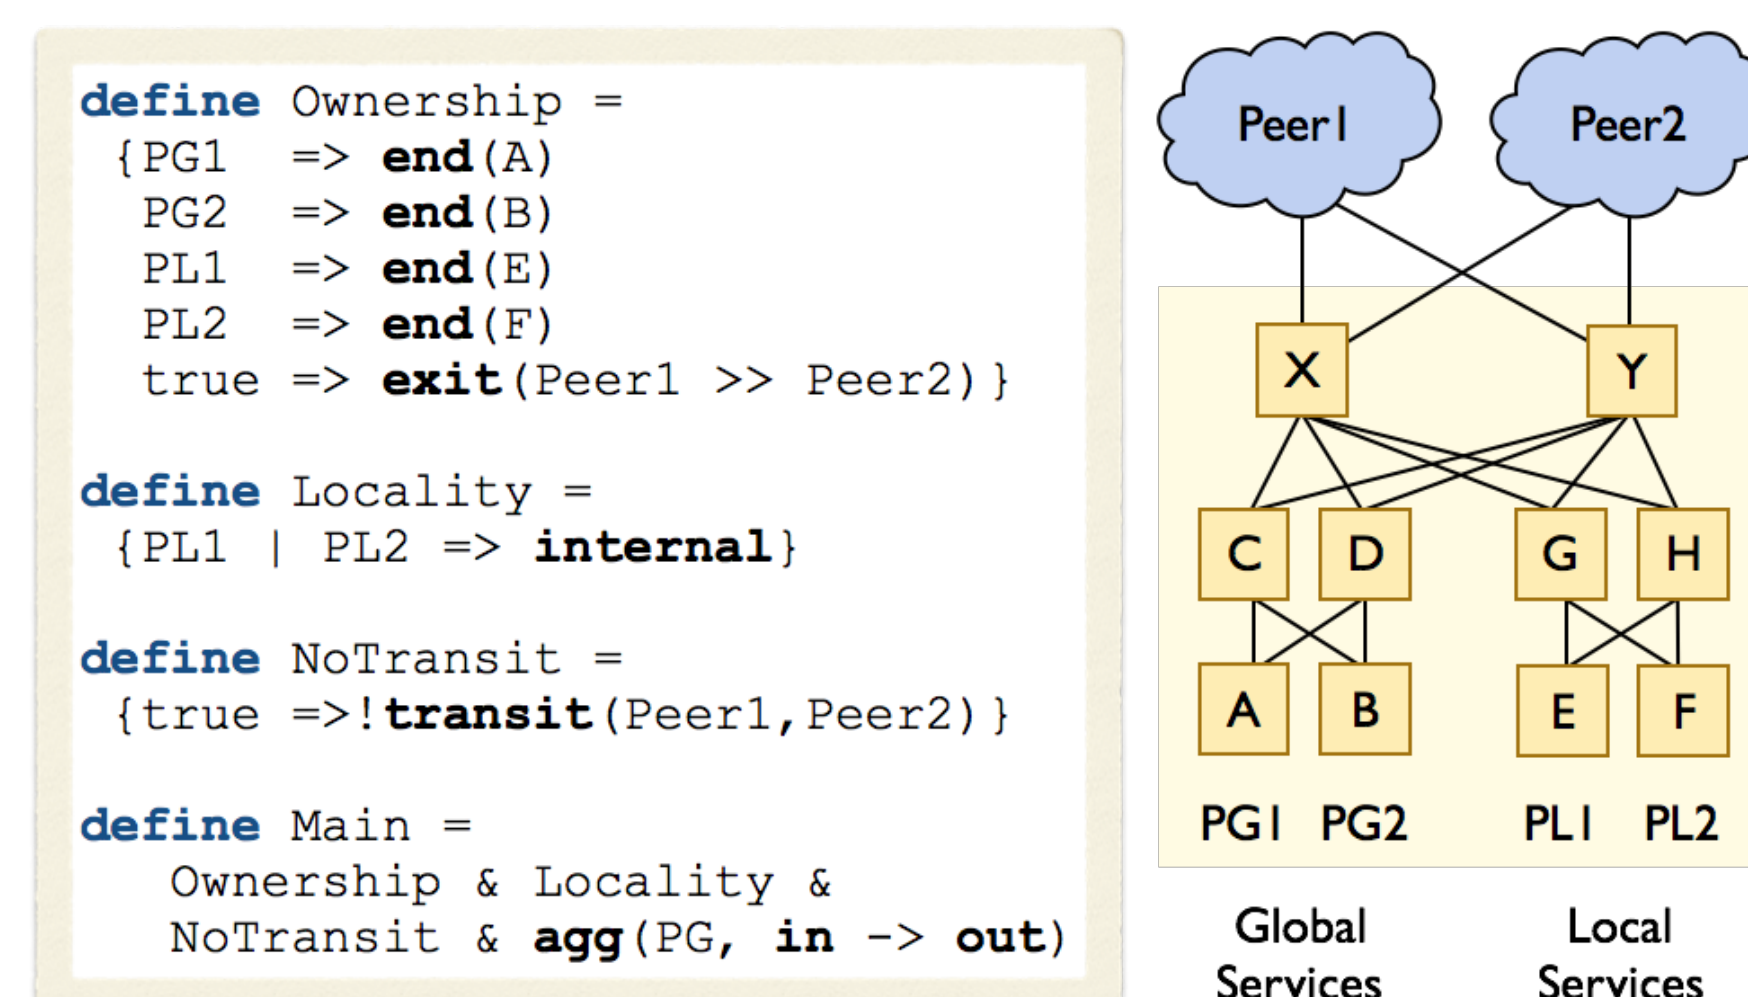


Figure 3. Configuring a datacenter with Propane

## Compilation

### Compilation Stages:

1. Compile Propane to per-destination state machines that associate paths with ranks. A lower rank means the path is preferred.
2. Combine this with the topology to get a graph representing all policy-compliant paths
3. After checking the graph for failure safety, we generate per-device BGP configs from this graph.

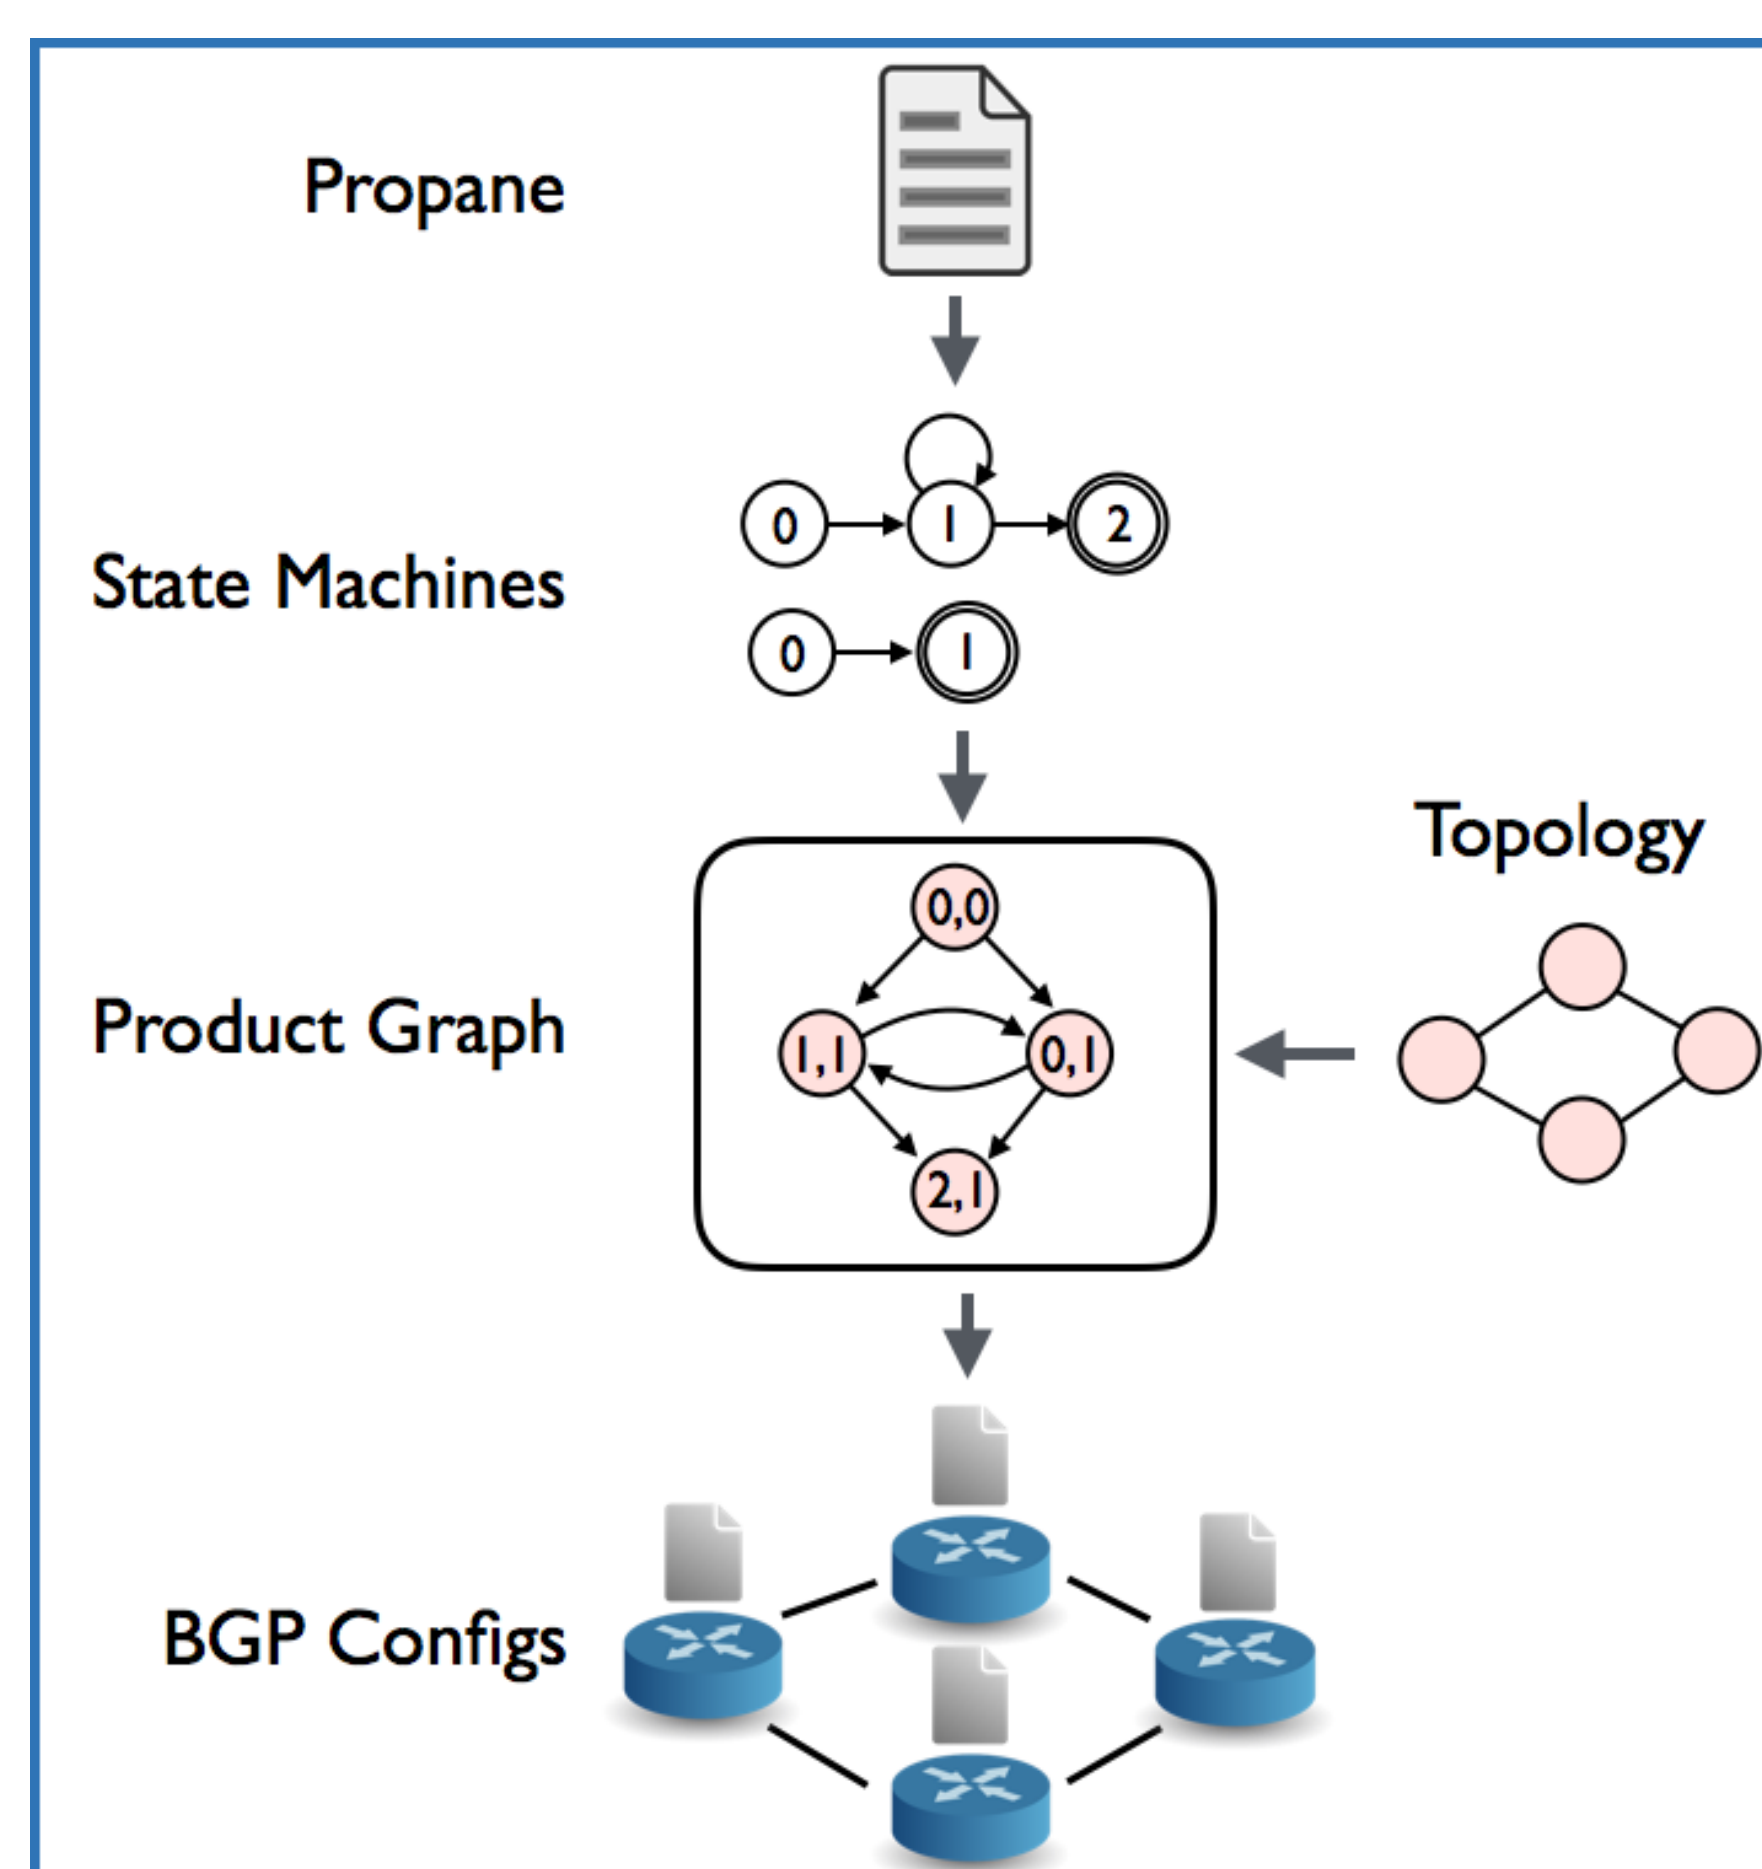


Figure 4. Compilation pipeline for the Propane language

### Product Graph:

- Intermediate representation for policies
- Combines topology and policy
- Graph that captures all policy-compliant paths
- Associates ranks with paths to a destination
- Used for safety analysis / failure analysis of

### Example: Creating a product graph

Add an edge in the Product Graph if

- The edge exists in the topology
- The state machines transition on the edge target

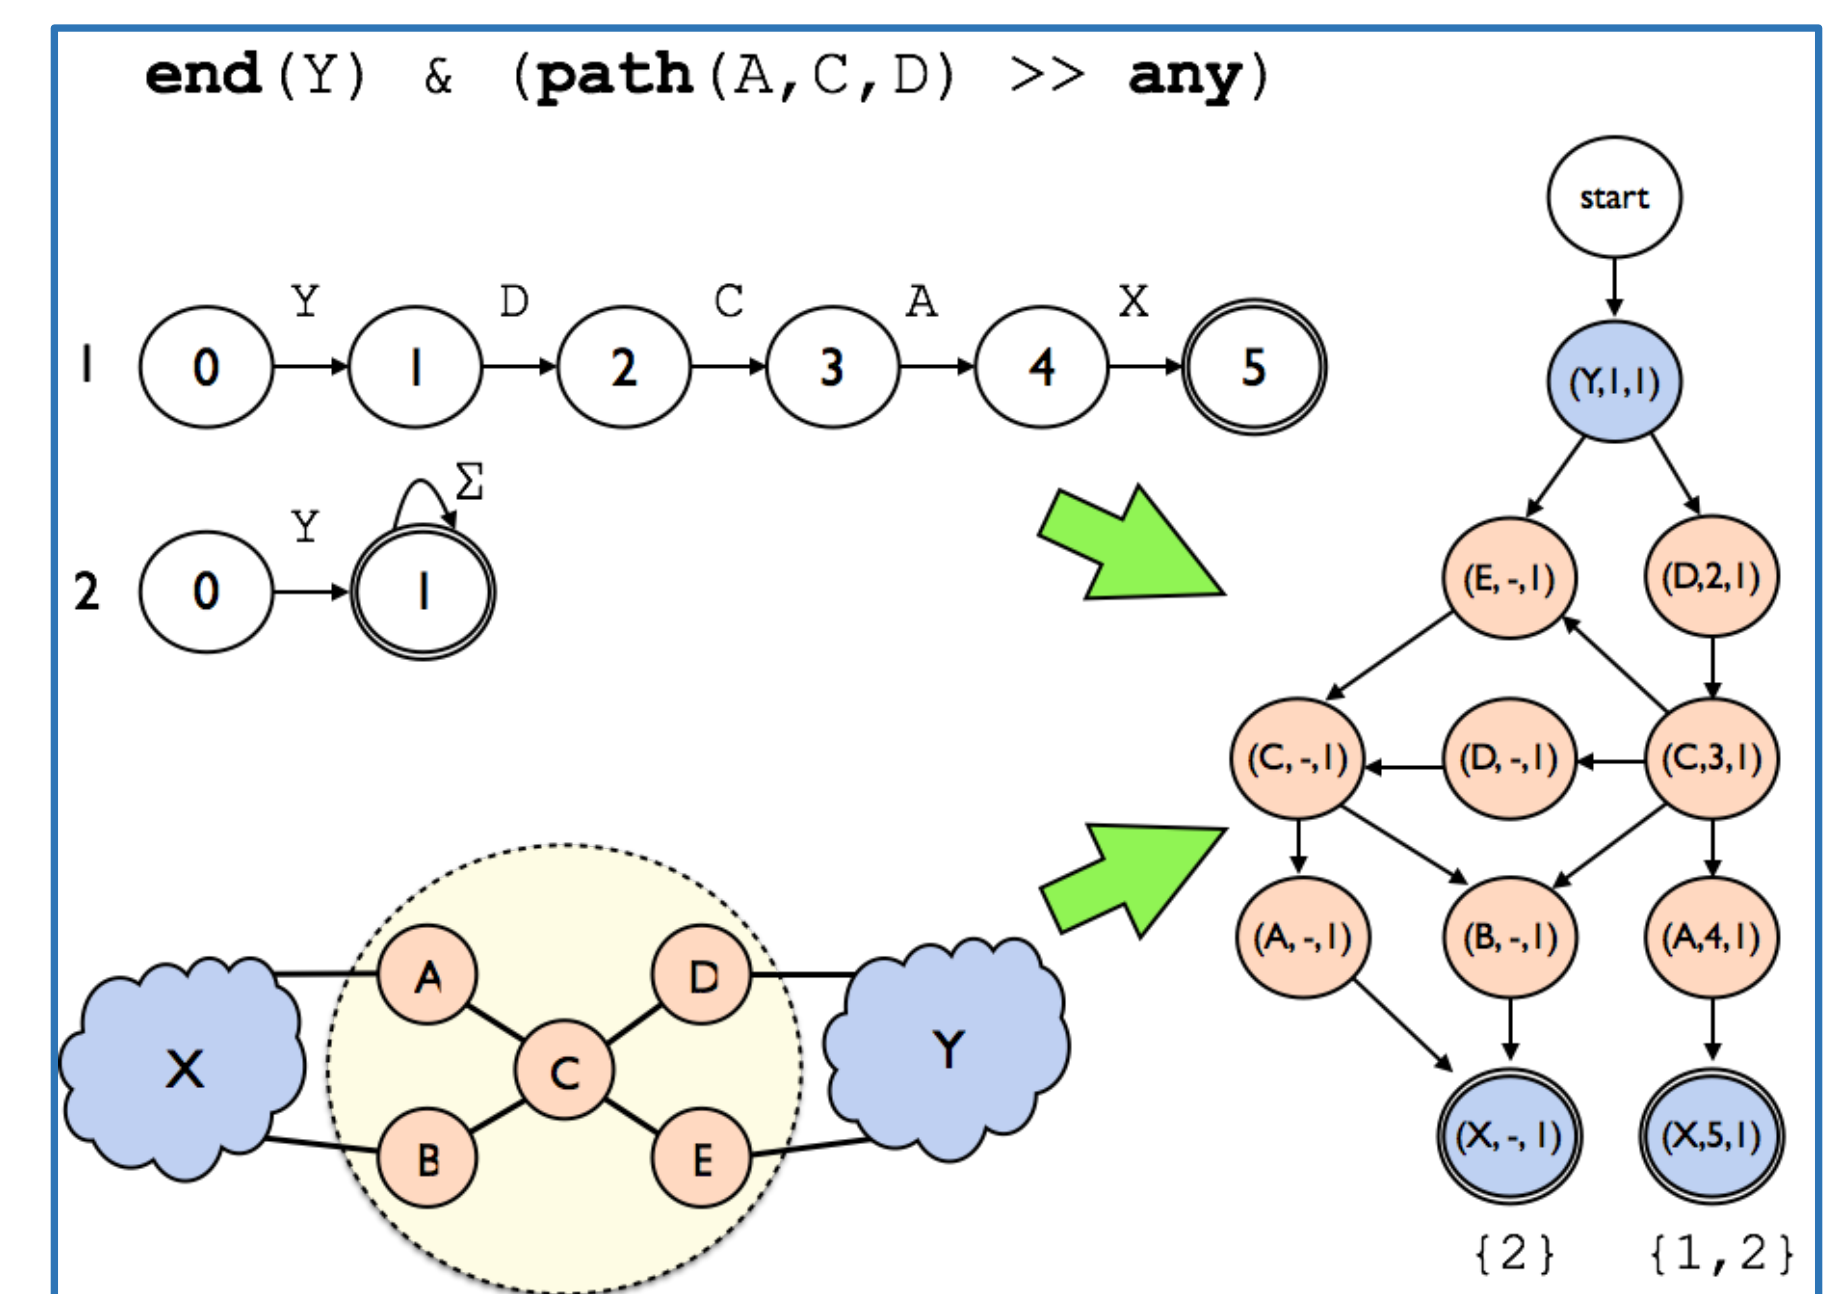


Figure 5. Product Graph construction

### Translation to BGP:

- Tag Product Graph state with BGP communities
- Incoming edges → import filters
- Outgoing edges → export filters
- BGP local preferences inferred from analysis

Ensures that BGP will always dynamically find **some best path** to each destination.

## Implementation

### Compiler:

- Compiler is 9500 lines of F# code
- Flags to configure using: no-export, MEDS, prepending, anycast, etc.
- Compilation parallelized across prefixes

### Benchmarks:

- Policy translated to Propane for datacenter and transit networks for a large cloud provider
- Topology scaled to determine the effect of topology size on compilation time

### Policy Size:

Lines of Propane: **31** for the datacenter, **43** for the transit network – orders of magnitude reduction.

### Compilation Time:

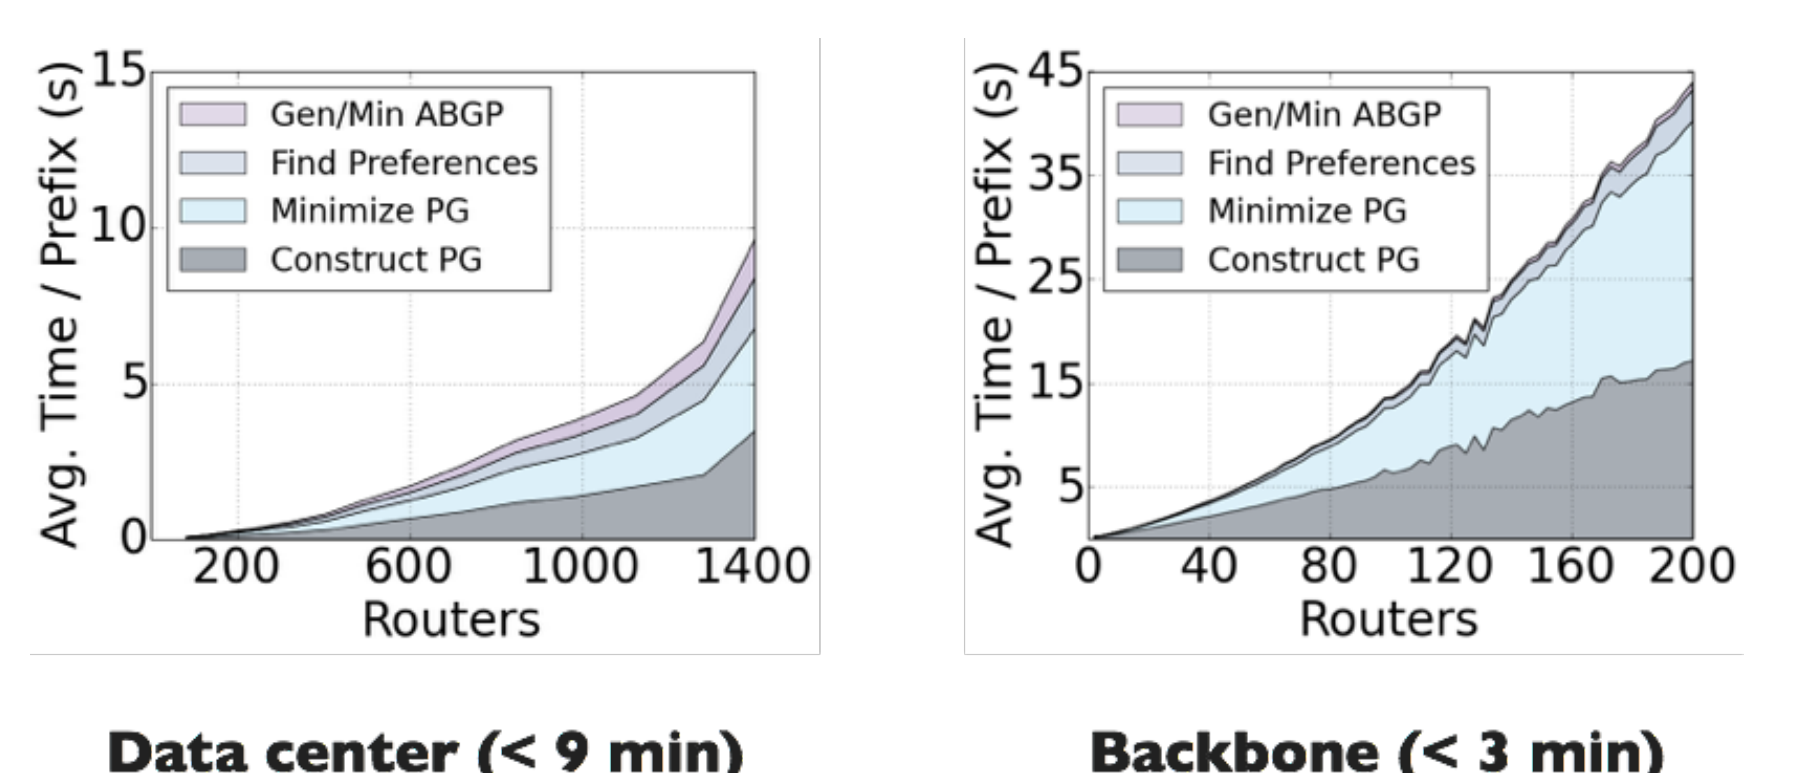


Figure 6. Compilation time for data center and backbone networks

## Contact

Ryan Beckett  
Princeton University  
rbeckett@princeton.edu  
(303) 956-6712

## References

1. M. Al-Fares, A. Loukissas, and A. Vahdat. A scalable commodity data center network architecture. In SIGCOMM, 2008.
2. N. Feamster and H. Balakrishnan. Detecting BGP configuration faults with static analysis. In NSDI, 2005.
3. R. Mahajan, D. Wetherall, and T. Anderson. Understanding BGP misconfiguration. In SIGCOMM, 2002.
4. Juniper. What's behind network downtime? <https://www-935.ibm.com/services/au/gts/pdf/200249.pdf>.
5. Yankee Group. As the value of enterprise networks escalates, so does the need for configuration management. <https://www.cs.princeton.edu/courses/archive/fall10/cos561/papers/Yankee04.pdf>.
6. J. McCauley, A. Panda, M. Casado, T. Koponen, and S. Shenker. Extending SDN to large-scale networks. In Open Networking Summit, 2013.
7. P. Berde, M. Gerola, J. Hart, Y. Higuchi, M. Kobayashi, T. Koide, B. Lantz, B. O'Connor, P. Radoslavov, W. Snow, and G. Parulkar. ONOS: towards an open, distributed SDN OS. In HotSDN, 2014.