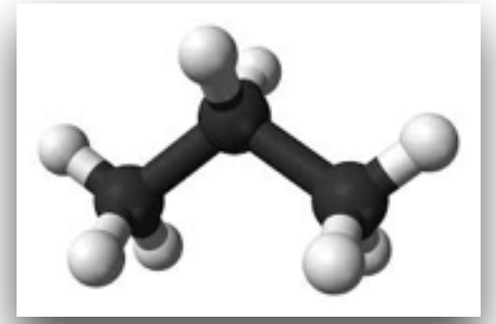

Propane

Programming Distributed Control Planes

Ryan Beckett (Princeton & MSR)
Ratul Mahajan (MSR)
Todd Millstein (UCLA)
Jitu Padhye (MSR)
David Walker (Princeton)

Configuring Networks is Error-Prone

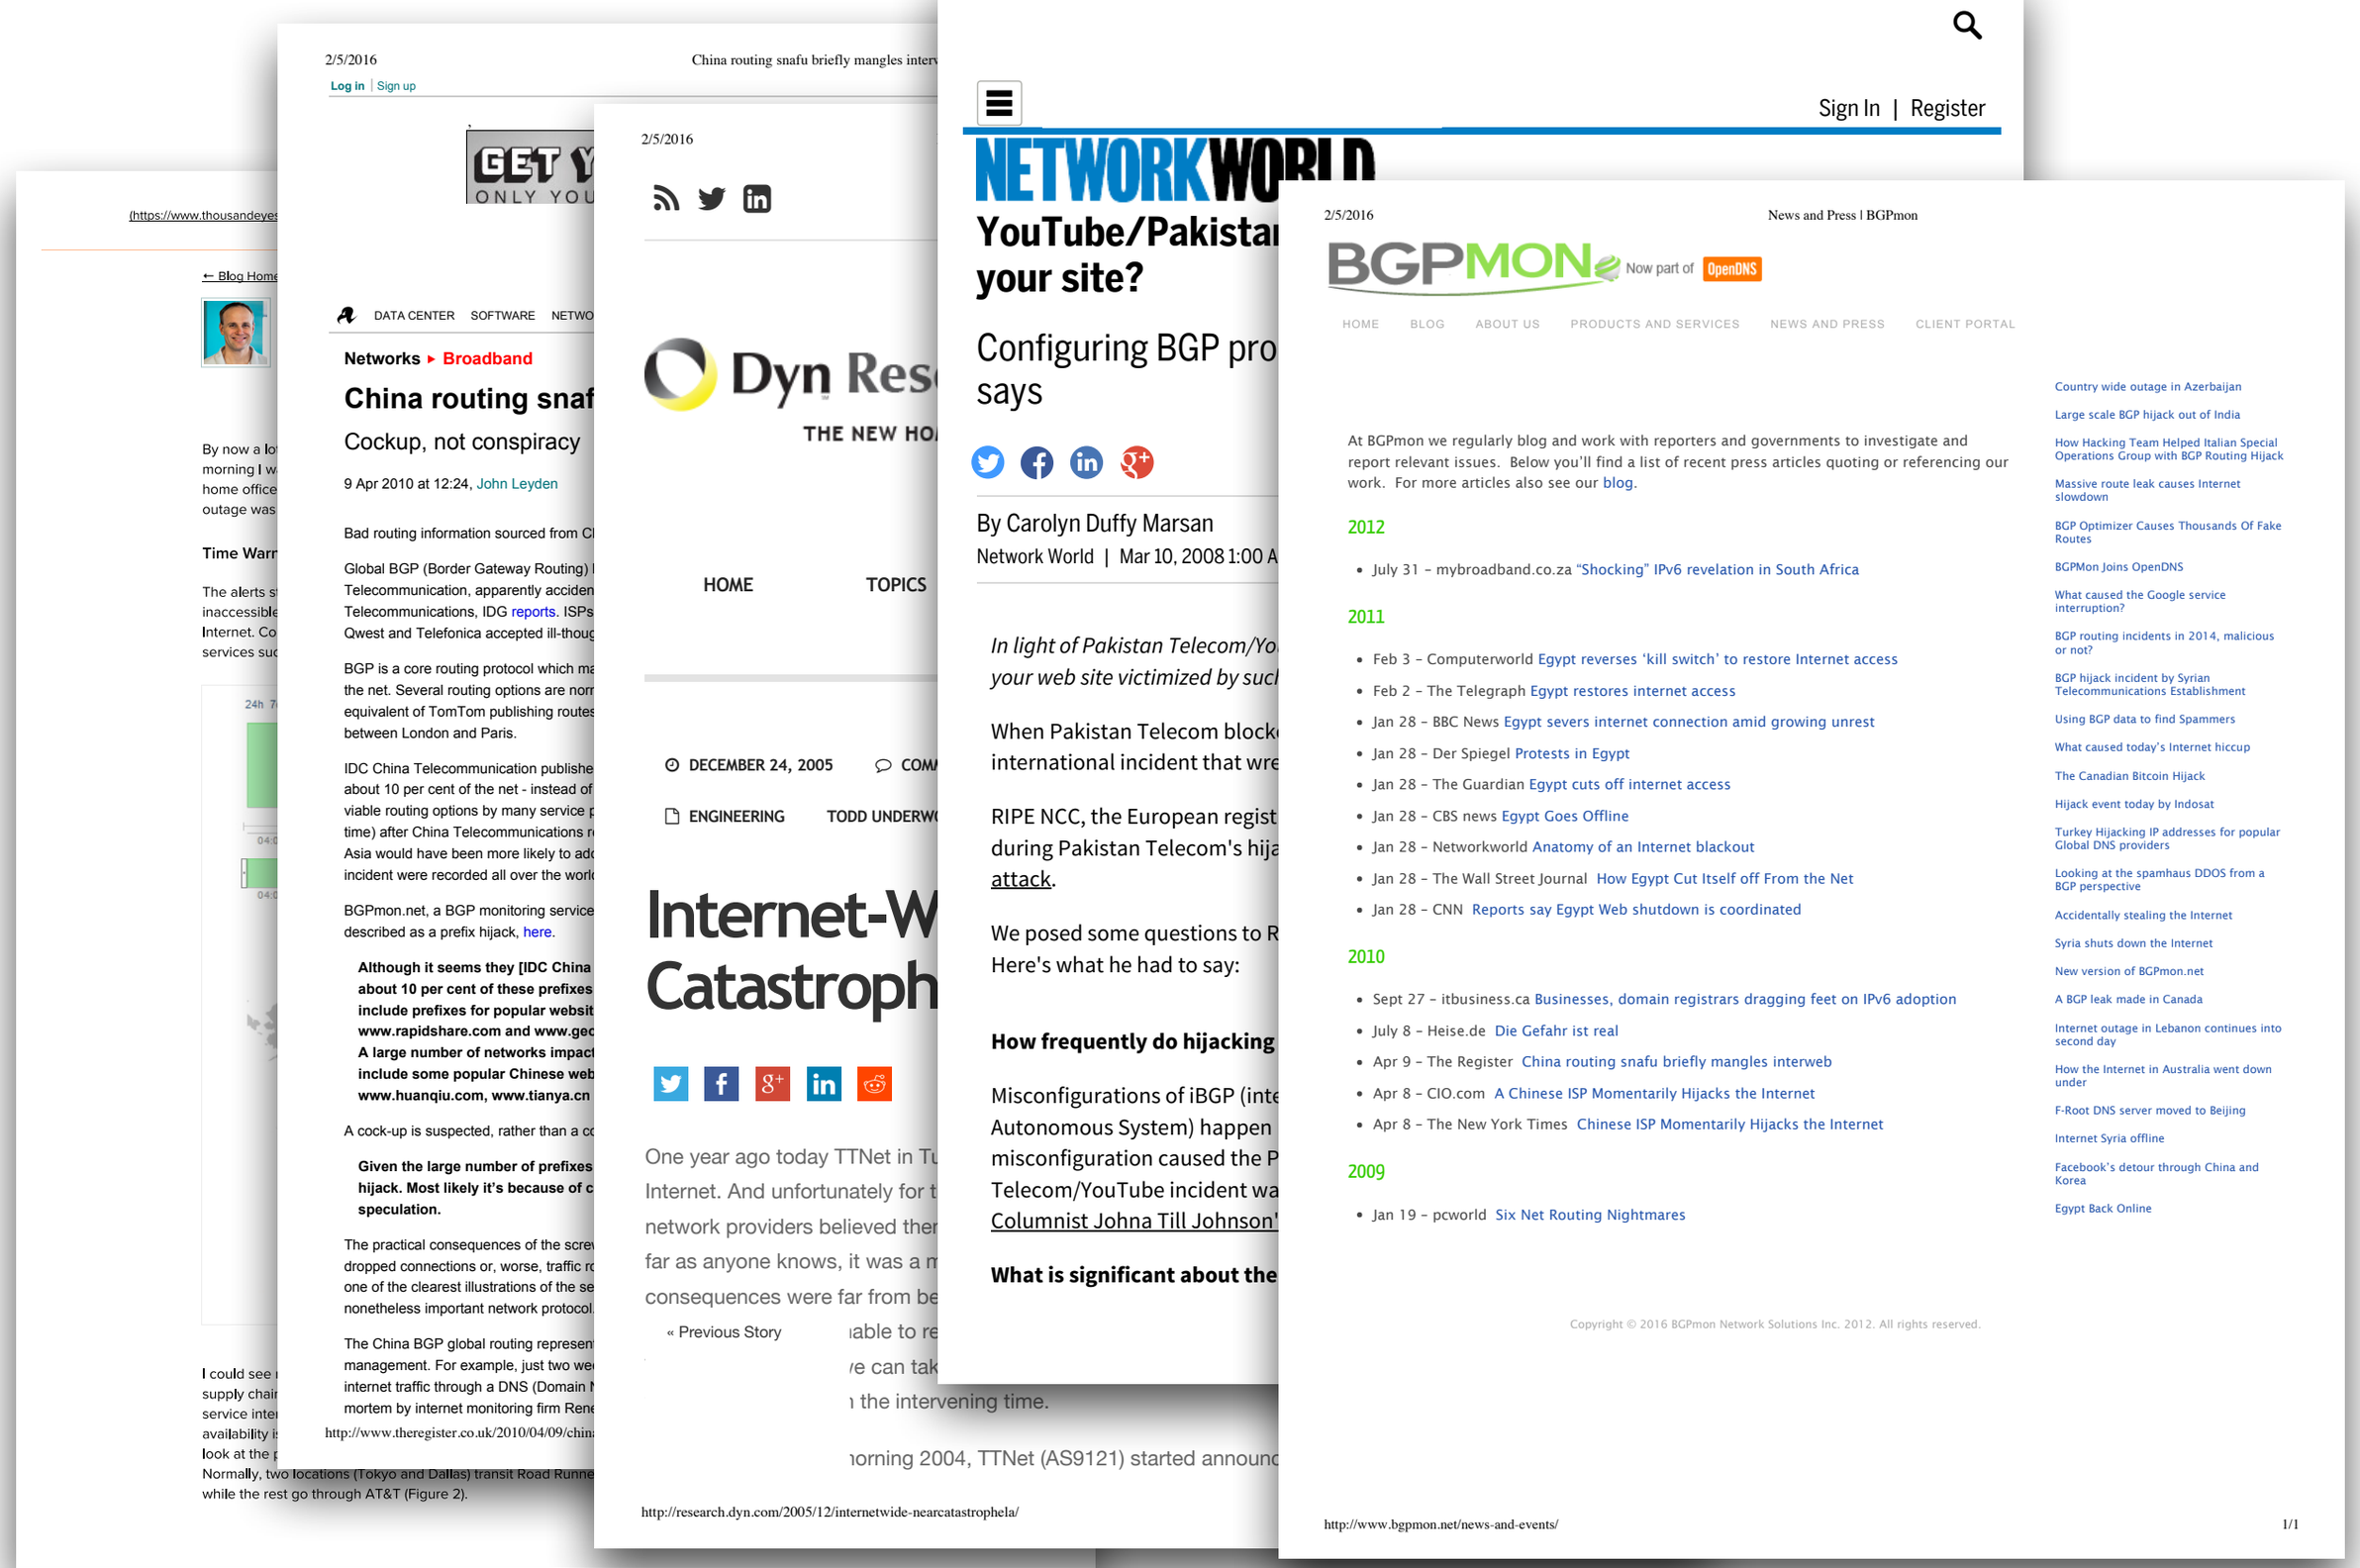

Objectives GAP **Mechanisms**

Objectives: Network-wide

- Prefer traffic to send traffic through customers over providers
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries

Mechanisms: Device-by-Device

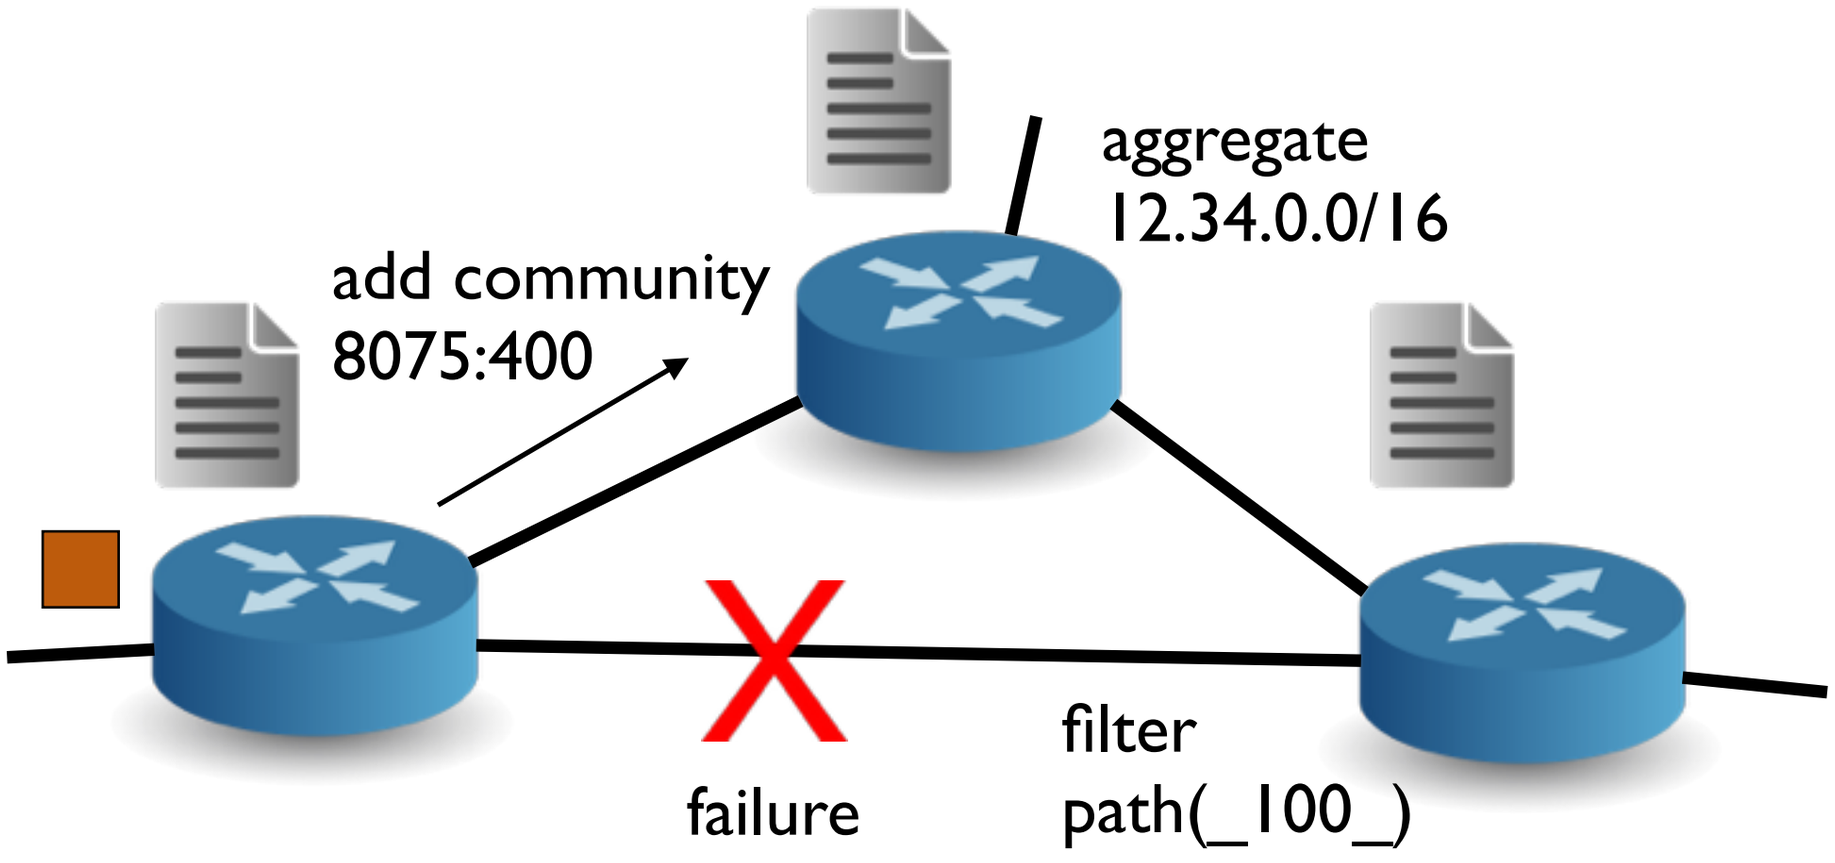

Propane:
Programming a Distributed Control Plane

1) Language for expressing high-level objectives with:

- Path constraints and relative preferences with fall-backs in case of failures
- Uniform abstractions for intra- and inter-domain routing

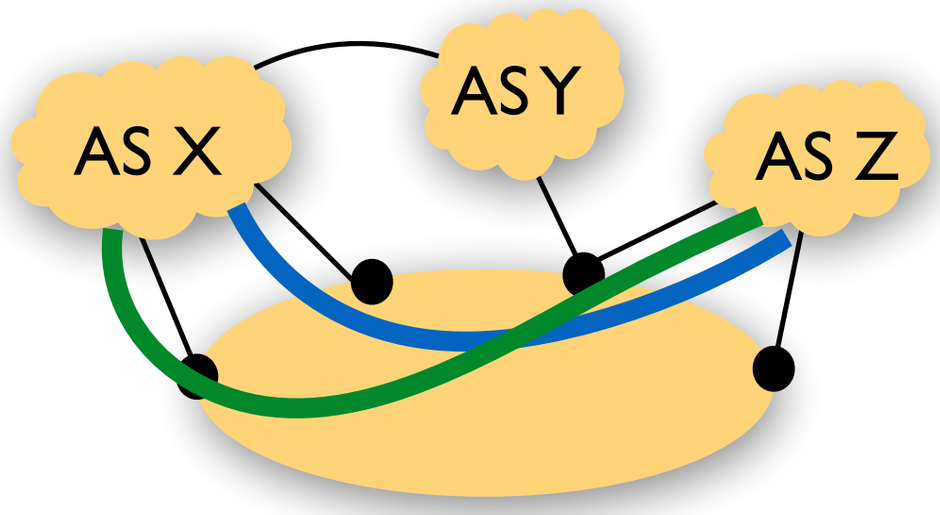

2) Compiler that generates low-level, distributed configs:

- Efficient algorithms to synthesize a set of policy-compliant BGP configs
- Failure analysis guarantees ***policy compliance*** under all failures

Example: A Data Center Network

A Data Center Network

Goals:

- P1: Announce global services externally as the aggregate PG
- P2: Do not announce local services externally
- P3: Prefer Backbone1 to Backbone2

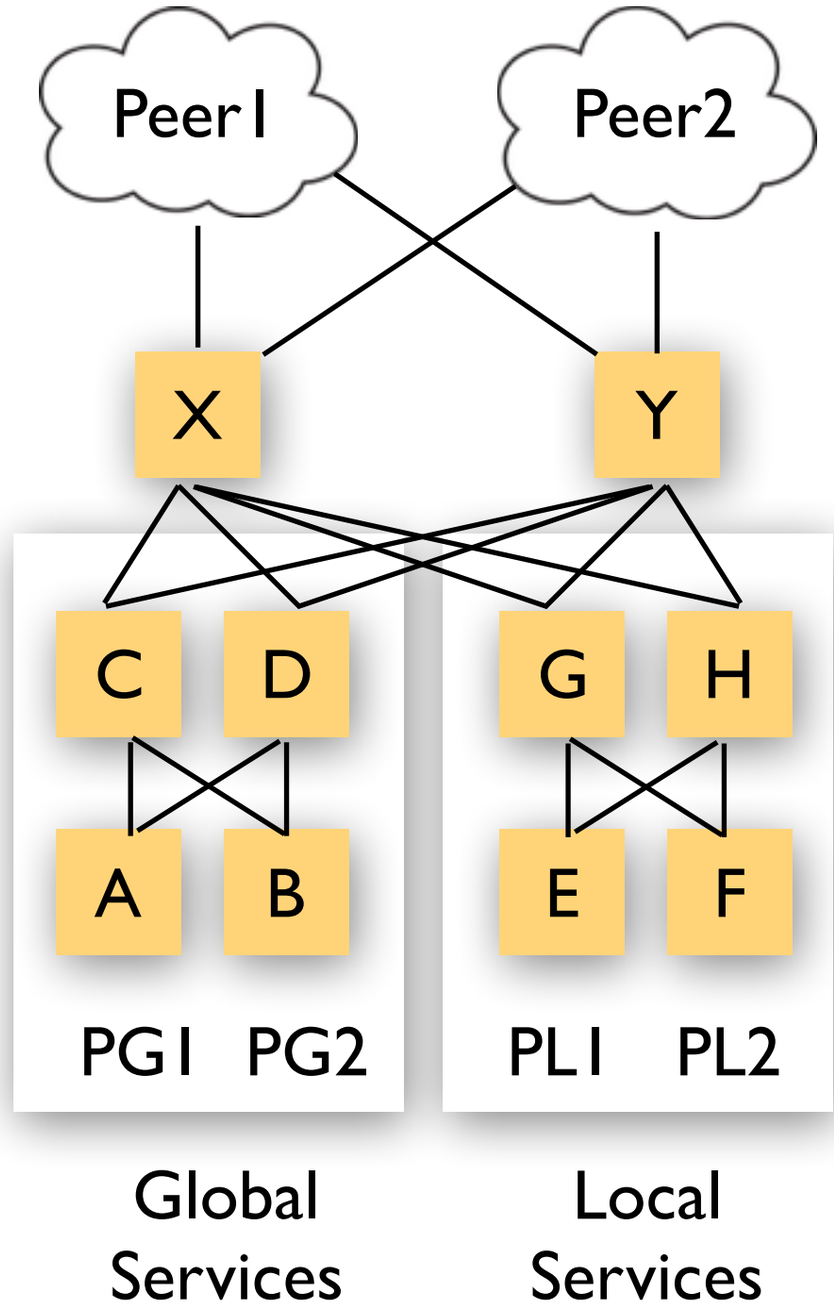

A Data Center Network

Goals:

- P1: Announce global services externally as the aggregate PG
- P2: Do not announce local services externally
- P3: Prefer Backbone1 to Backbone2

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X, Y do not need to know which prefixes are local vs global
- aggregate to PG if announce is subset of PG

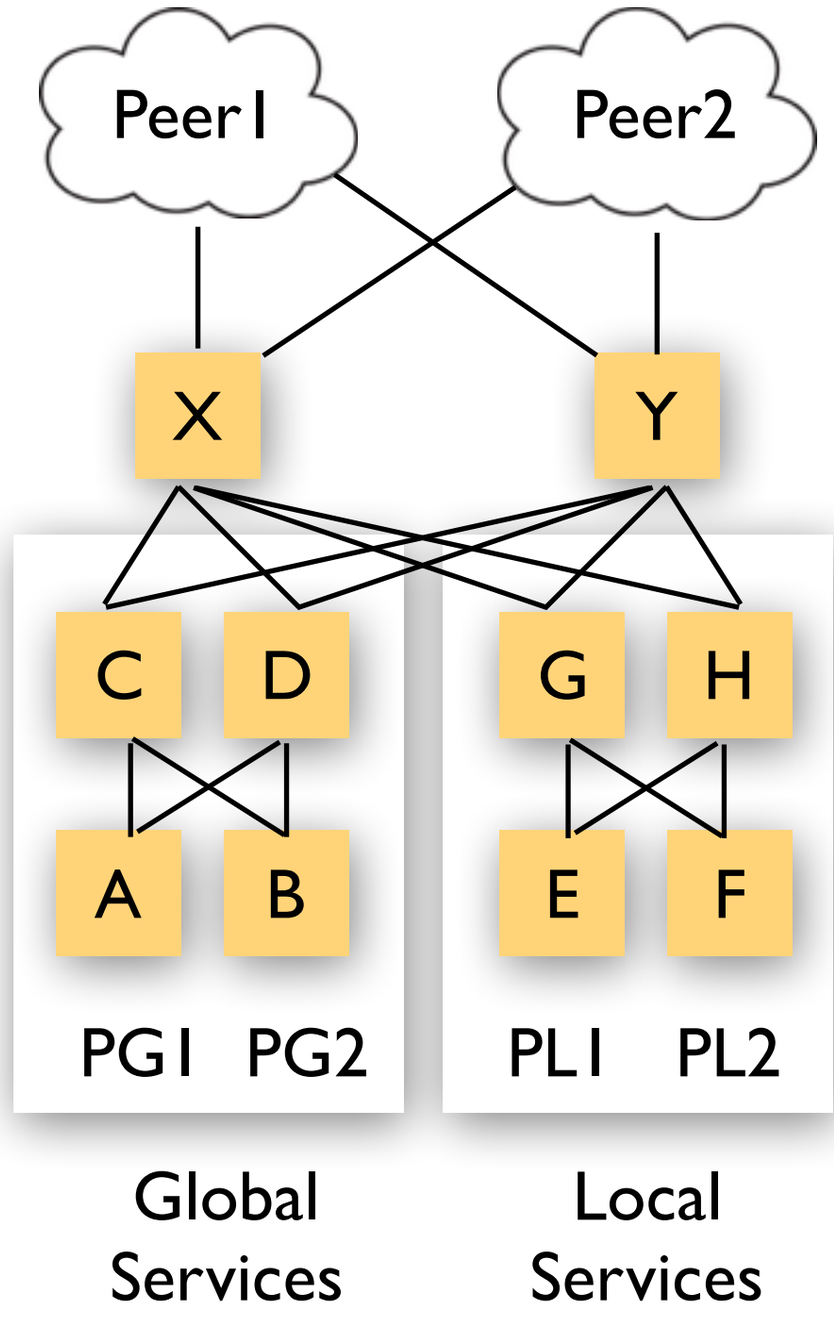

A Data Center Network

Goals:

- P1: Announce global services externally as the aggregate PG
- P2: Do not announce local services externally
- P3: Prefer Backbone1 to Backbone2

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X, Y do not need to know which prefixes are local vs global
- aggregate to PG if announce is subset of PG

Consider X-G, X-H Failure:

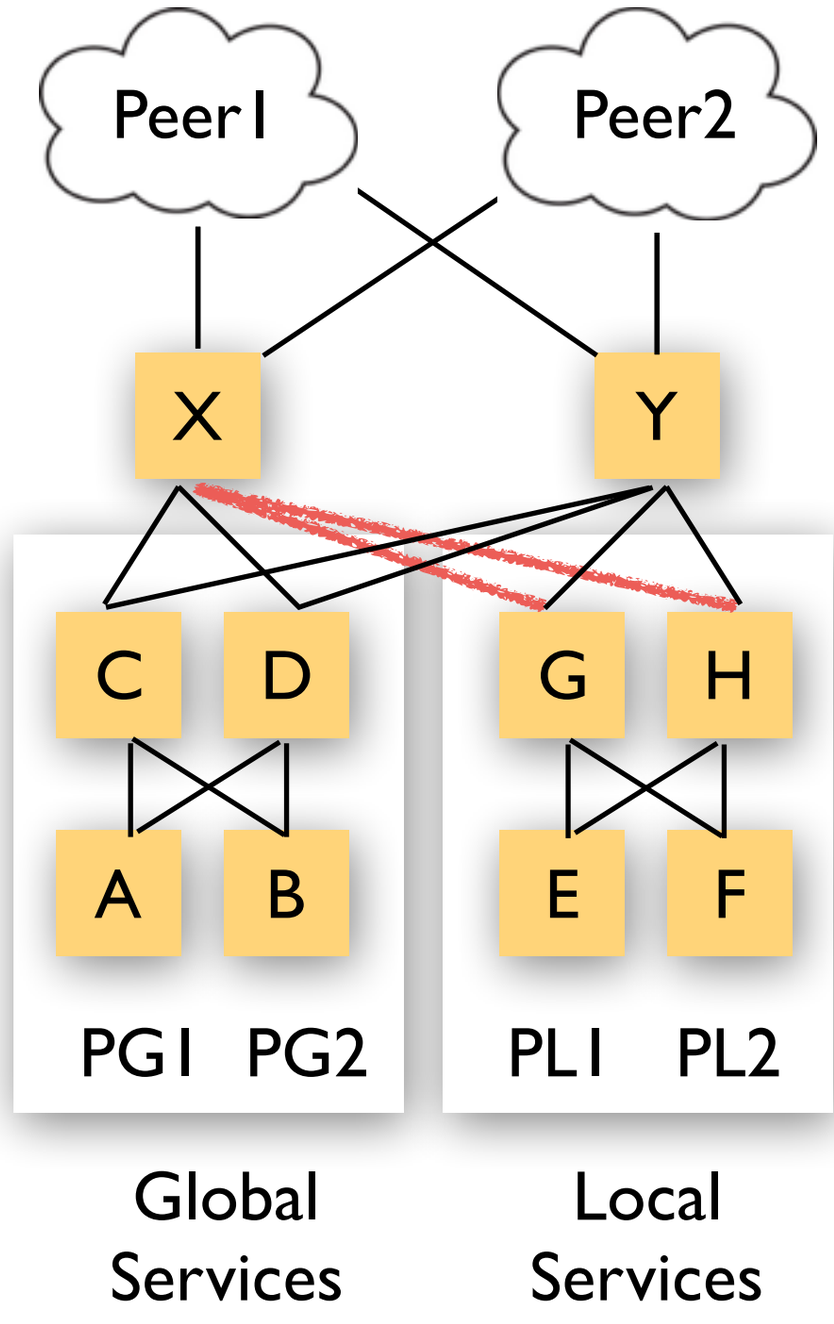

A Data Center Network

Goals:

- P1: Announce global services externally as the aggregate PG
- P2: Do not announce local services externally
- P3: Prefer Backbone1 to Backbone2

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X, Y do not need to know which prefixes are local vs global
- aggregate to PG if announce is subset of PG

Consider X-G, X-H Failure:

- PL* announcements travel H-Y-D-X
- PL* announcements are then leaked

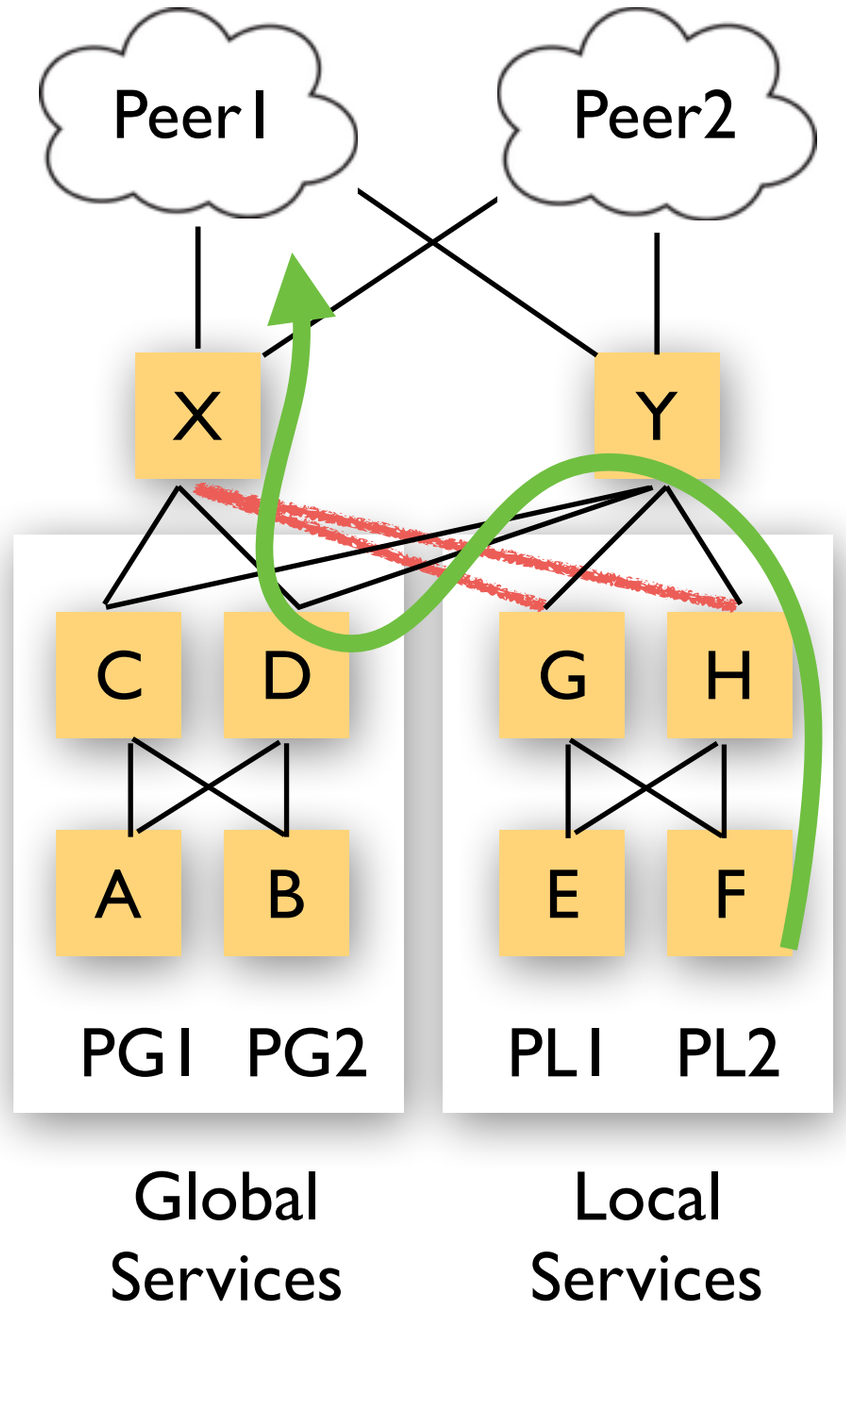

A Data Center Network

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X,Y do not need to know which prefixes are local vs global
- aggregate to PG if announce is subset of PG
- **disallow "valley" paths**

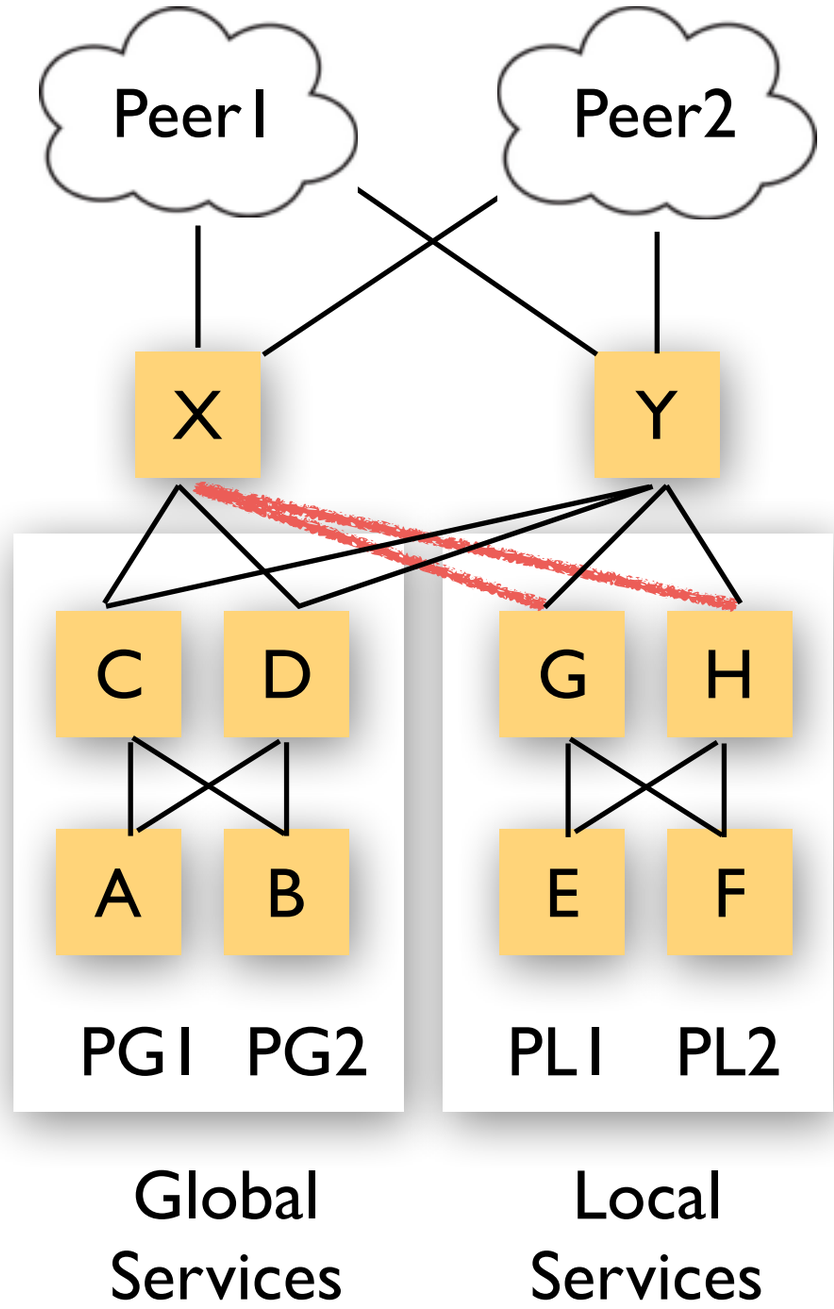

A Data Center Network

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X,Y do not need to know which prefixes are local vs global
  - aggregate to PG if announce is subset of PG
  - **disallow "valley" paths**

Consider D-A, X-C Failure:

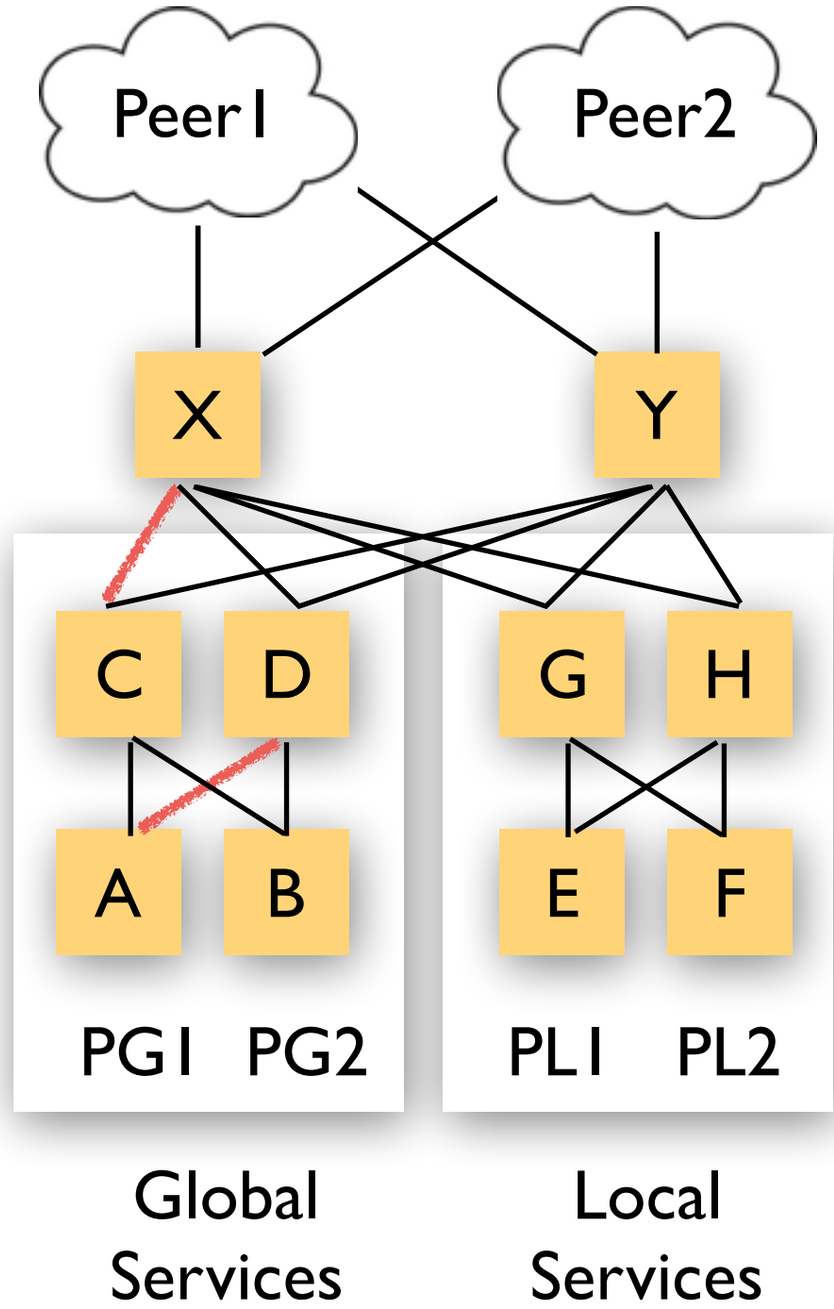

A Data Center Network

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X, Y do not need to know which prefixes are local vs global
  - aggregate to PG if announce is subset of PG
  - **disallow "valley" paths**

Consider D-A, X-C Failure:

- X and Y will hear PG2

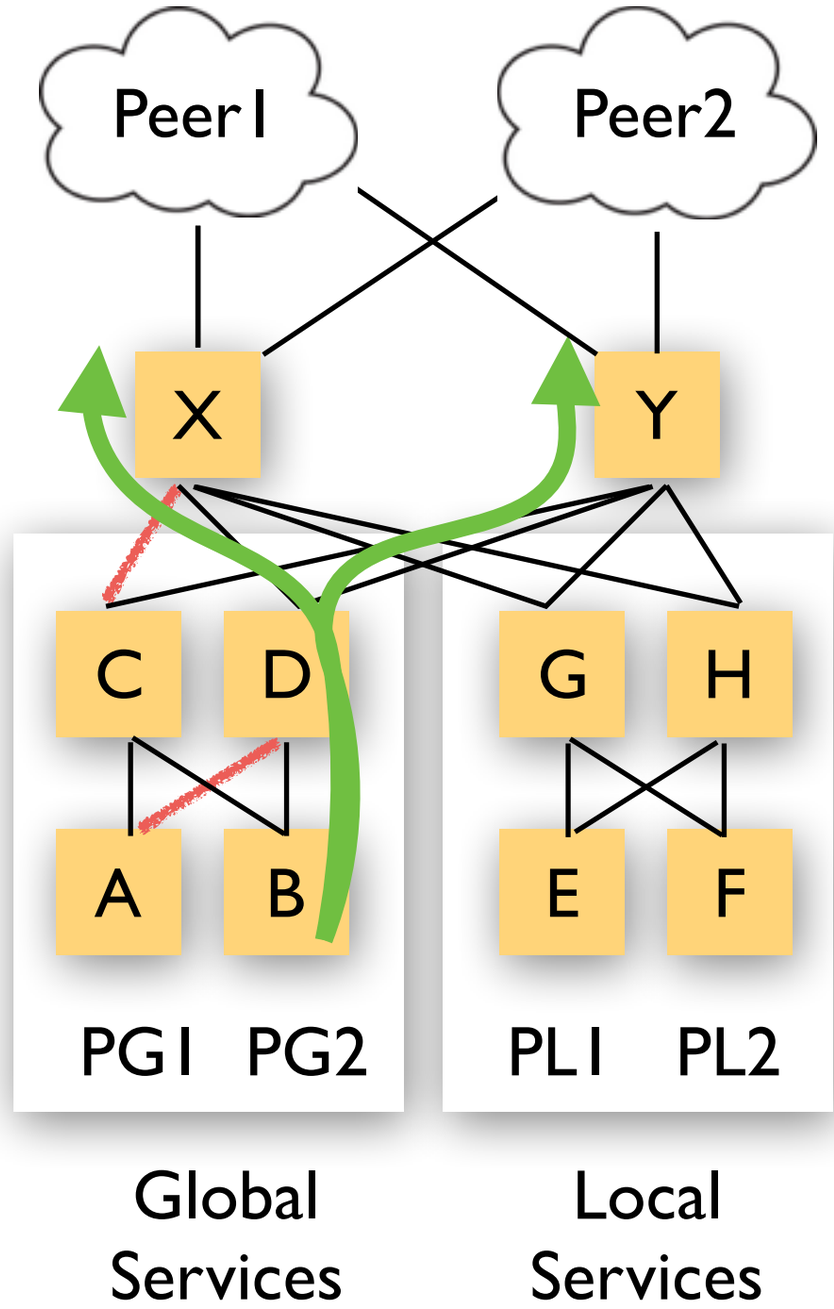

A Data Center Network

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X, Y do not need to know which prefixes are local vs global
  - aggregate to PG if announce is subset of PG
  - **disallow "valley" paths**

Consider D-A, X-C Failure:

- X and Y will hear PG2
- X and Y will announce aggregate PG

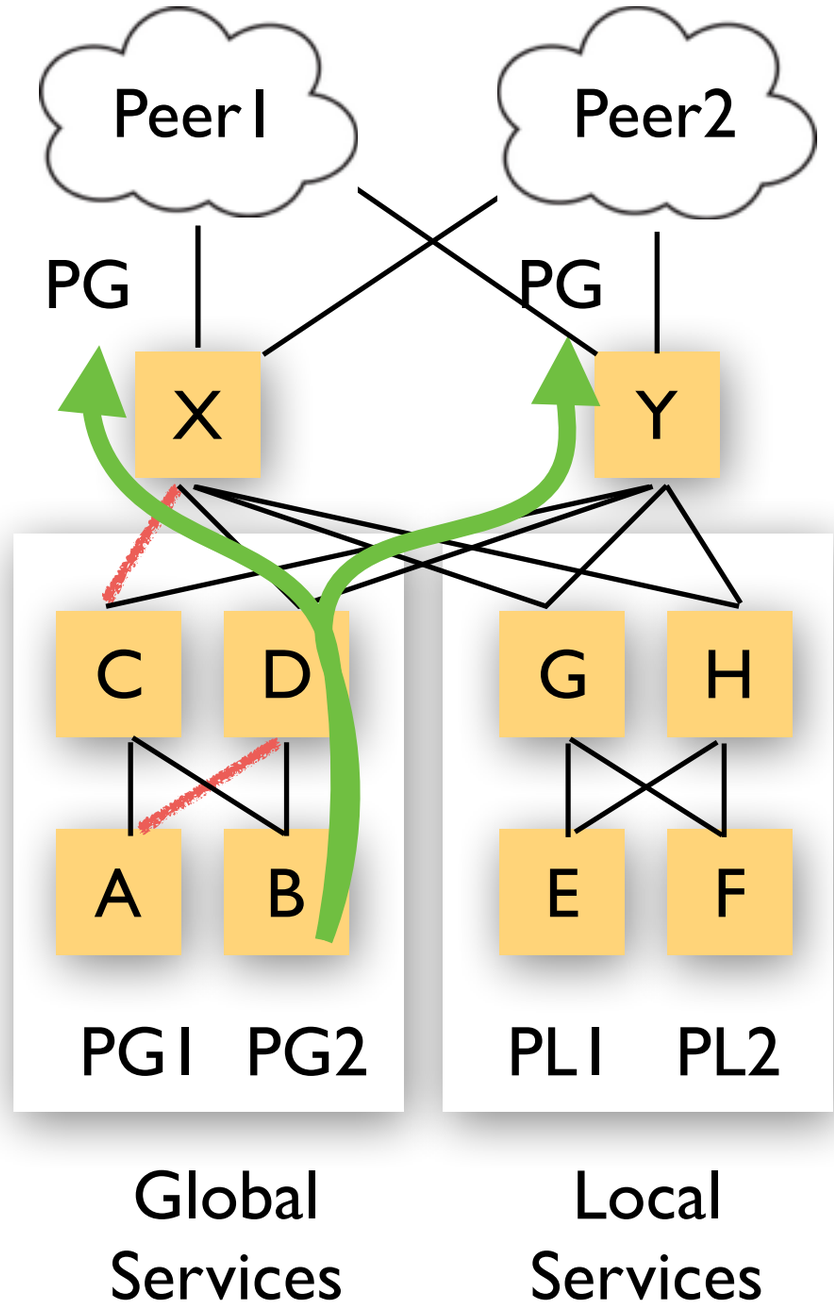

A Data Center Network

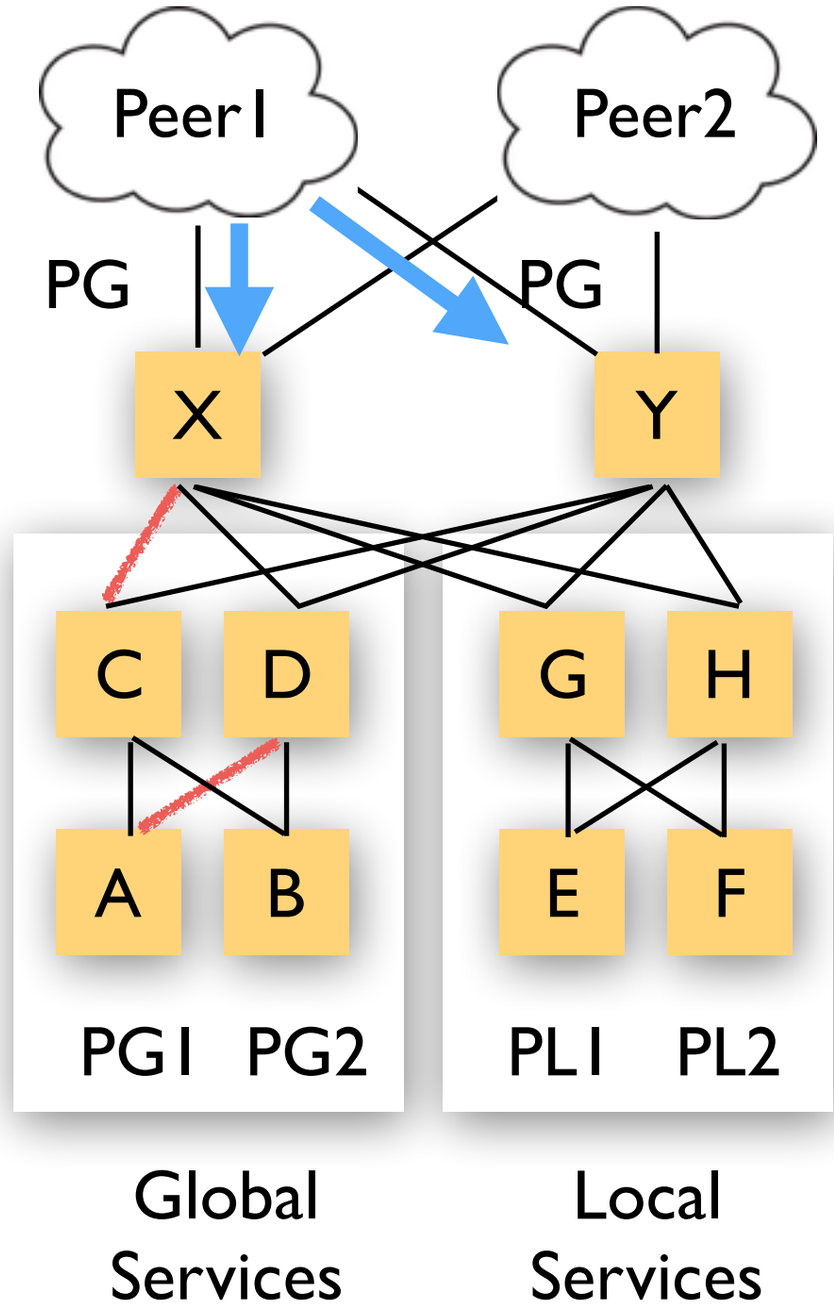

Implementation Techniques for X, Y:

- do export announce's from C, D outside
- do *not* export announce's from G,H outside
  - appeal: X,Y do not need to know which prefixes are local vs global
- aggregate to PG if announce is subset of PG
- **disallow "valley" paths**

Consider D-A, X-C Failure:

- X and Y will hear PG2
- X and Y will announce aggregate PG
- But PG1 is inaccessible through X because there is no valley routing
- An aggregation-induced black hole is created [See Le et al, CoNext '11]

Quick Demo

- Originate prefixes for each TOR router
- Do not announce local services externally
- Aggregation on global prefixes
- Prefer Peer1 over Peer2
- Prevent transit between Back1 and Back2

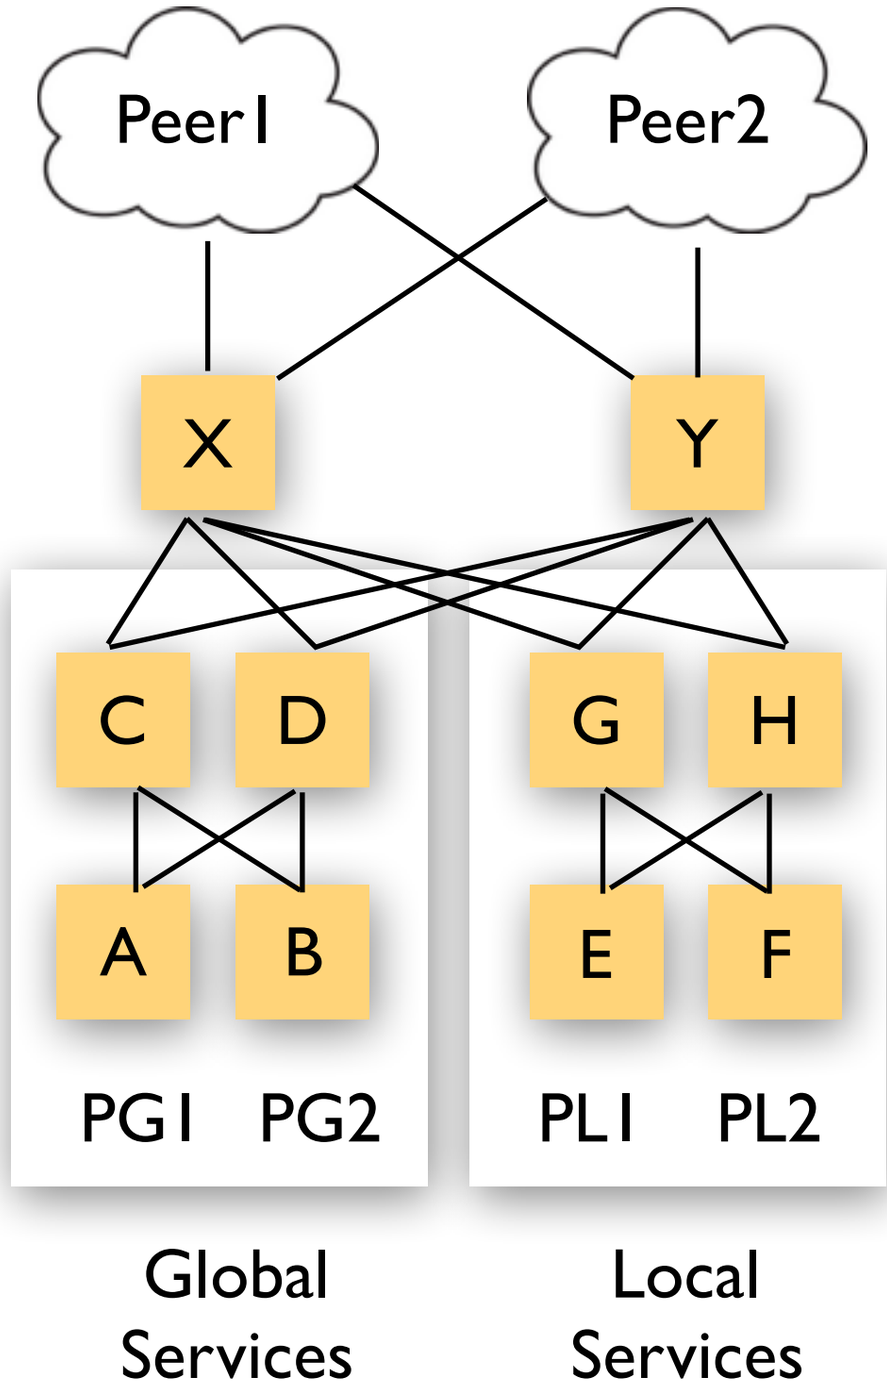

Compiling Propane

Propane Compiler

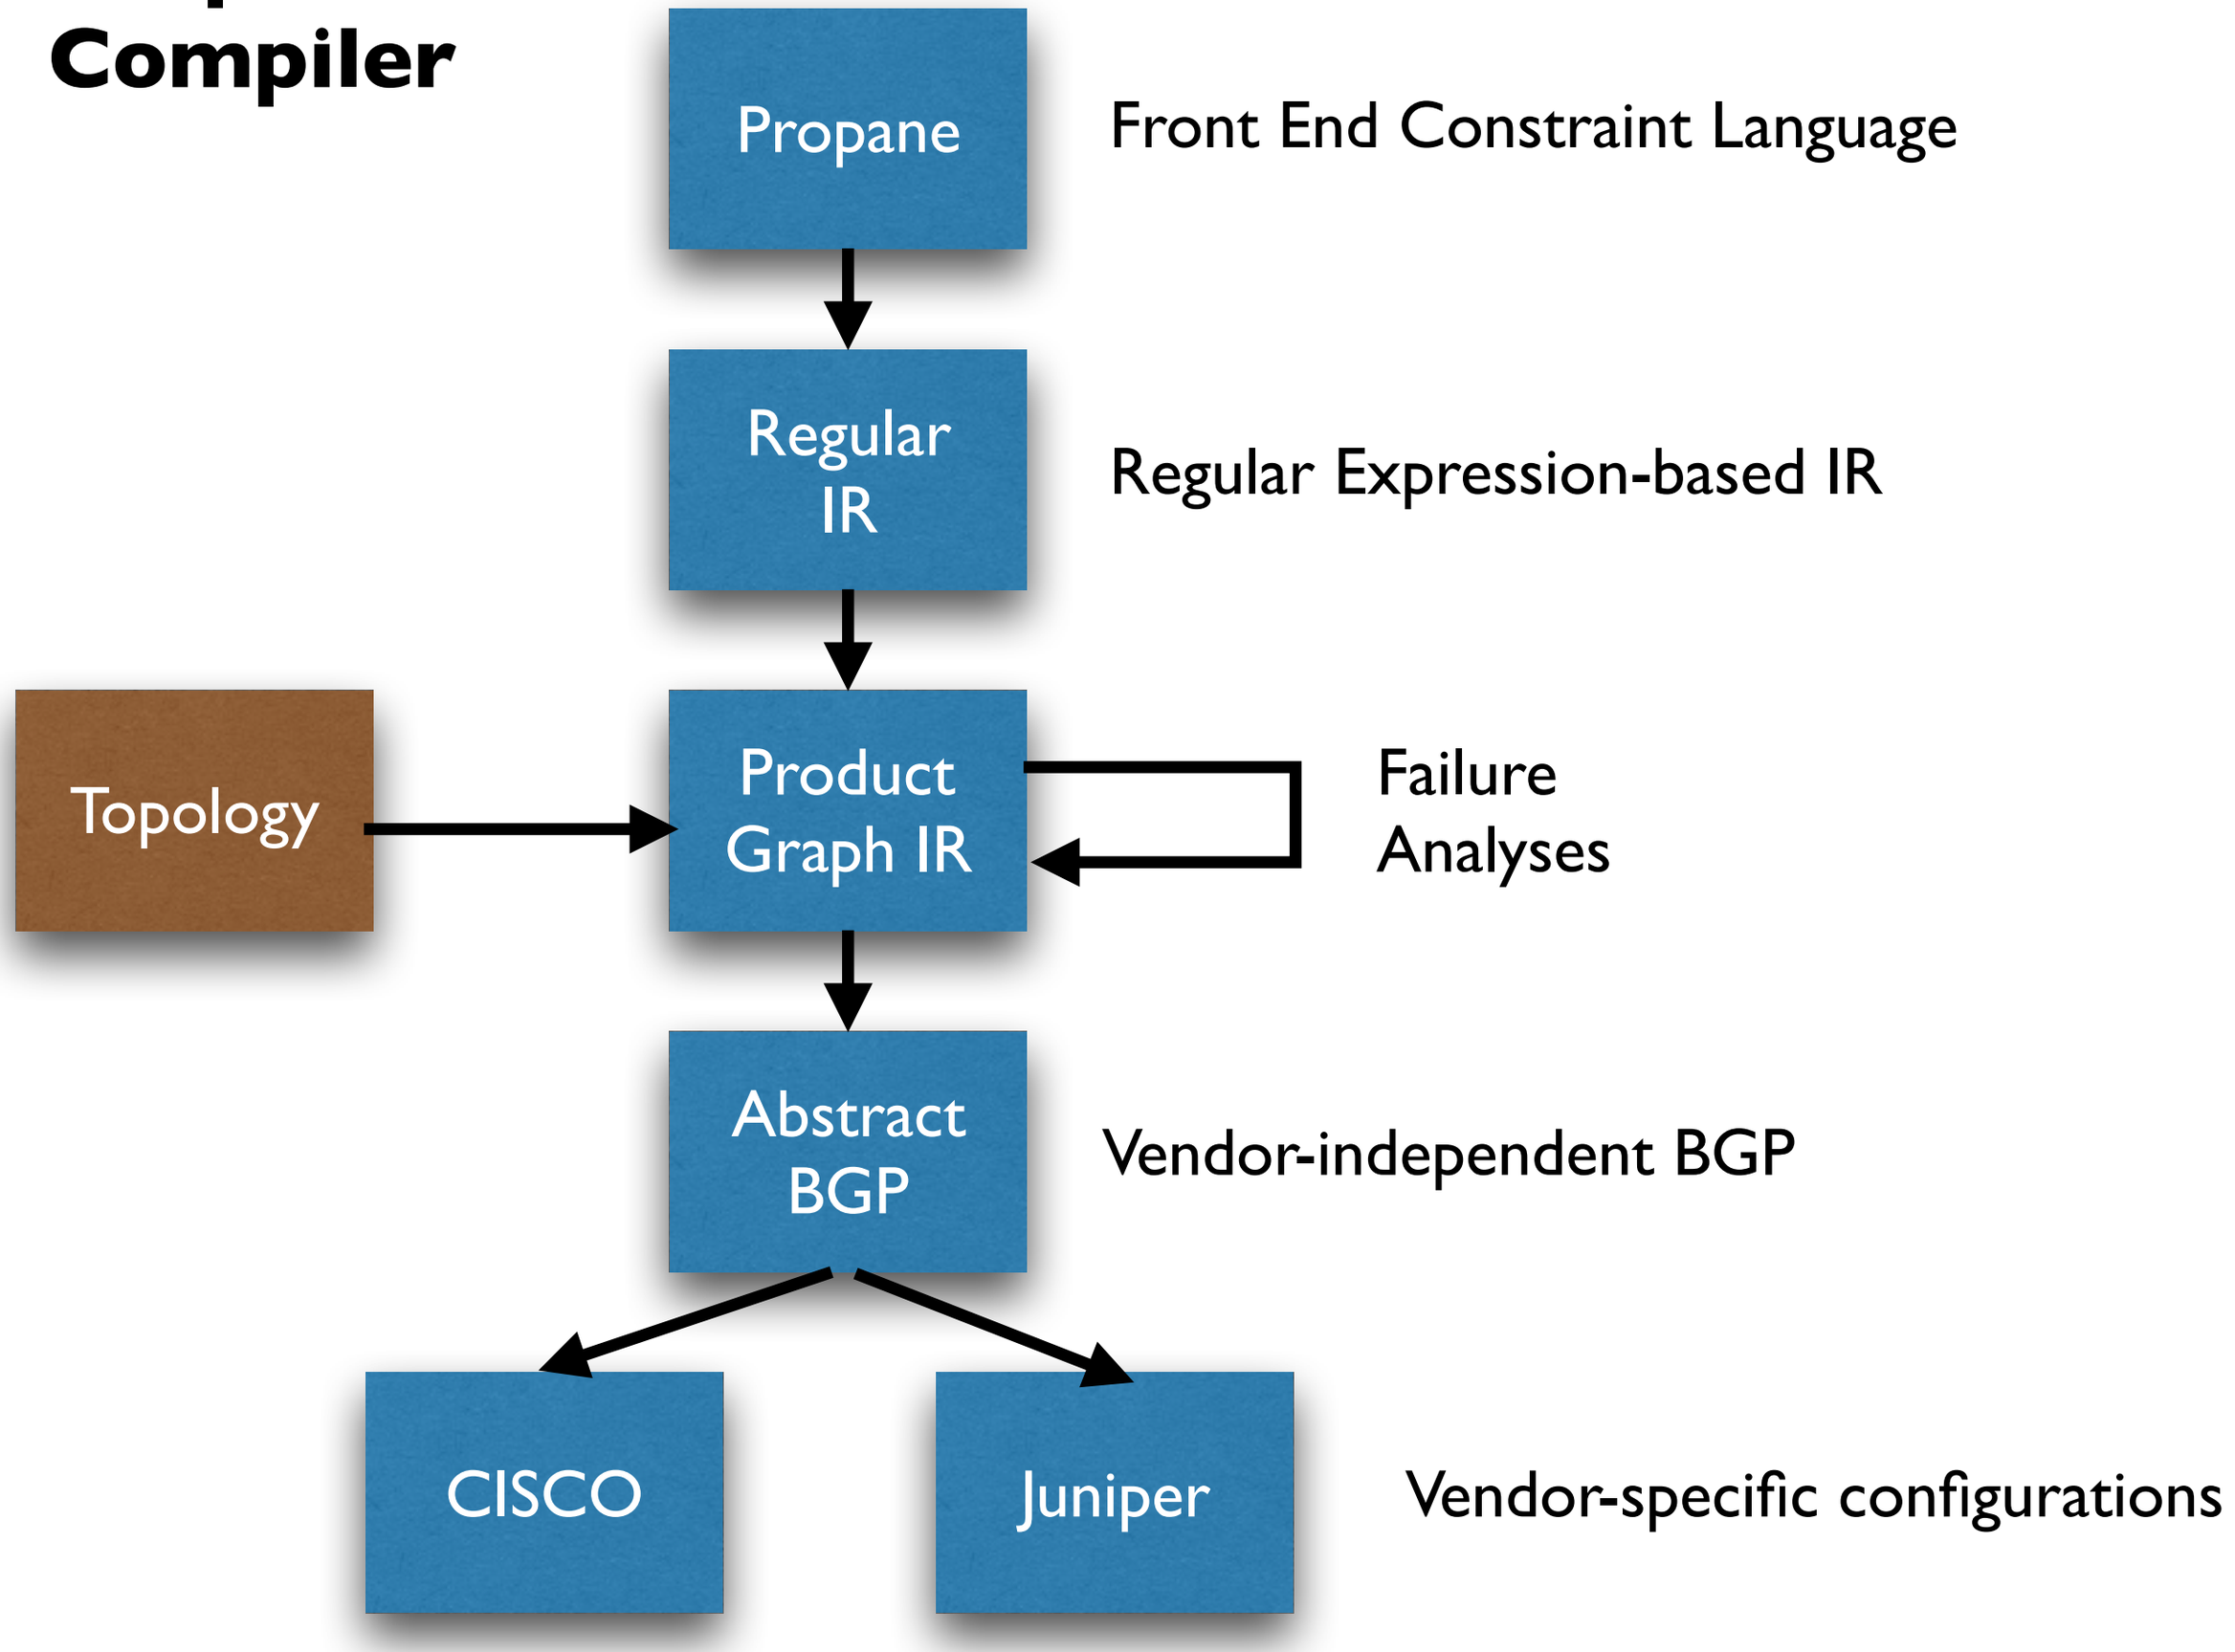

Propane Regular IR

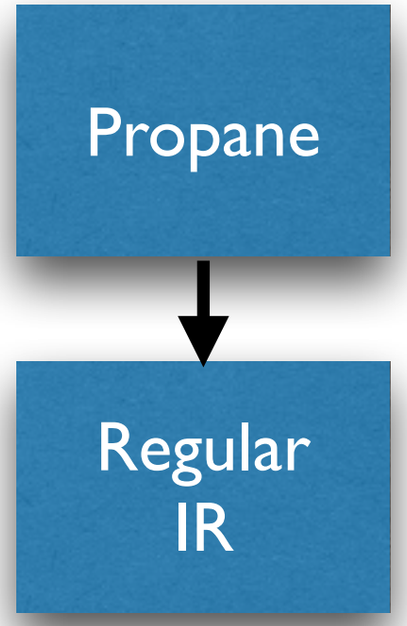

Expand constraints in to regular expressions. EG:

$$\mathbf{end}(X) = (\Sigma^*.X)$$

$$\mathbf{exit}(X) = (\mathbf{out}^*.\mathbf{in}^*.(X \cap \mathbf{in}).\mathbf{out}+)\cup (\mathbf{out}^*.\mathbf{in}+.(X \cap \mathbf{out}).\mathbf{out}^*)$$

A few other simple transformations:

- conjunction of constraints ==> intersection of regular expressions
- conjunction of policies ==> prefix-by-prefix intersection
- nested preferences lifted: (x >> y) . z ==> (x.z) >> (y.z)

Constructing the Product Graph (PG)

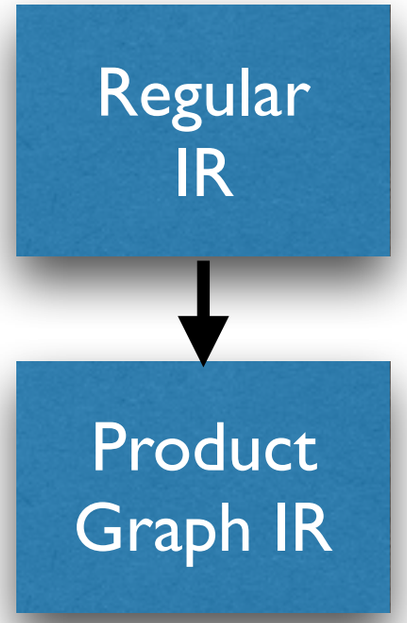

(a) Topology

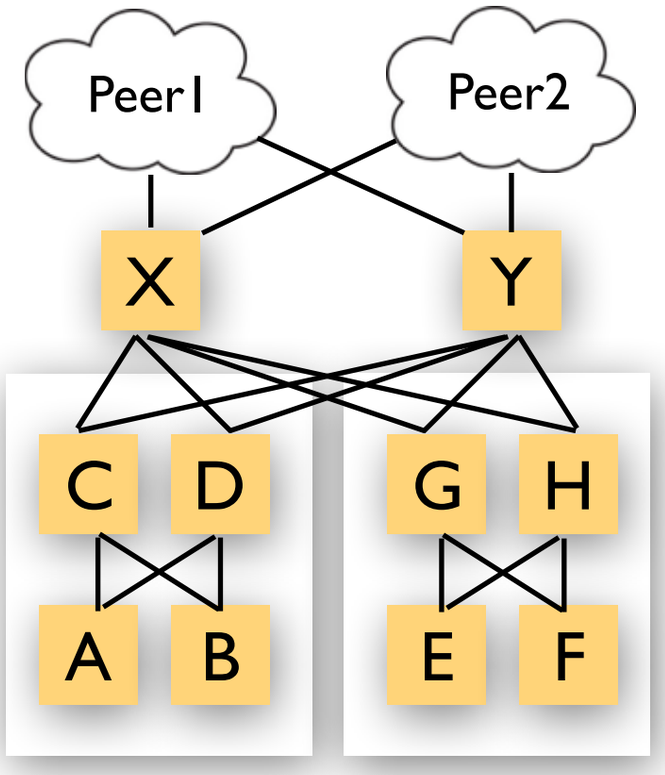

(b) Policy

1)

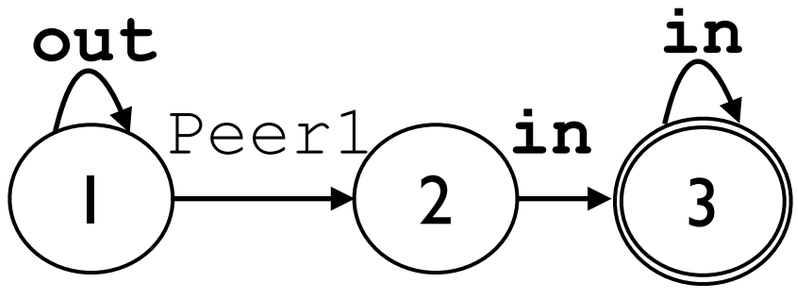

2)

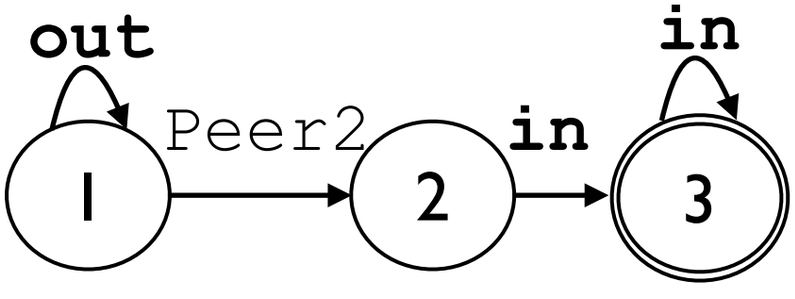

Constructing the Product Graph (PG)

(a) Topology

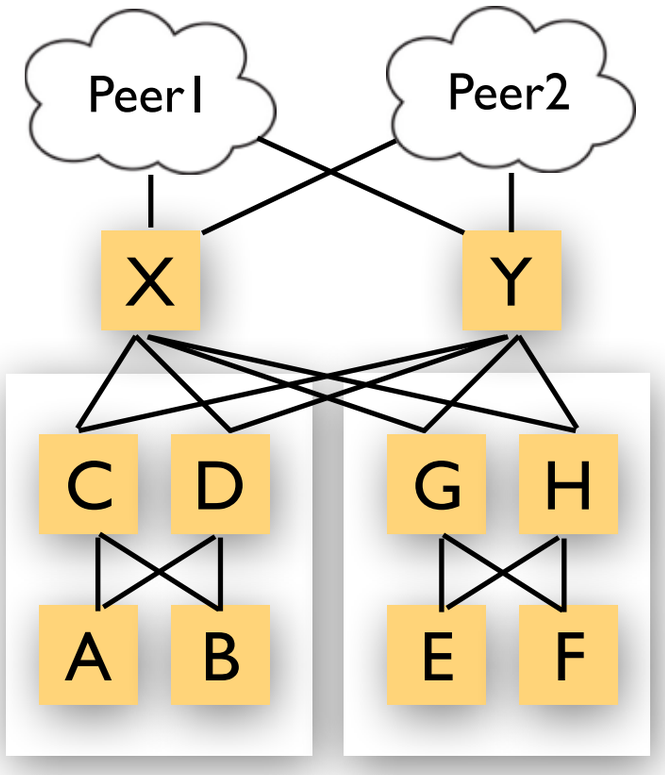

(b) Policy

1)

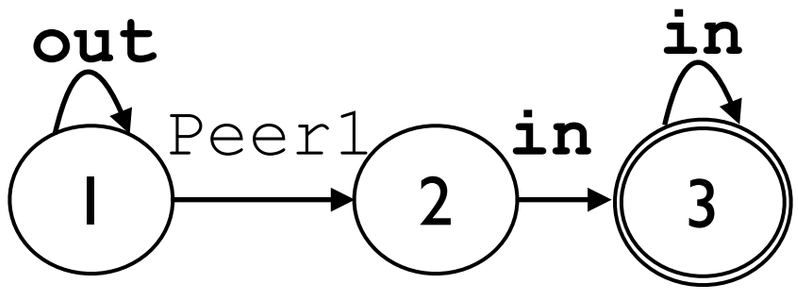

2)

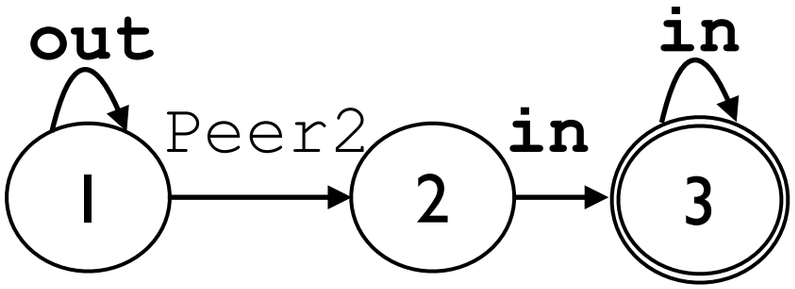

General Idea:
PG represents locations reachable in the topology while following the policy

Each PG node contains:

- topology node (S)
- state of automaton 1 (a)
- state of automaton 2 (b)
- set of preferences achieved

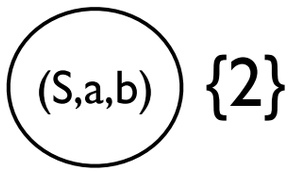

Two PG nodes are connected if topology nodes are connected and the automata make the specified transition

Constructing the Product Graph (PG)

(a) Topology

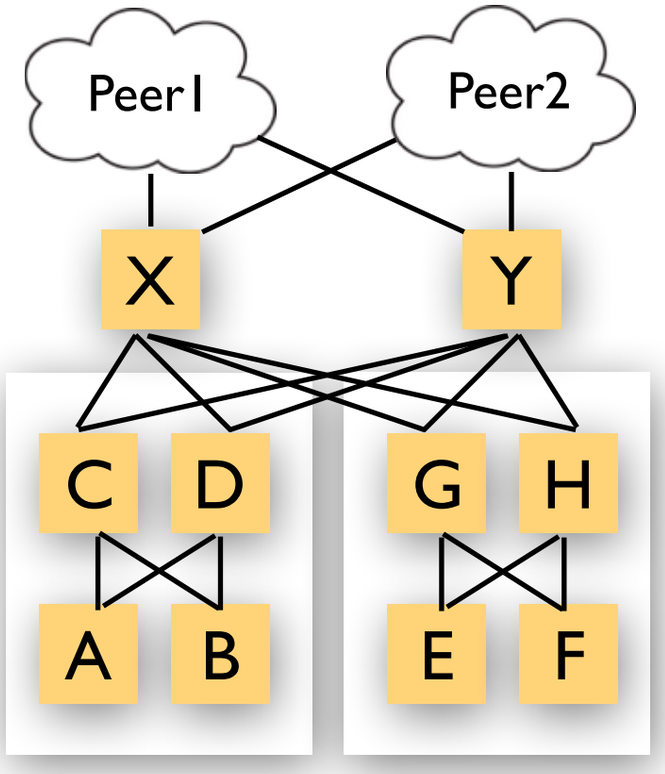

(b) Policy

1)

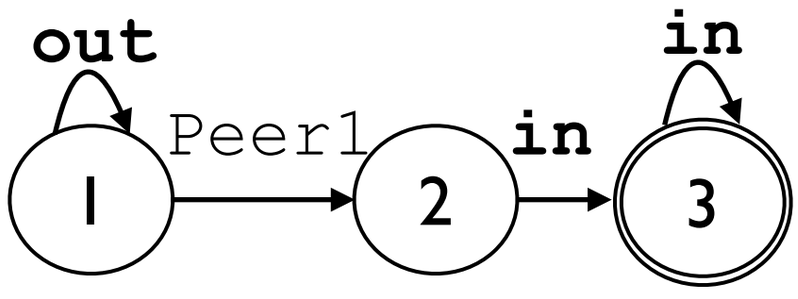

2)

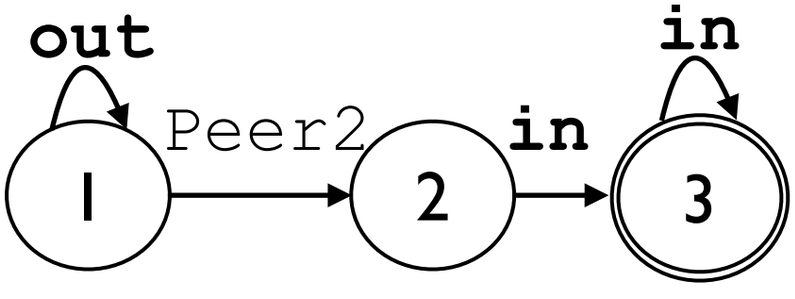

(c) Product Graph

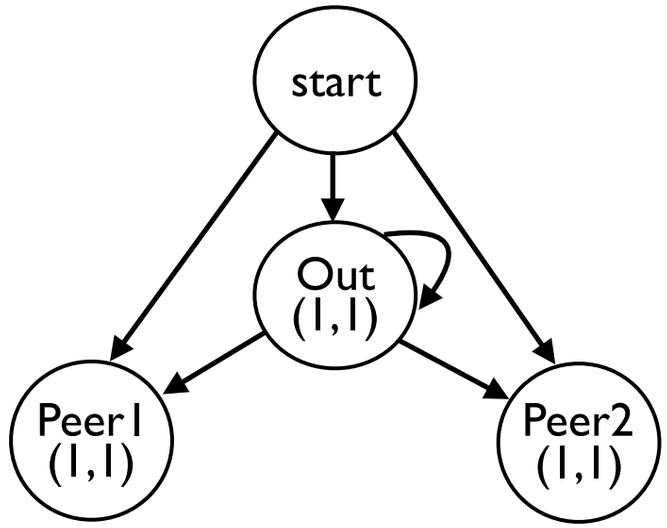

Constructing the Product Graph (PG)

(a) Topology

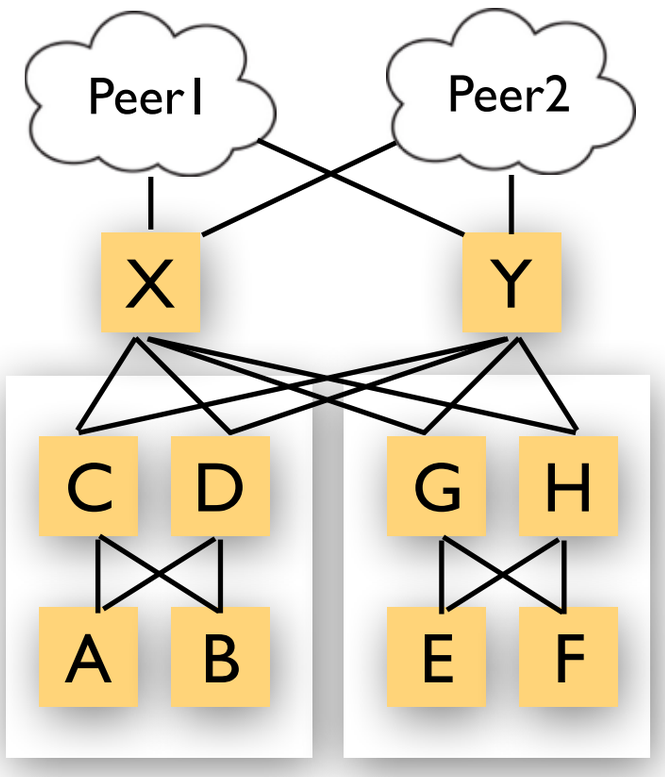

(b) Policy

1)

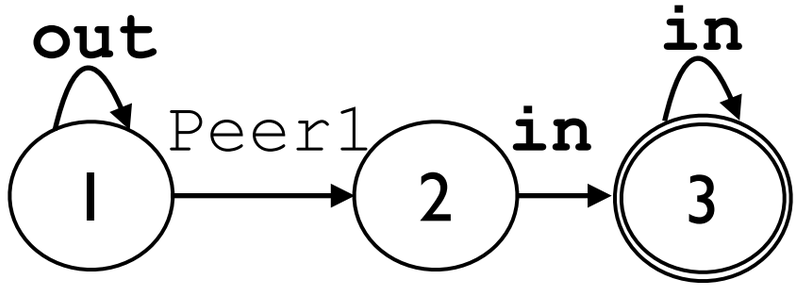

2)

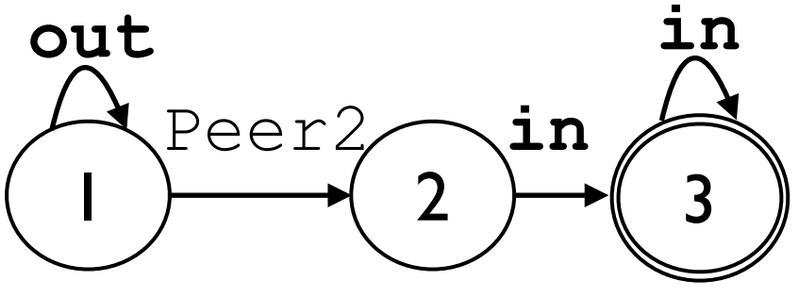

(c) Product Graph

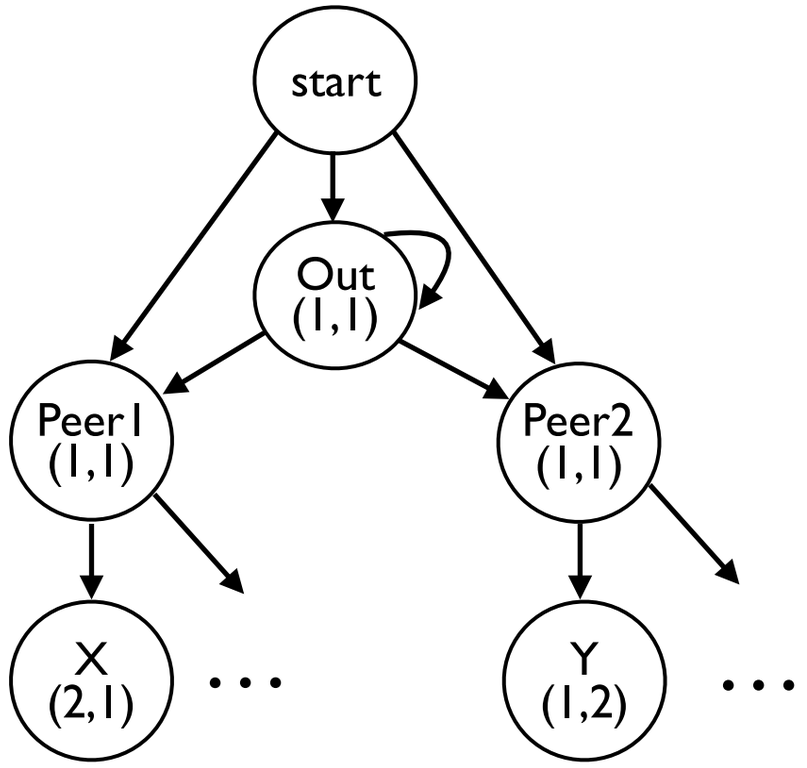

Constructing the Product Graph (PG)

(a) Topology

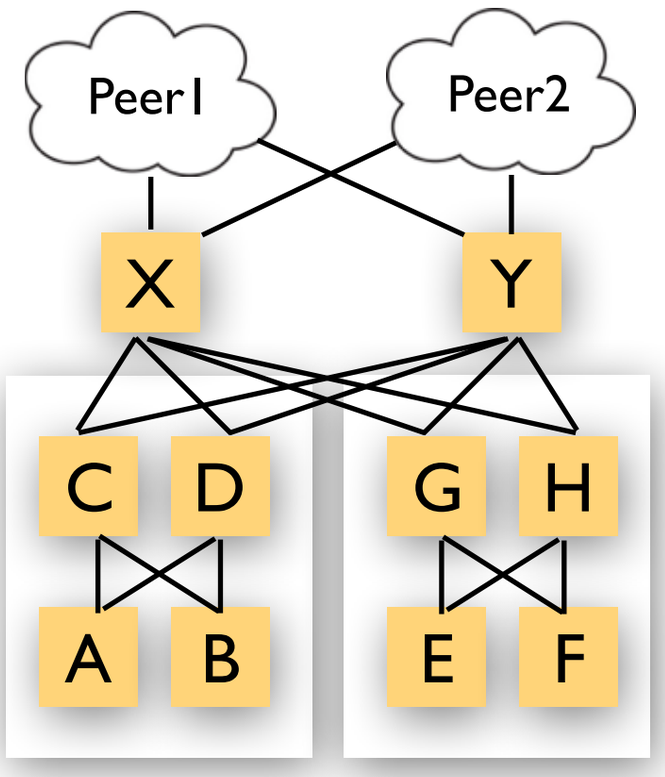

(b) Policy

1)

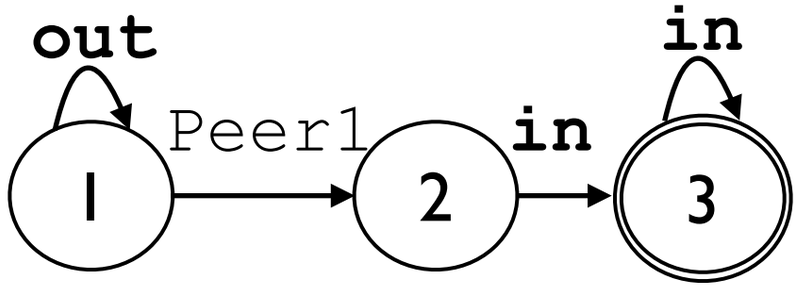

2)

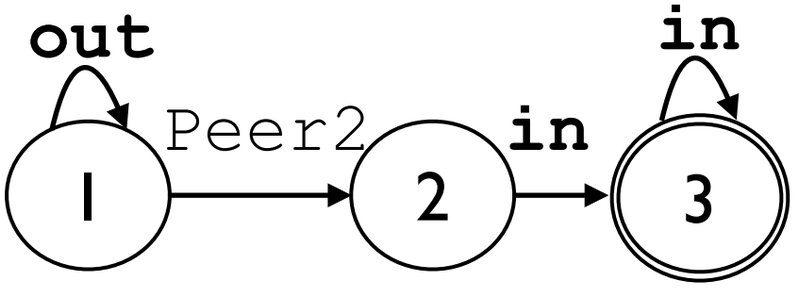

(c) Product Graph

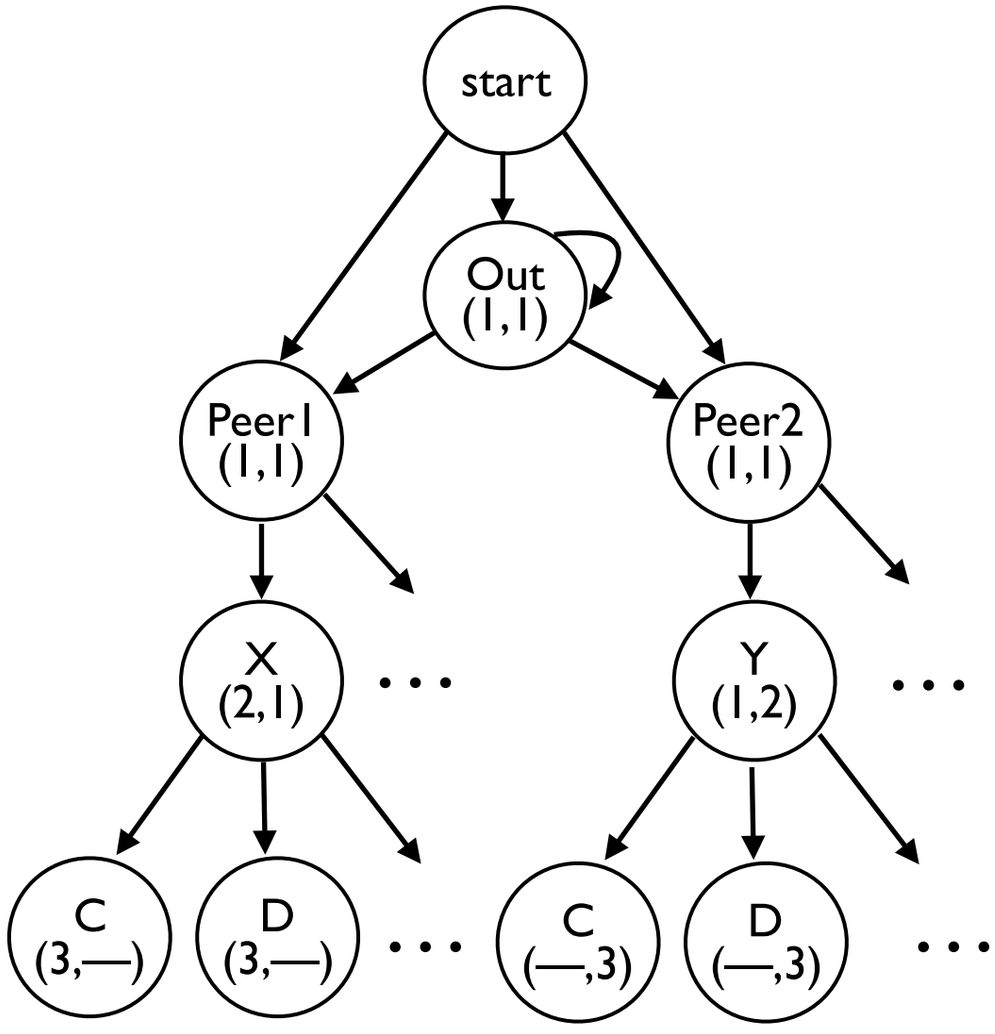

Constructing the Product Graph (PG)

(a) Topology

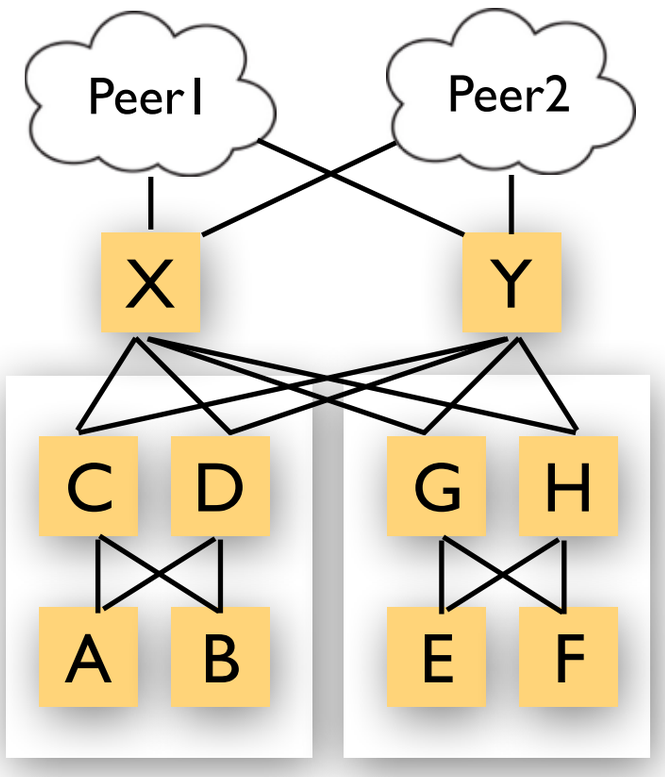

(b) Policy

1)

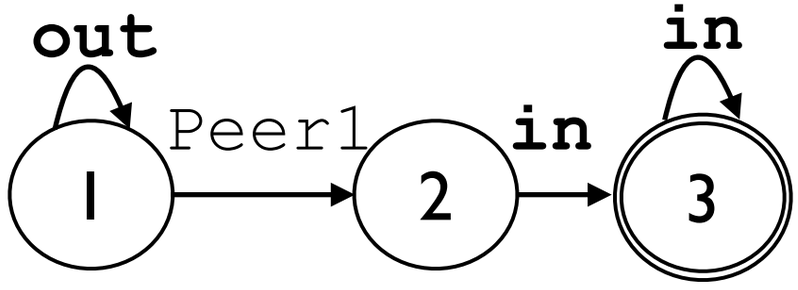

2)

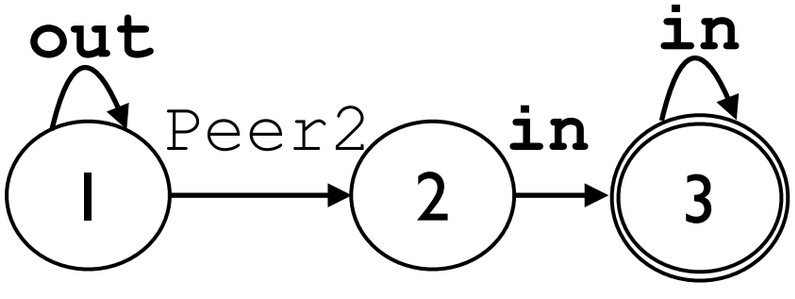

(c) Product Graph

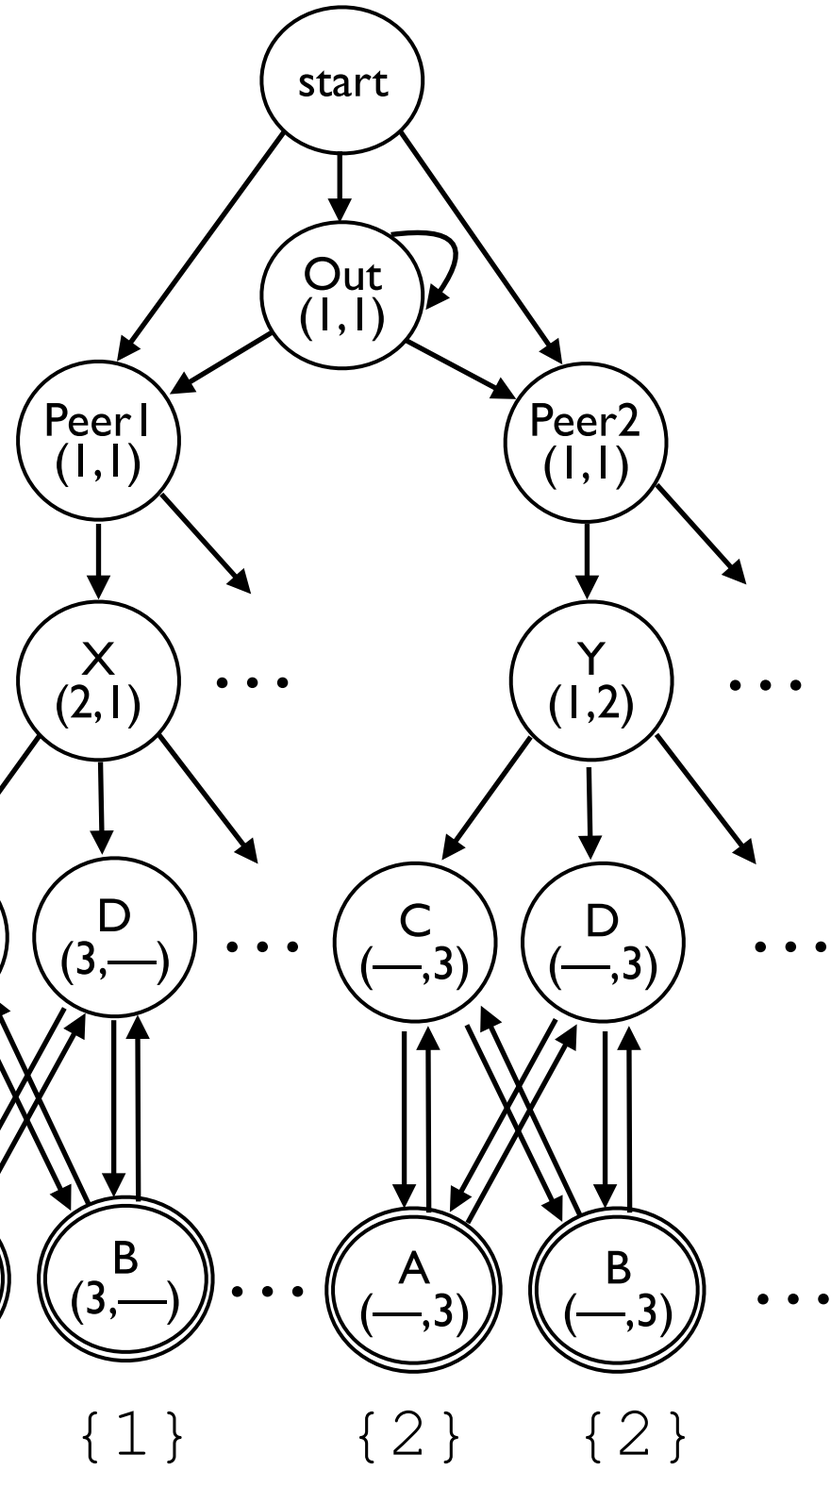

Compilation to BGP:

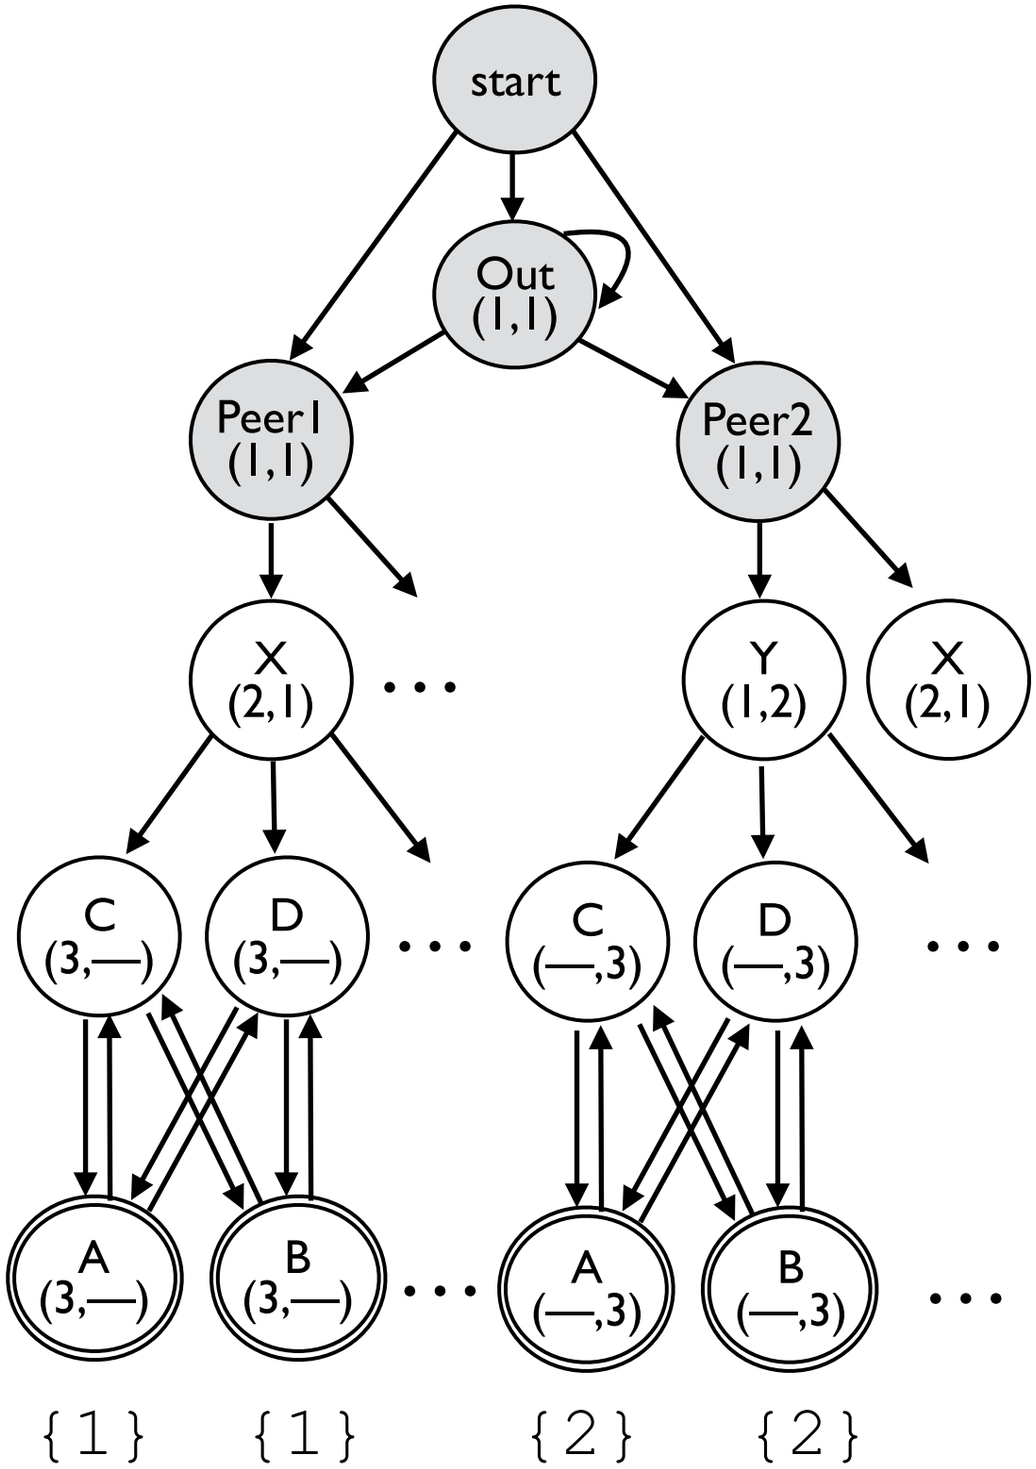

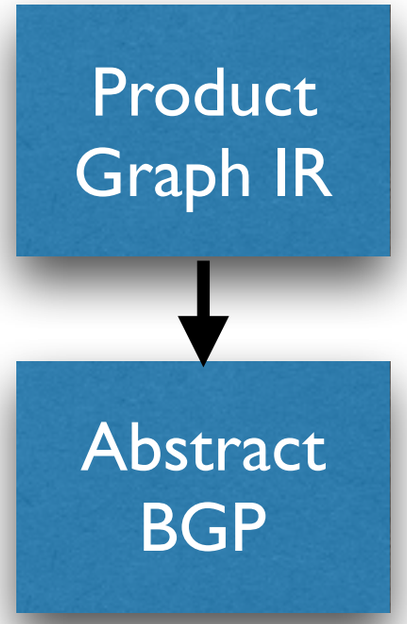

Idea:

- Filter import messages according to incoming PG edges.
- For each internal location, decide which announcements to prefer, forward messages along PG edges
- Use a community value to tag the state of the automata
- For each external location, do nothing

Compilation to BGP:

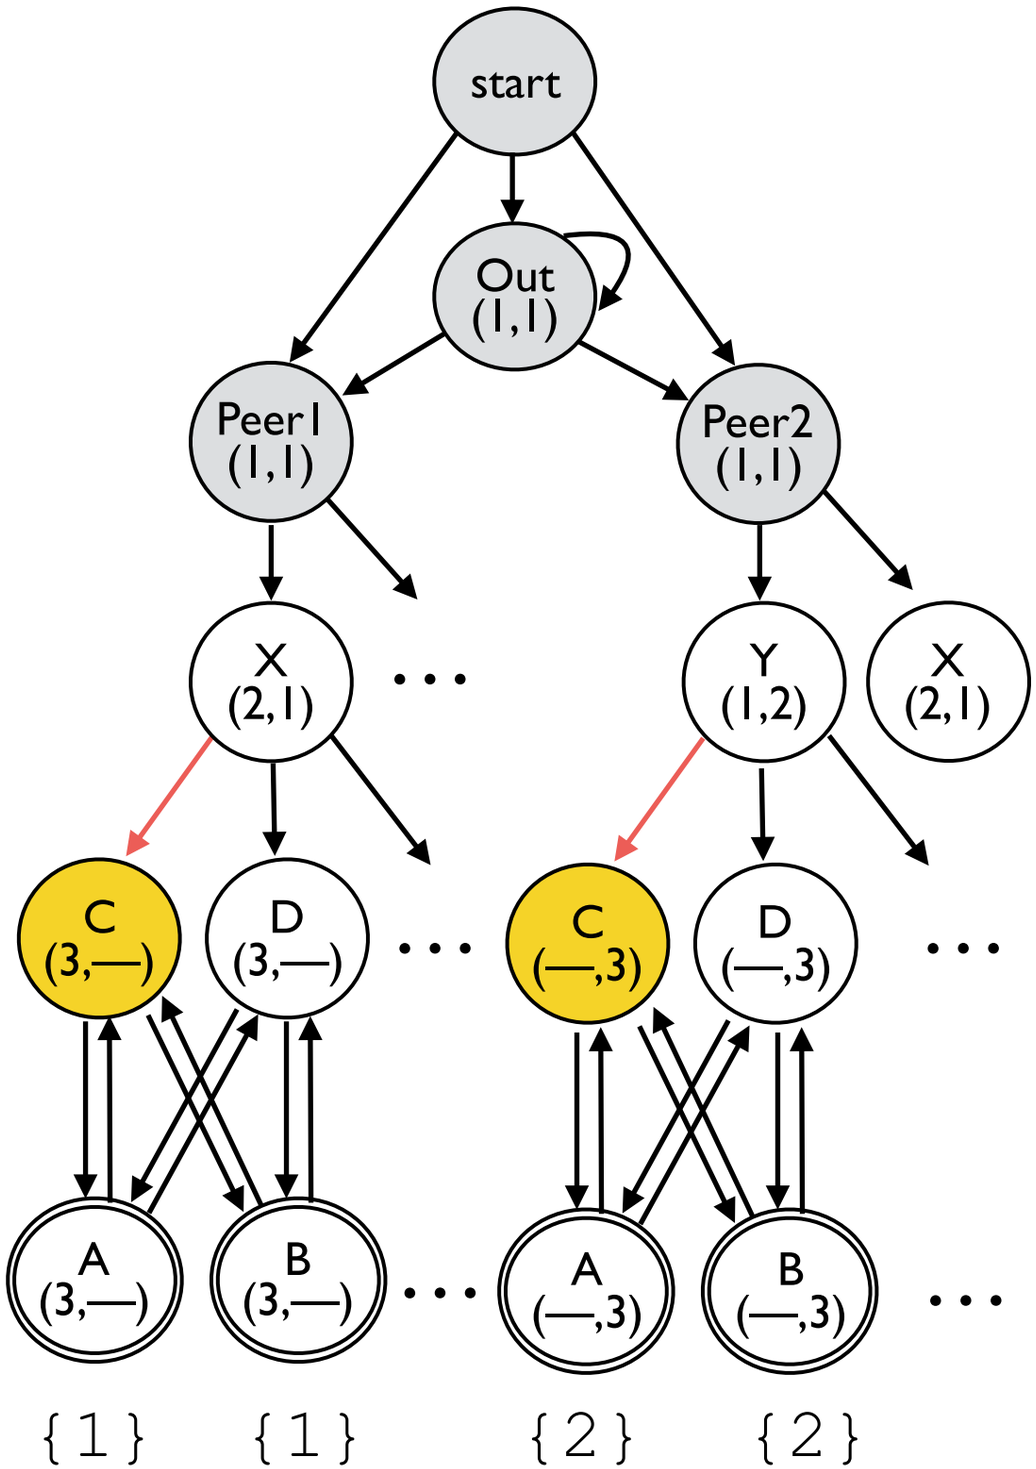

Router C

allow peer=X comm=(2,1)

peer ←{A,B} comm ← (3, —)

allow peer=Y comm=(1,2)

peer ←{A,B} comm ← (—, 3)

Compilation to BGP:

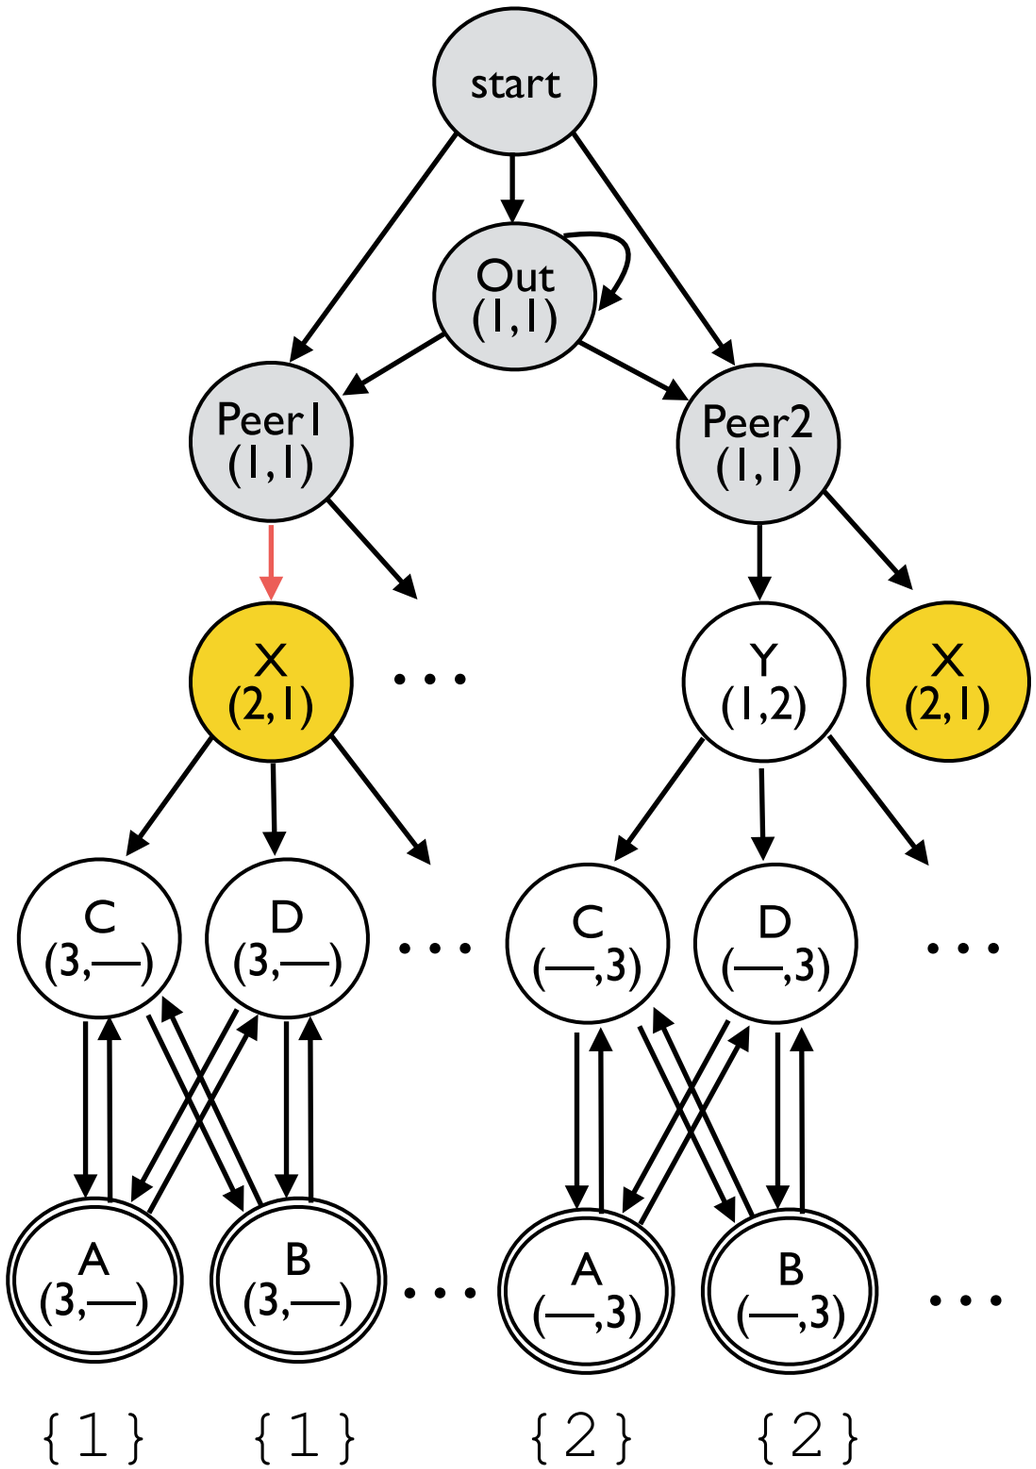

Router C

allow peer=X comm=(2,1)
peer ←{A,B} comm ← (3, —)
allow peer=Y comm=(1,2)
peer ←{A,B} comm ← (—, 3)

Router X

allow regex(Peer1 . out*)
peer ←{C,D,G,H} comm ← (2, 1)
allow regex(Peer2 . out*)
peer ←{C,D,G,H} comm ← (1, 2)

Compilation to BGP:

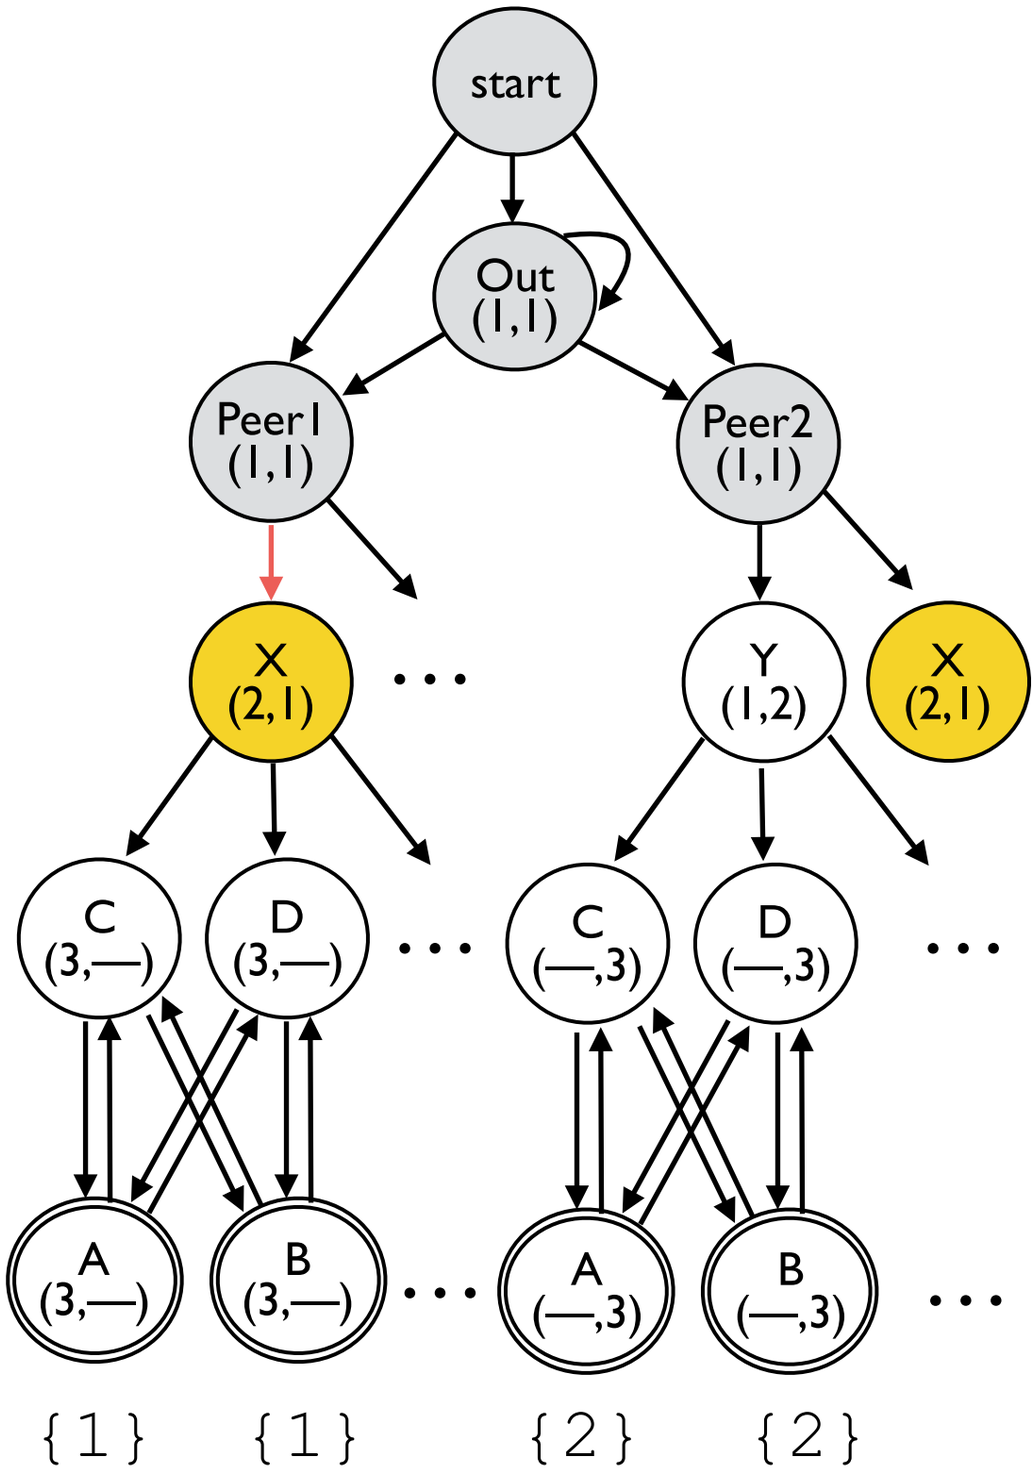

Router C

allow peer=X comm=(2,1)
peer ←{A,B} comm ← (3, —)
allow peer=Y comm=(1,2)
peer ←{A,B} comm ← (—, 3)

Graph Analysis

Router X

allow regex(Peer1 . out*) **with** lp = 100
peer ←{C,D,G,H} comm ← (2, 1)
allow regex(Peer2 . out*) **with** lp = 99
peer ←{C,D,G,H} comm ← (1, 2)

Implementation (5,500 lines of F#)

Benchmarks:

- data center policies (~1600 routers) and backbone policies (~200 routers, many peers/router) from a large cloud provider
- policy from English docs
- Ignoring prefix, customer group and ownership definitions:
  - 31 lines for data center
  - 43 lines for backbone

Scaling (8 core Windows machine):

- 10s/pfx (mean) for largest data center
- 45s/pfx (mean) for largest backbone
- 3 minutes total for the backbone
- 9 minutes total for the data center

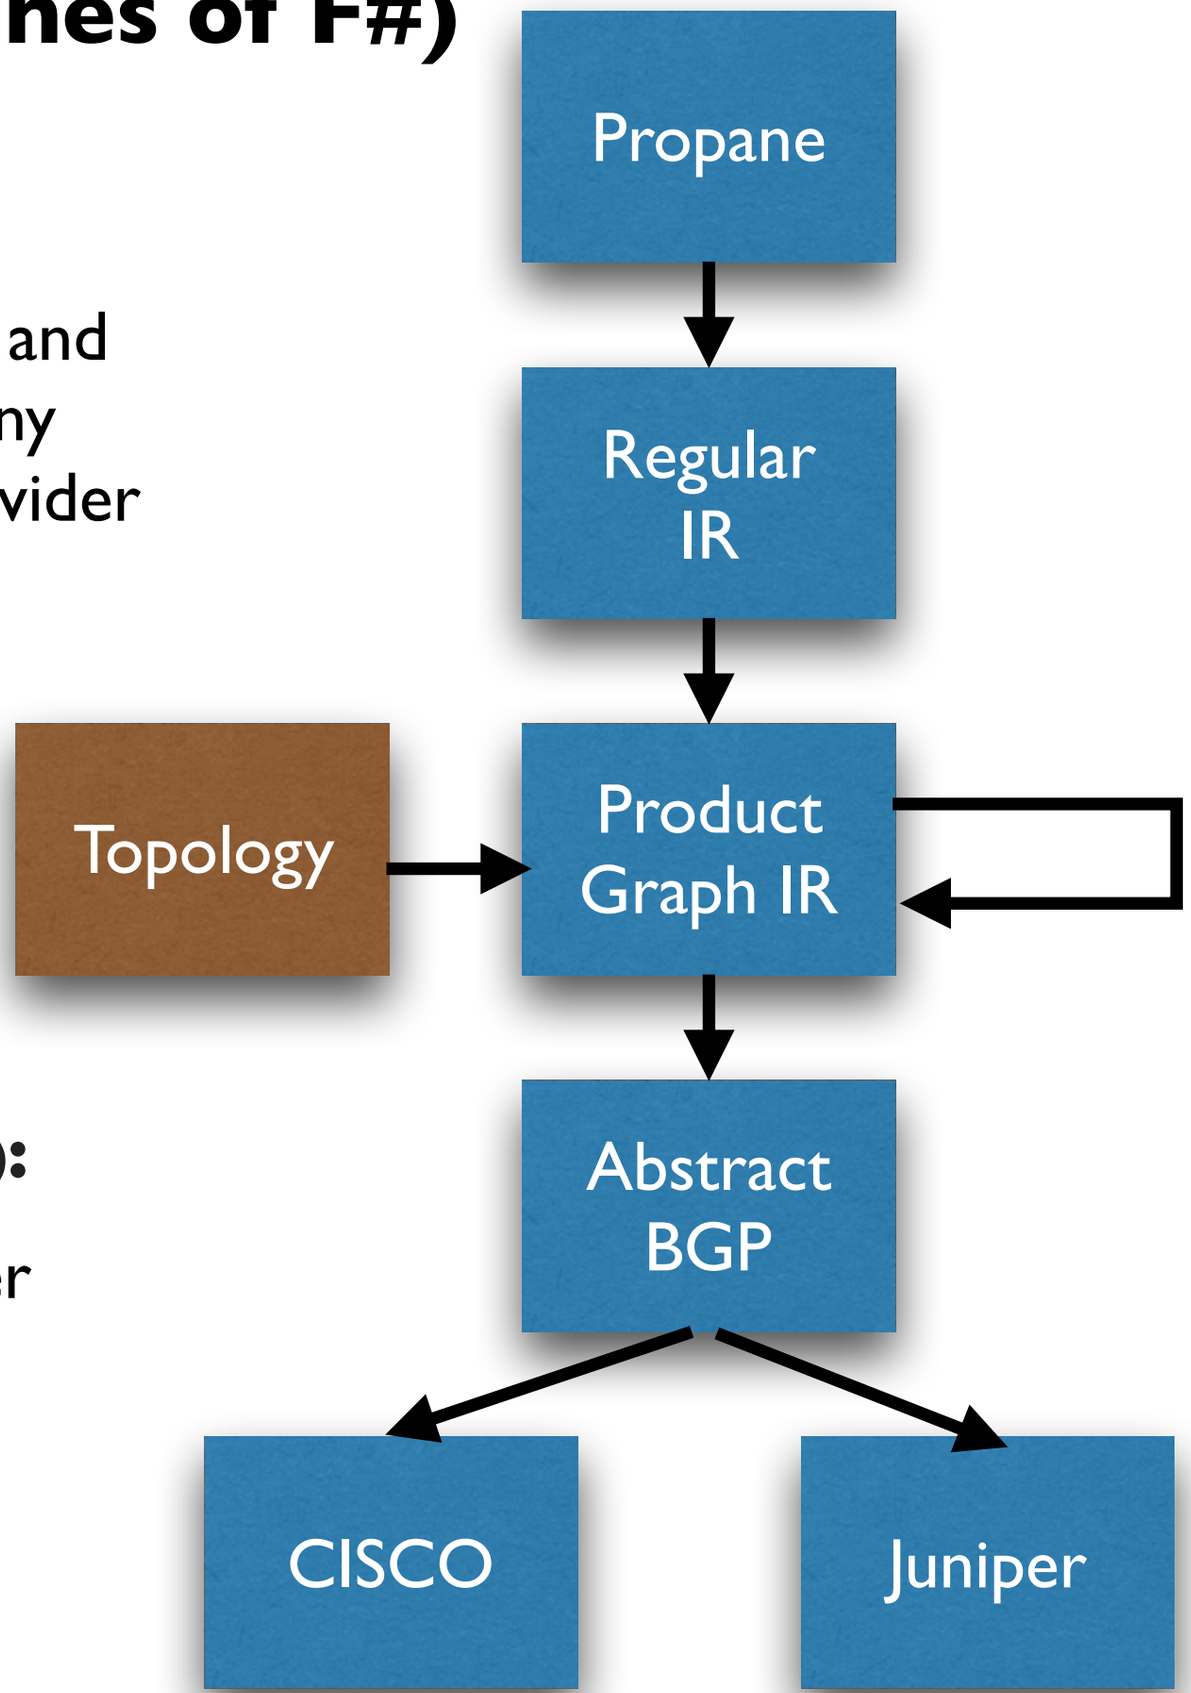

Thanks